

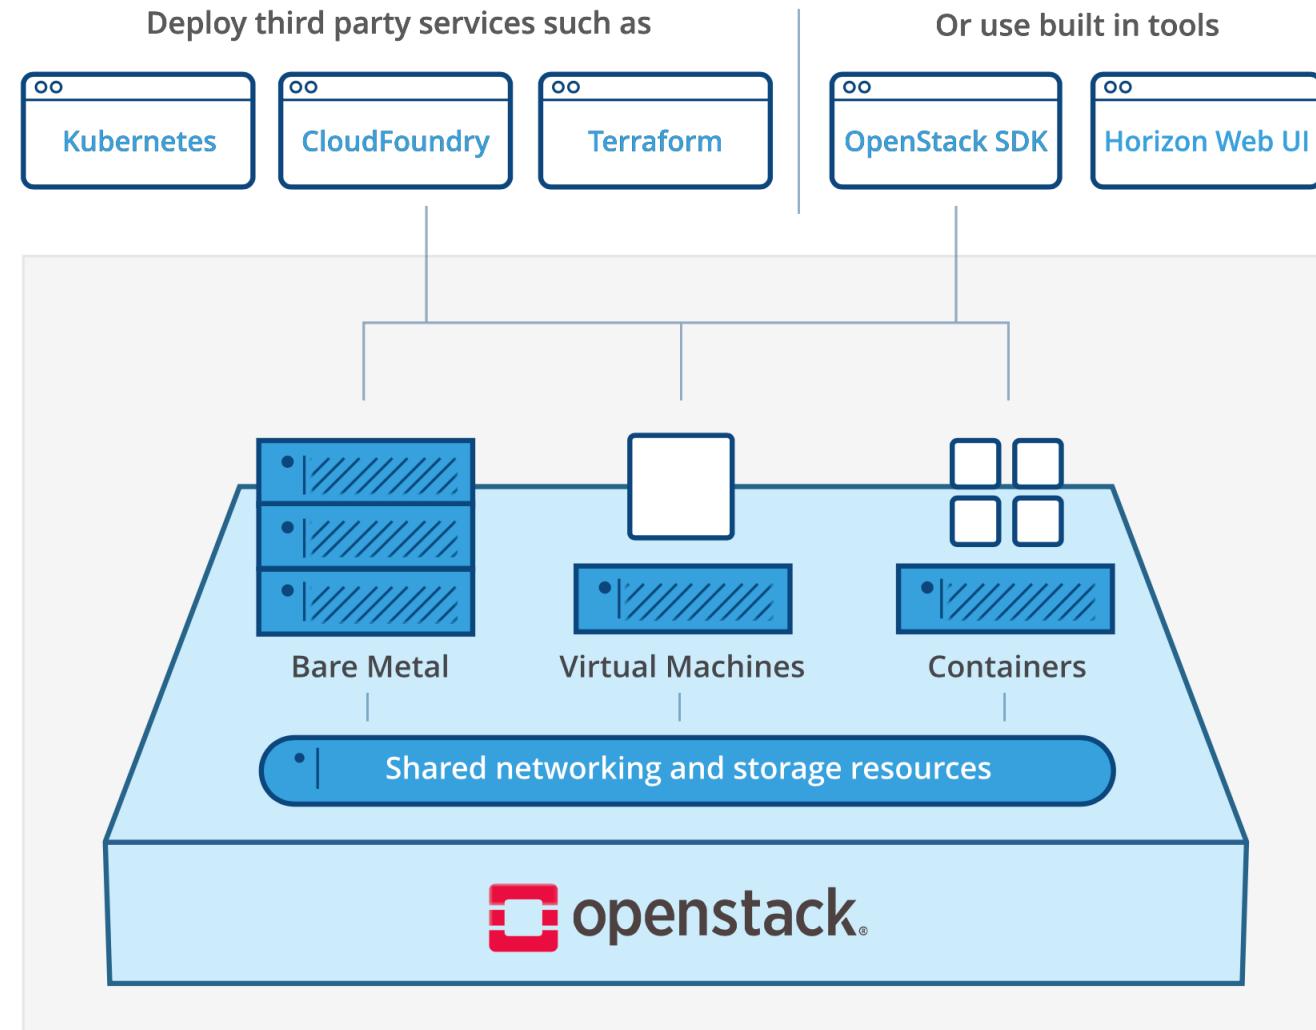
# Introdução ao Openstack

Aula Teórica nº2

2020/2021

# Open source SW for creating private and public clouds

- controls large pools of compute, storage, and networking resources throughout a datacenter
- managed through a [dashboard](#) or via the [OpenStack API](#)
- Built by a growing community of contributors
- Opportunities for vendors to develop their own solutions and services



# OpenStack Who's

- Who's building OpenStack?
  - Originally founded by Rackspace Hosting and NASA in July 2010
  - Adopted by Ubuntu in 2011
  - Now driven by OpenStack Foundation (established in September, 2012)
- Who's backing OpenStack?
  - More than 200 companies - 87000+ people from 180 countries
- Who's using OpenStack?
  - Enterprise, service providers, SMBs, researchers...
  - In private public and Telco cloud

# Companies Supporting The OpenStack Foundation

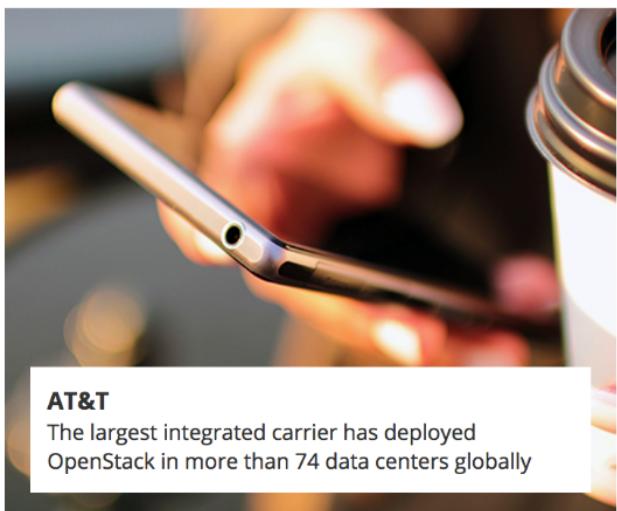
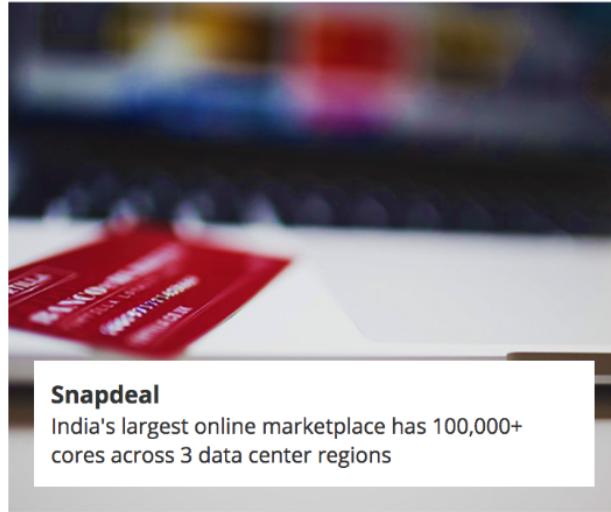
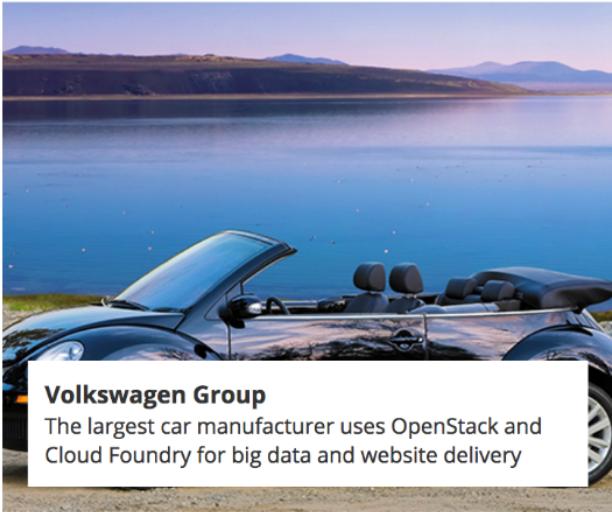
## Platinum members



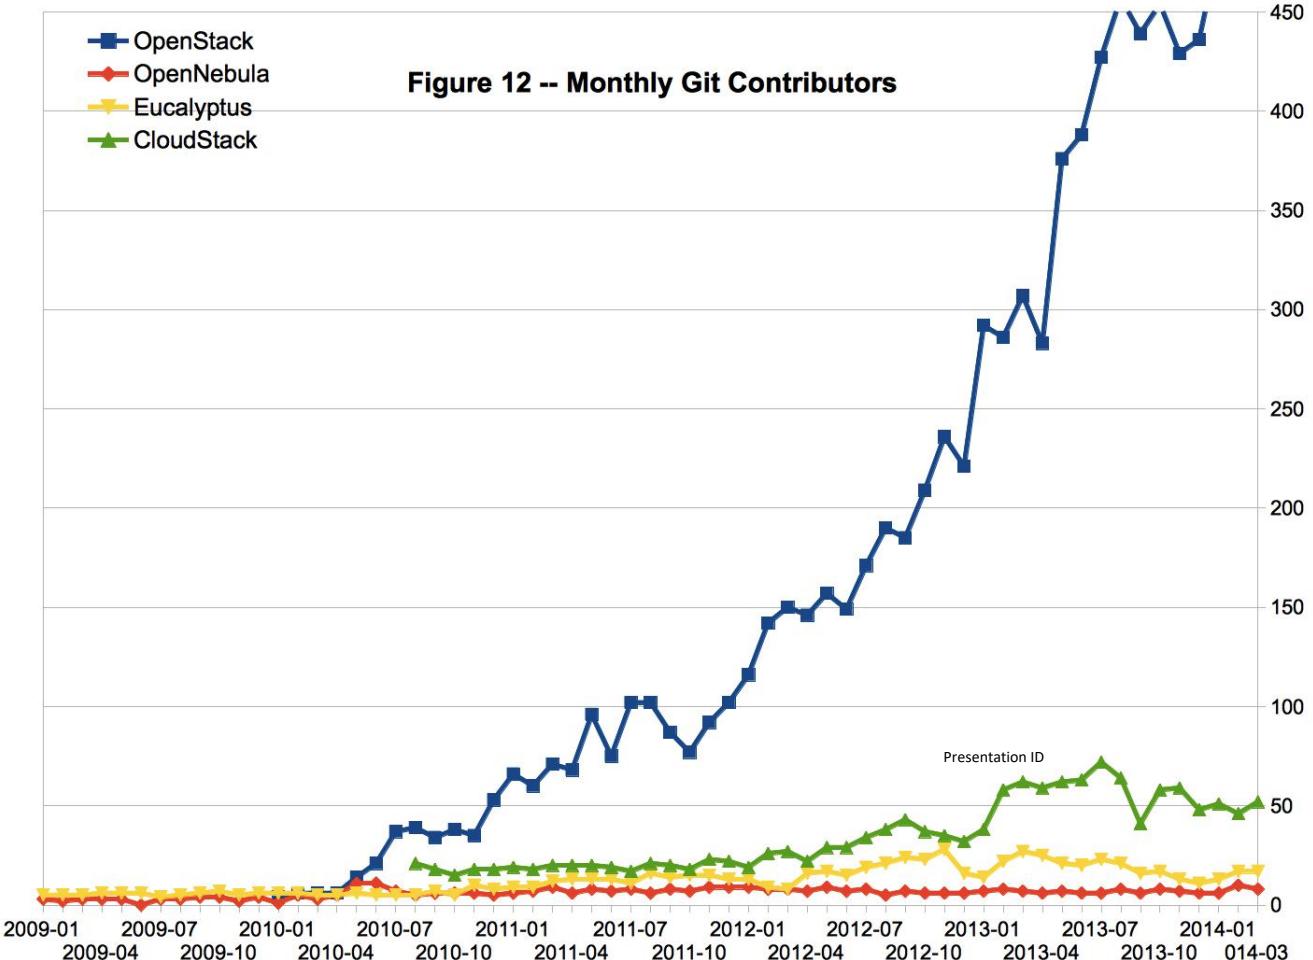
## Gold members



# OpenStack - Industries



# Open Source Cloud SW Analysis



## Initial Release:

- OpenStack, October 2010
- CloudStack, May 2010
- Eucalyptus, May 2008
- OpenNebula, March 2008

~ 10 years

## Latest Releases:

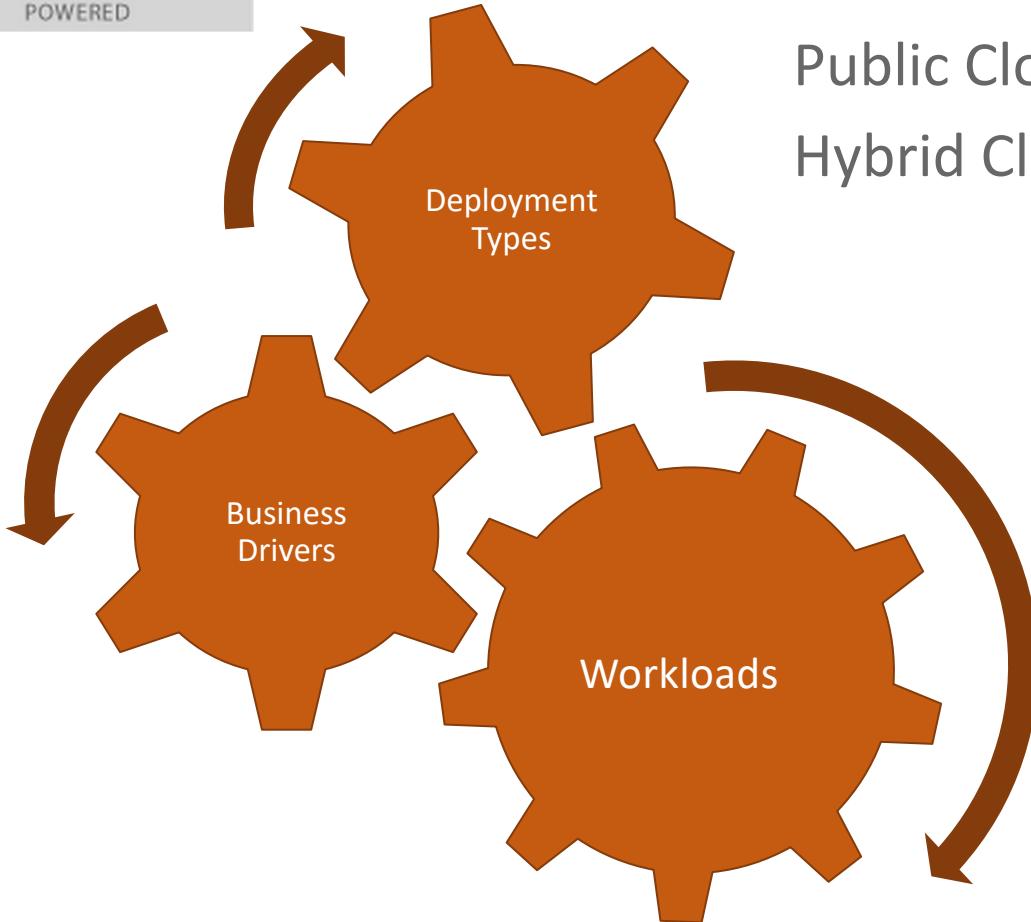
- Openstack, October 2020
- CloudStack, January 2021
- Eucalyptus, April 2018
- OpenNebula, February 2021

# OpenStack Releases

Series	Status	Initial Release Date	Series	Status	Initial Release Date	Next Phase	EOL Date
Xena	<a href="#">Future</a>	2021-10-06 <i>estimated (schedule)</i>	Newton	<a href="#">End Of Life</a>	2016-10-06		2017-10-25
Wallaby	<a href="#">Development</a>	2021-04-14 <i>estimated (schedule)</i>	Mitaka	<a href="#">End Of Life</a>	2016-04-07		2017-04-10
Victoria	<a href="#">Maintained</a>	2020-10-14	Liberty	<a href="#">End Of Life</a>	2015-10-15		2016-11-17
Ussuri	<a href="#">Maintained</a>	2020-05-13	Kilo	<a href="#">End Of Life</a>	2015-04-30		2016-05-02
Train	<a href="#">Maintained</a>	2019-10-16	Juno	<a href="#">End Of Life</a>	2014-10-16		2015-12-07
Stein	<a href="#">Extended Maintenance (see <a href="#">note</a> below)</a>	2019-04-10	Icehouse	<a href="#">End Of Life</a>	2014-04-17		2015-07-02
Rocky	<a href="#">Extended Maintenance (see <a href="#">note</a> below)</a>	2018-08-30	Havana	<a href="#">End Of Life</a>	2013-10-17		2014-09-30
Queens	<a href="#">Extended Maintenance (see <a href="#">note</a> below)</a>	2018-02-28	Grizzly	<a href="#">End Of Life</a>	2013-04-04		2014-03-29
Pike	<a href="#">Extended Maintenance (see <a href="#">note</a> below)</a>	2017-08-30	Folsom	<a href="#">End Of Life</a>	2012-09-27		2013-11-19
Ocata	<a href="#">Extended Maintenance (see <a href="#">note</a> below)</a>	2017-02-22	Essex	<a href="#">End Of Life</a>	2012-04-05		2013-05-06
			Diablo	<a href="#">End Of Life</a>	2011-09-22		2013-05-06
			Cactus	<a href="#">End Of Life</a>	2011-04-15		
			Bexar	<a href="#">End Of Life</a>	2011-02-03		
			Austin	<a href="#">End Of Life</a>	2010-10-21		

# OpenStack Trends

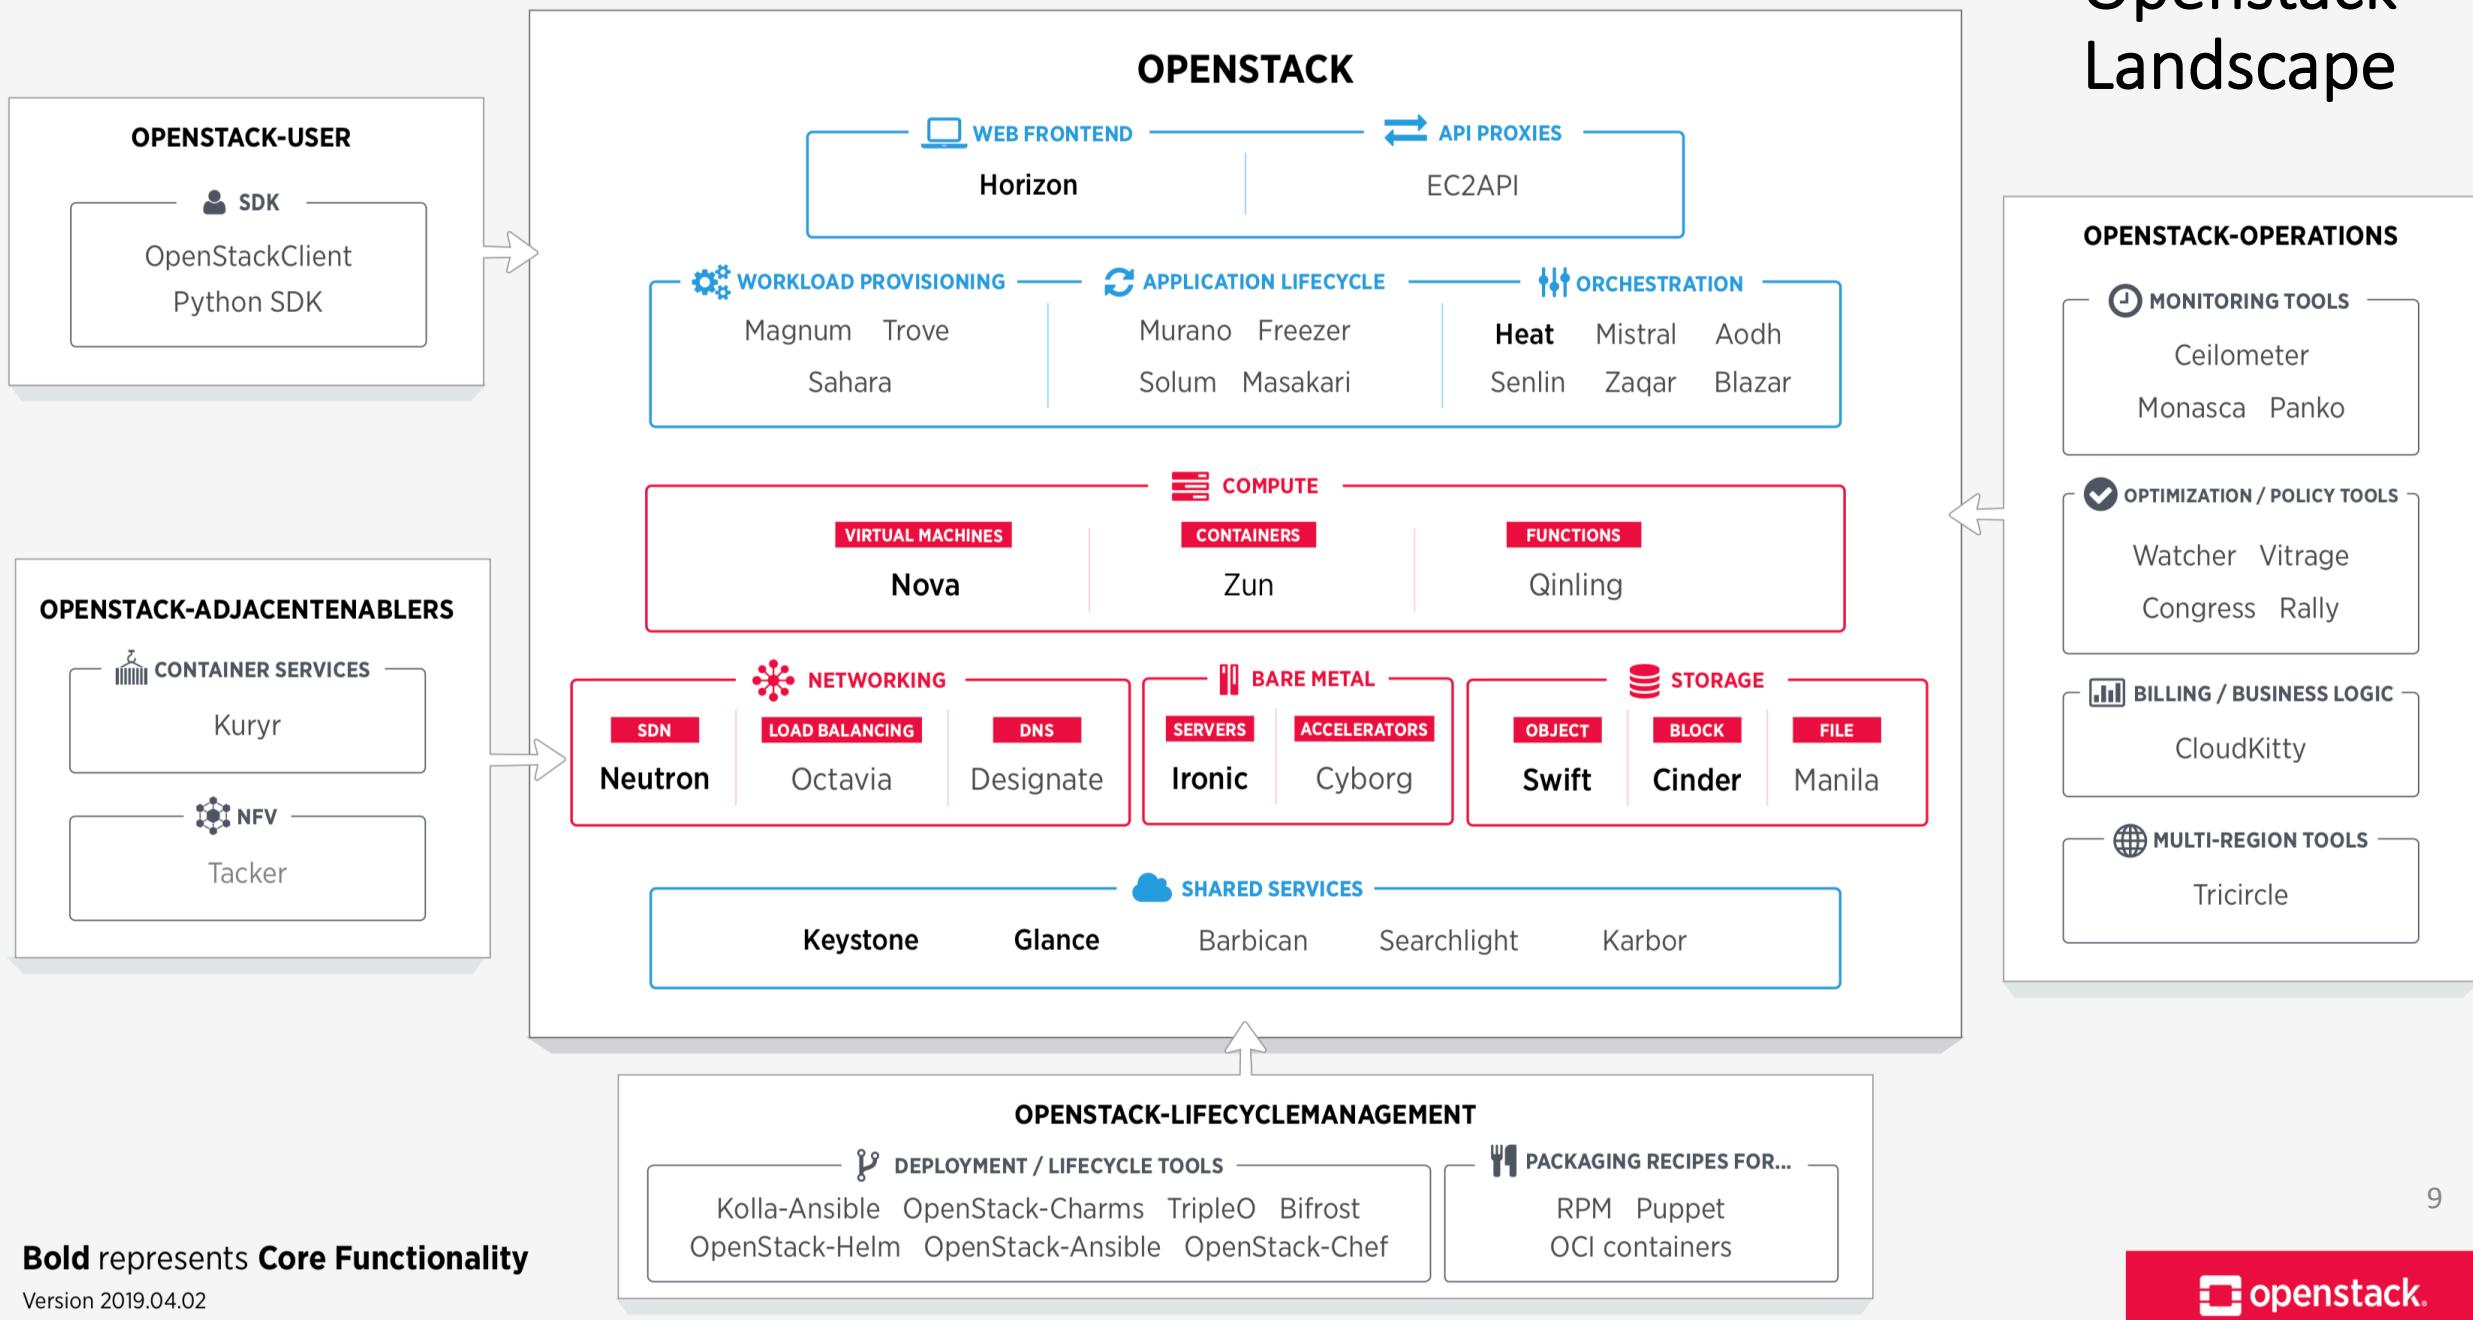
Ability to Innovate  
Open Technology  
Cost Savings  
Avoiding vendor lock-in  
Operational efficiency



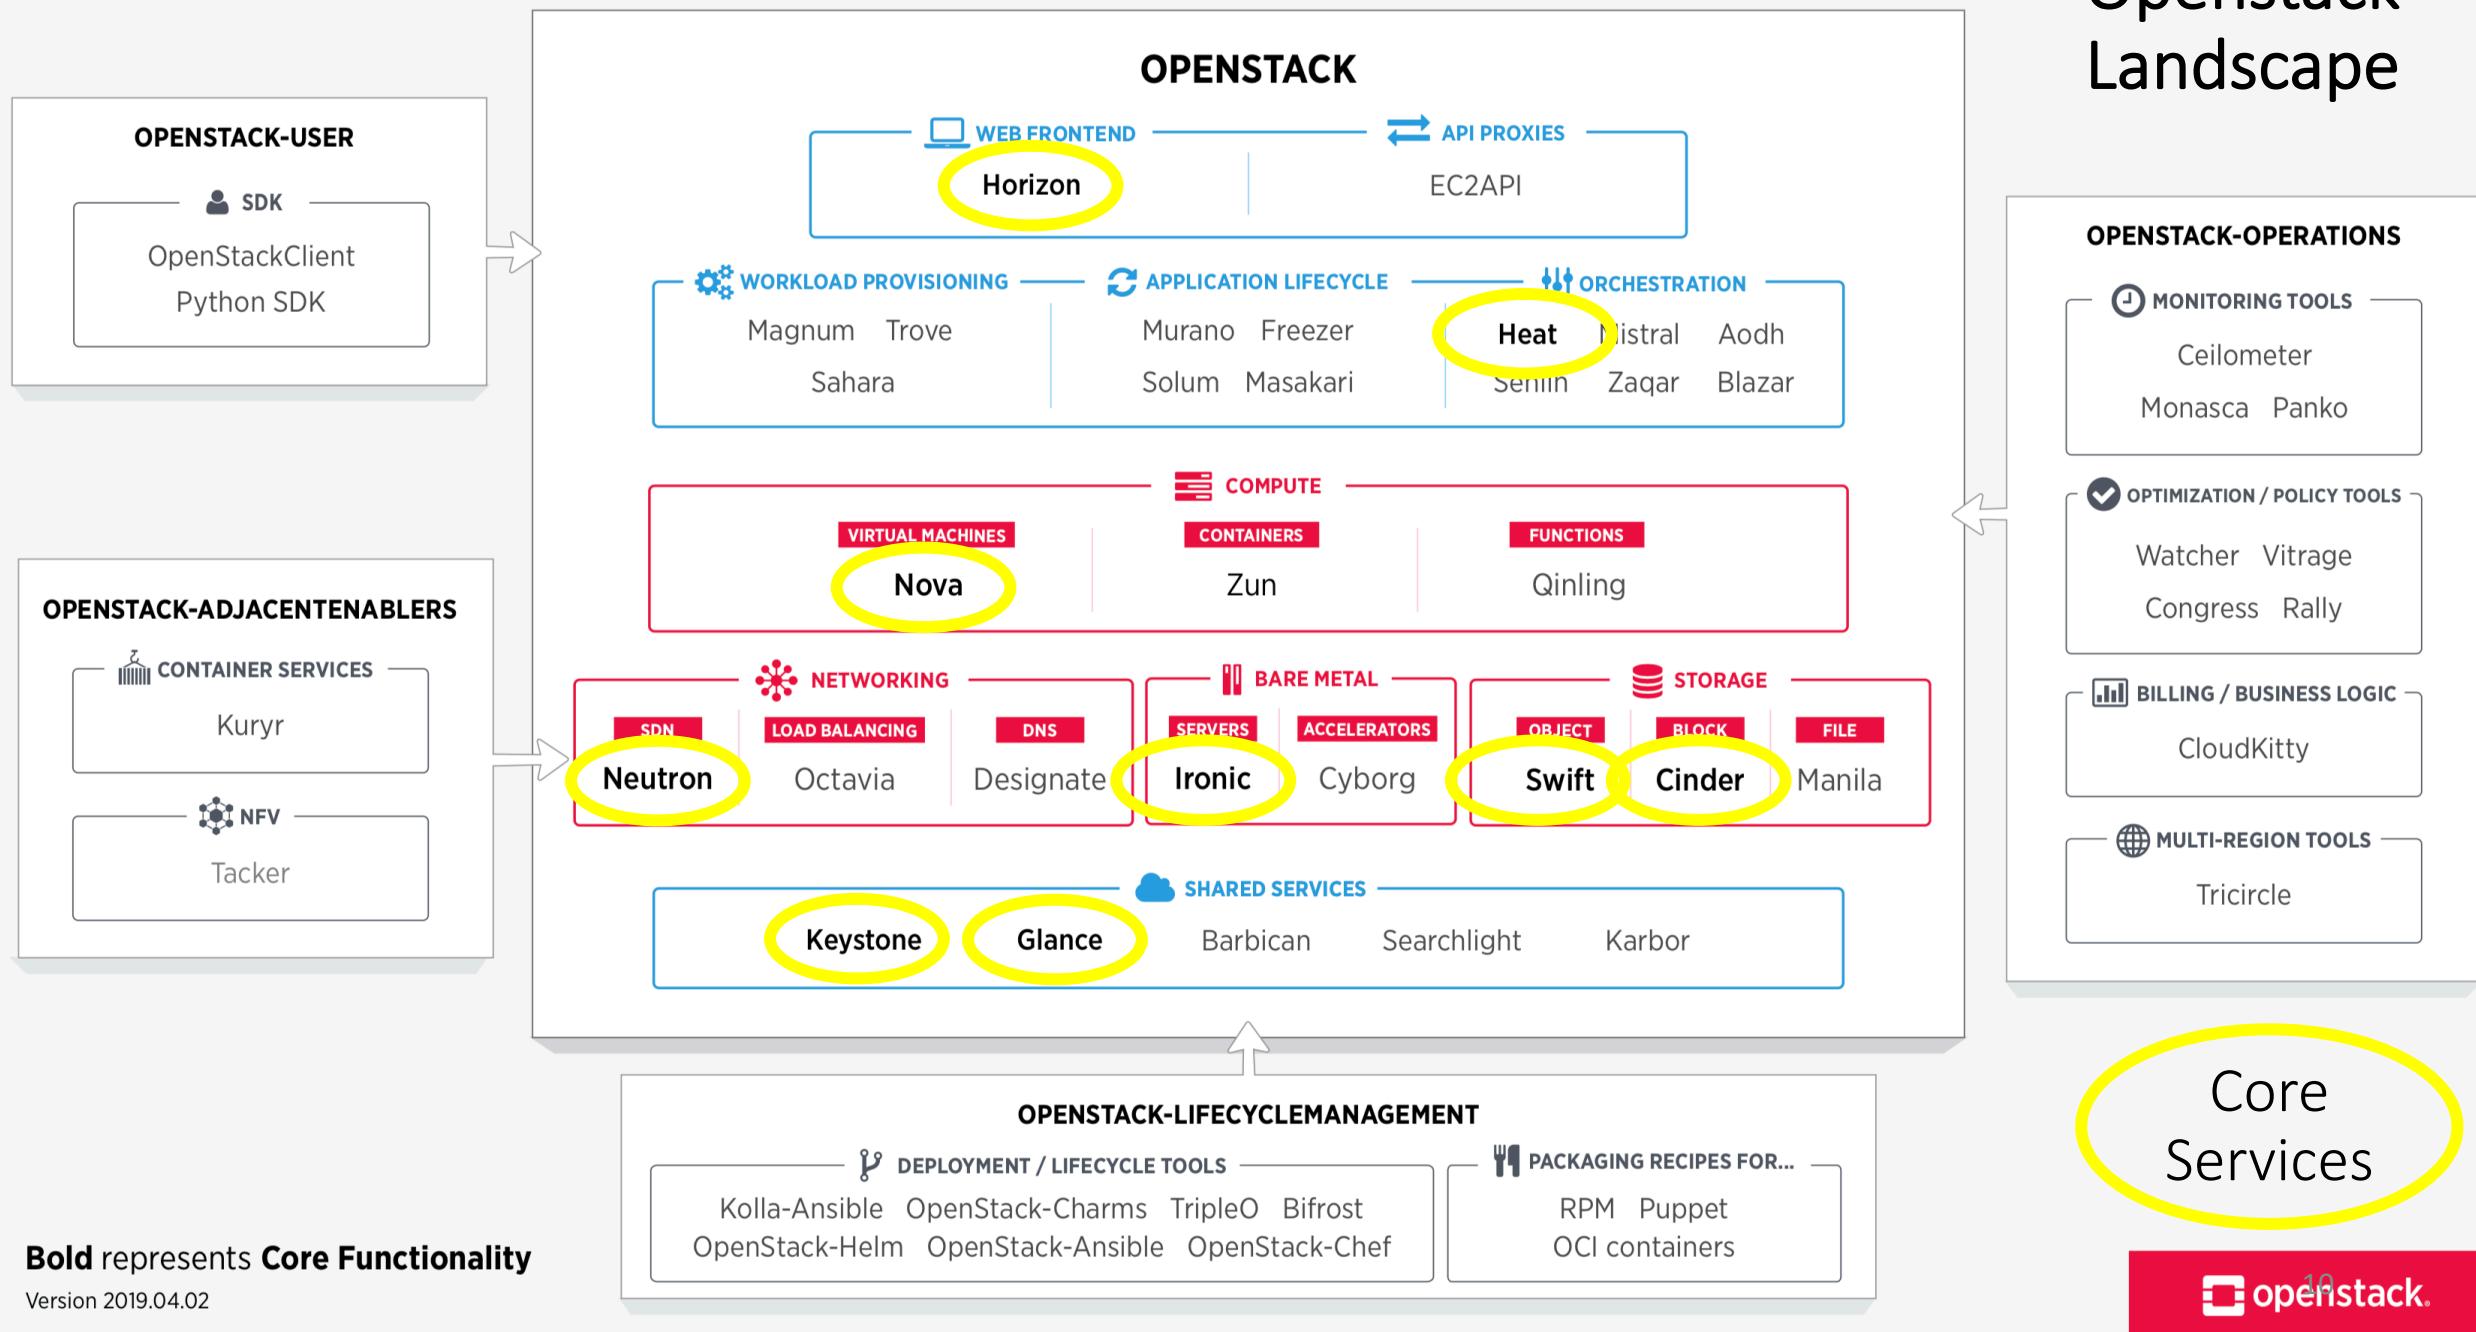
On-Premise Private Cloud  
Hosted Private Clouds  
Public Cloud  
Hybrid Cloud

Web Services  
QA/Test environment  
Databases/Datastores  
Enterprise Applications  
Research/HPC

# Openstack Landscape



# Openstack Landscape



# OpenStack Core Services



Nova

**Austin, 2010**

Compute provisioning service



Keystone

**Essex, 2012**

Identity (client authentication) service



Cinder

**Folsom, 2012**

Block storage service (LVM or plugin drivers for other storage)



Swift

**Austin, 2010**

Highly available, distributed object/BLOB store service



Horizon

**Essex, 2012**

web based user interface (dashboard) to OpenStack services



Heat

**Havana, 2013**

Heat orchestrates the infrastructure resources based on templates



Glance

**Bexar, 2011**

Image services include discovering, registering, and retrieving virtual machine images



Neutron

**Folsom, 2012**

SDN networking project focused on delivering networking-as-a-service (NaaS)



Ironic

**Kilo, 2015**

Bare metal provisioning service

# OpenStack Services

## Compute

	<b>NOVA</b>	Compute Service
	<b>ZUN</b>	Containers Service
	<b>QINLING</b>	Functions Service

## Bare Metal

	<b>IRONIC</b>	Bare Metal Provisioning Service
	<b>CYBORG</b>	Accelerators resource management

## Storage

	<b>SWIFT</b>	Object store
	<b>CINDER</b>	Block Storage
	<b>MANILA</b>	Shared filesystems

## Networking

	<b>NEUTRON</b>	Networking
	<b>OCTAVIA</b>	Load balancer
	<b>DESIGNATE</b>	DNS service

## Orchestration

	<b>HEAT</b>	Orchestration
	<b>SENLIN</b>	Clustering service
	<b>MISTRAL</b>	Workflow service
	<b>ZAQAR</b>	Messaging Service
	<b>BLAZAR</b>	Resource reservation service
	<b>AODH</b>	Alarming Service

## Workload Provisioning

	<b>MAGNUM</b>	Container Orchestration Engine Provisioning
	<b>SAHARA</b>	Big Data Processing Framework Provisioning
	<b>TROVE</b>	Database as a Service

## Application Lifecycle

	<b>MASAKARI</b>	Instances High Availability Service
	<b>MURANO</b>	Application Catalog
	<b>SOLUM</b>	Software Development Lifecycle Automation
	<b>FREEZER</b>	Backup, Restore, and Disaster Recovery

## Monitoring Tools

	<b>CEILOMETER</b>	Metering & Data Collection Service
	<b>PANKO</b>	Event, Metadata Indexing Service
	<b>MONASCA</b>	Monitoring

## Optimization/policy tools

	<b>WATCHER</b>	Optimization Service
	<b>VITRAGE</b>	Root Cause Analysis service
	<b>CONGRESS</b>	Governance
	<b>RALLY</b>	Benchmark service

## Billing / Business Logic

	<b>CLOUDKITTY</b>	Billing and chargebacks
---	-------------------	-------------------------

## Multi-Region Tools

	<b>TRICIRCLE</b>	Networking Automation for Multi-Region Deployments
---	------------------	--

## Containers

	<b>KURYR</b>	OpenStack Networking integration for containers
---	--------------	---

## NFV

	<b>TACKER</b>	NFV Orchestration
---	---------------	-------------------

## API Proxies

	<b>EC2API</b>	EC2 API proxy
---	---------------	---------------

## Web Frontend

	<b>HORIZON</b>	Dashboard
---	----------------	-----------

## Swift add-ons

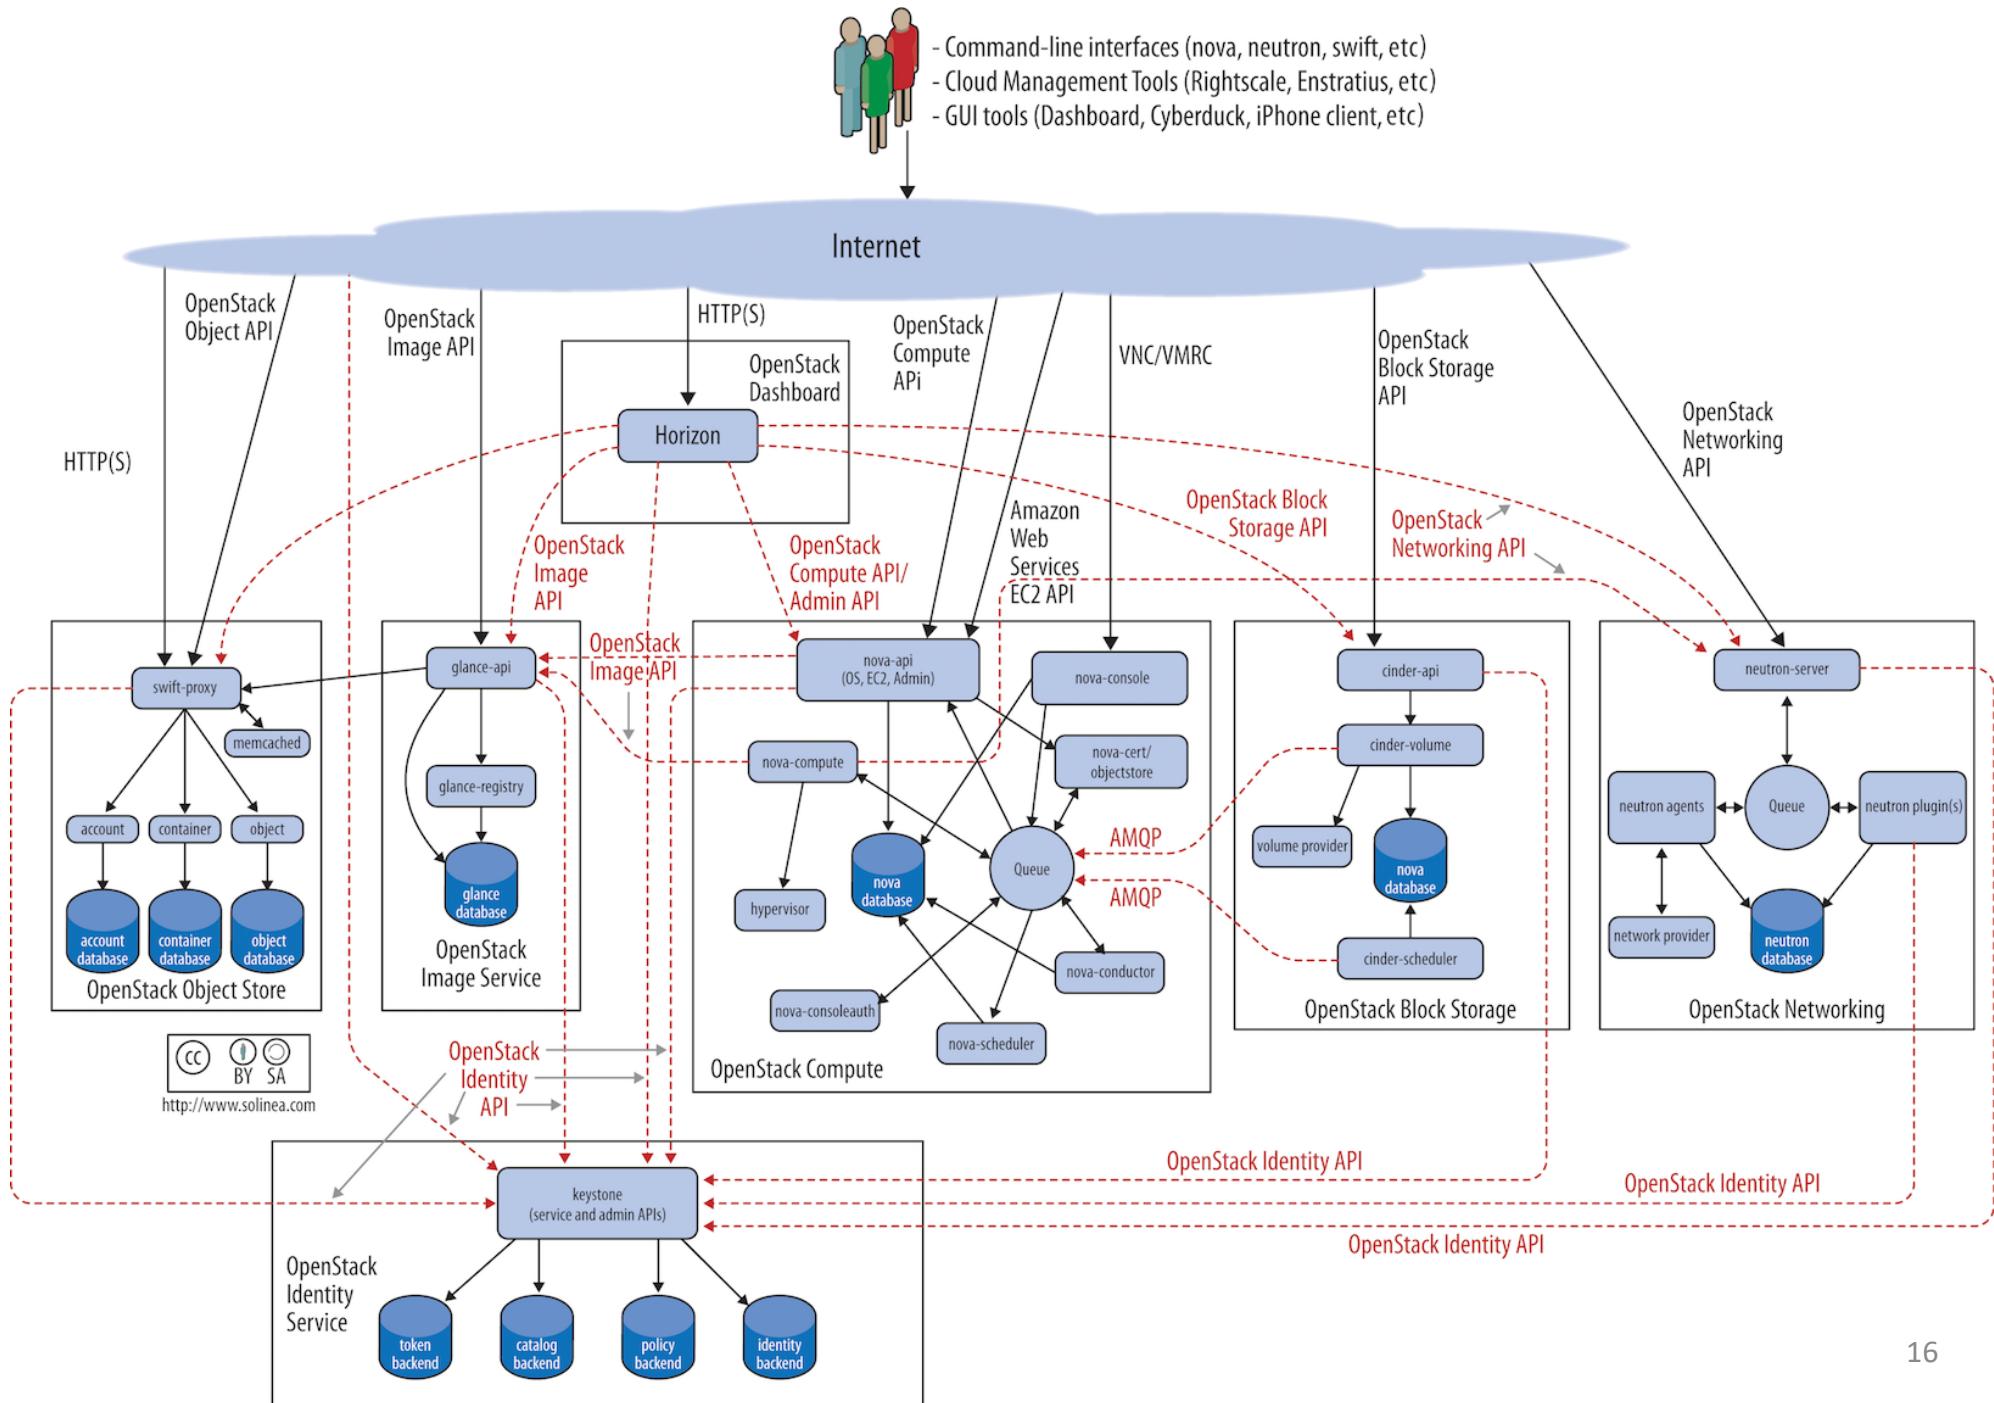
	<b>STORLETS</b>	Computable object storage
---	-----------------	---------------------------

OpenStack  
Services

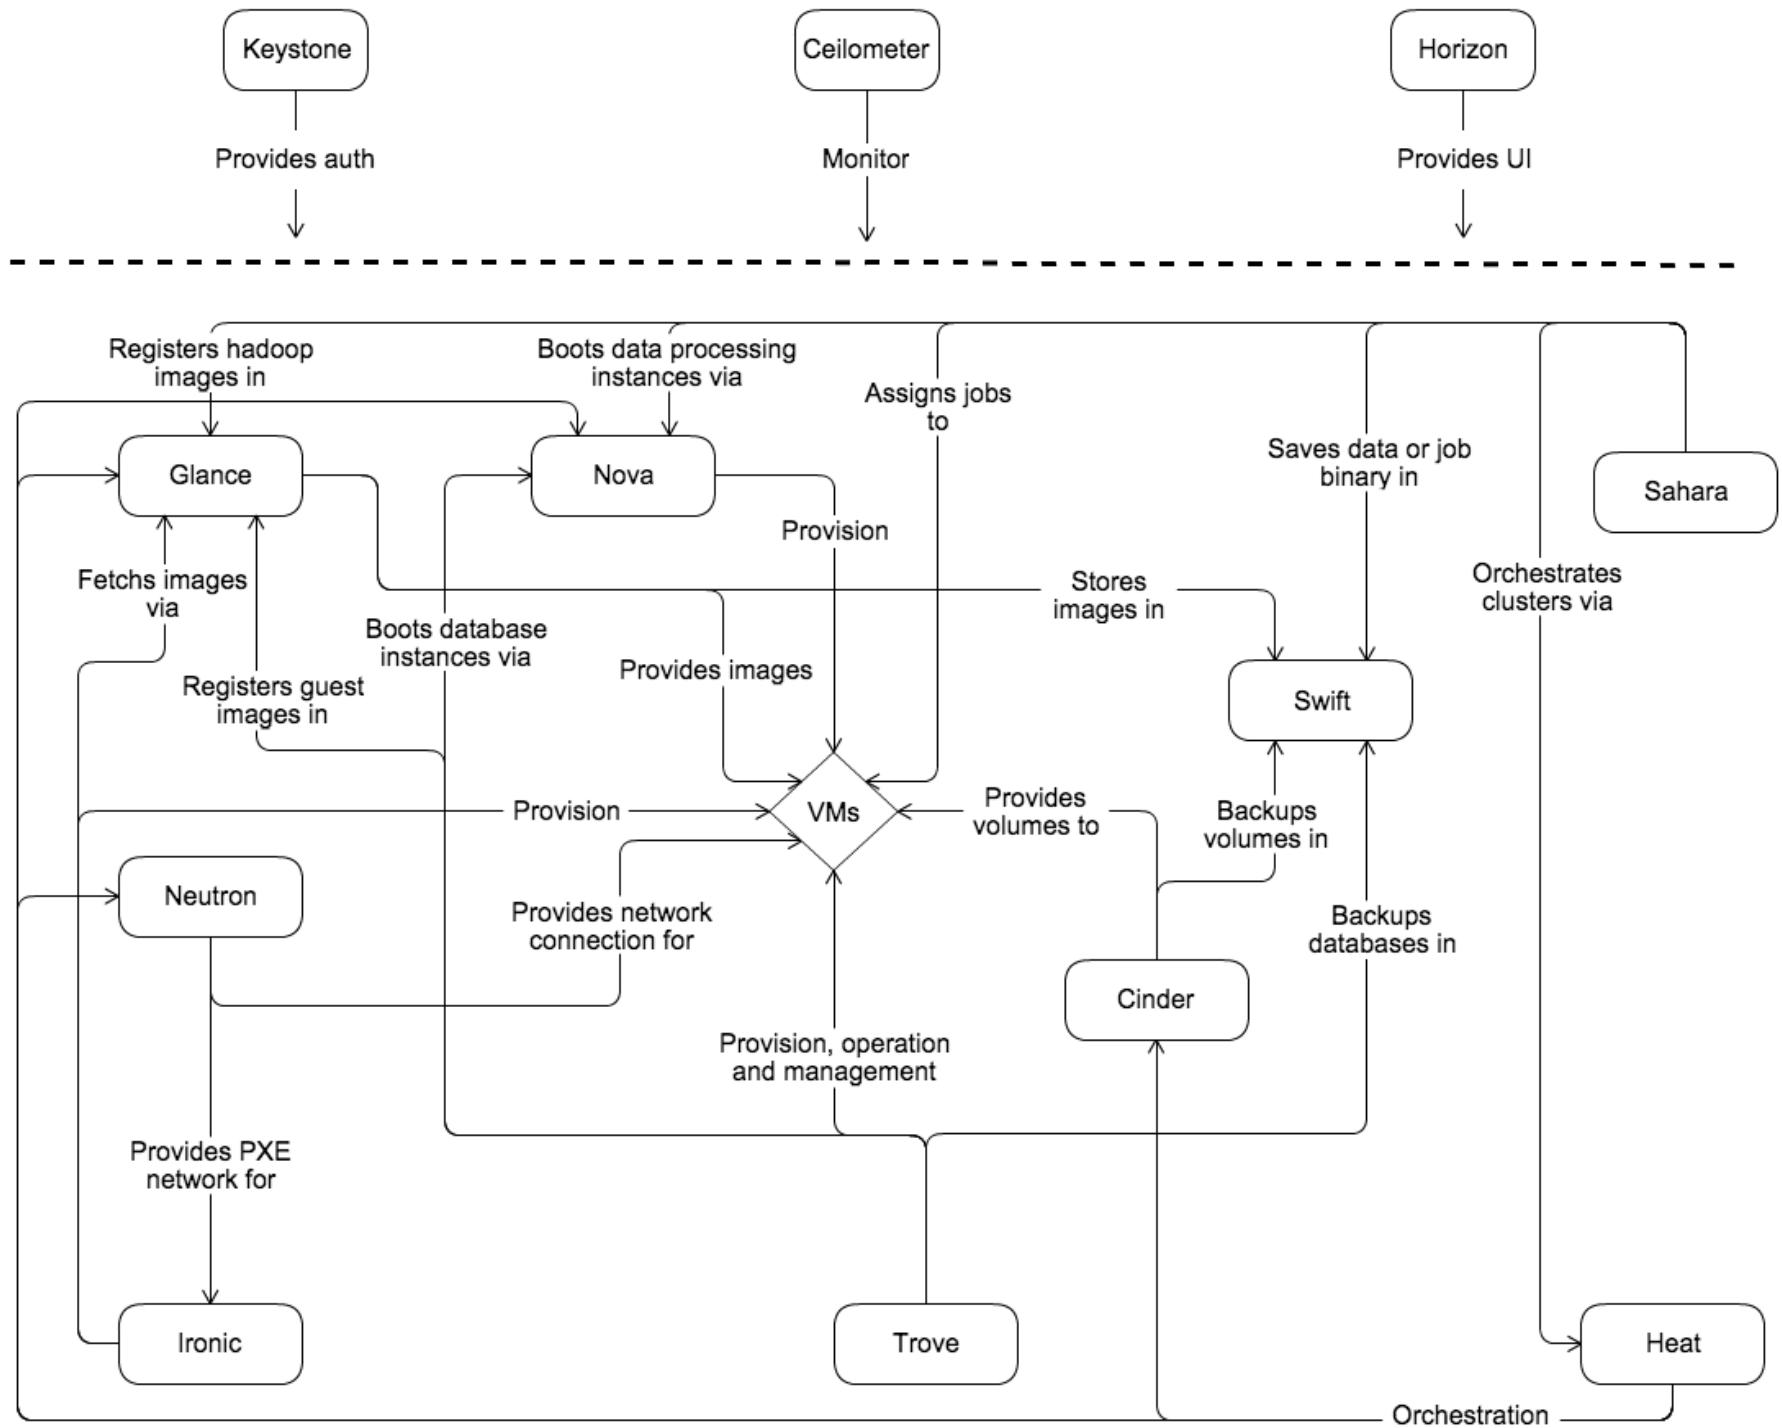
# OpenStack Terminology

- **Instance** – Virtual Machine
- **Image** – Operating system in a file that is used to launch virtual machine, multiple formats (qcow2, ISO, etc.)
- **Application Programming Interface (API)** – Interface (routines, protocols, tools) for building/accessing an application
- **Message Queue** – Hub for passing messages between daemons/services
- **Volume** – Persistent block storage for instances
- **Project ("Tenant")** – Provides logical separation among OpenStack users
- **Flavors** – Pre-created bundles of computing resources
- **Fixed IP** – Associated to an instance on start-up (internal IP)
- **Floating IP** – Public facing IP address

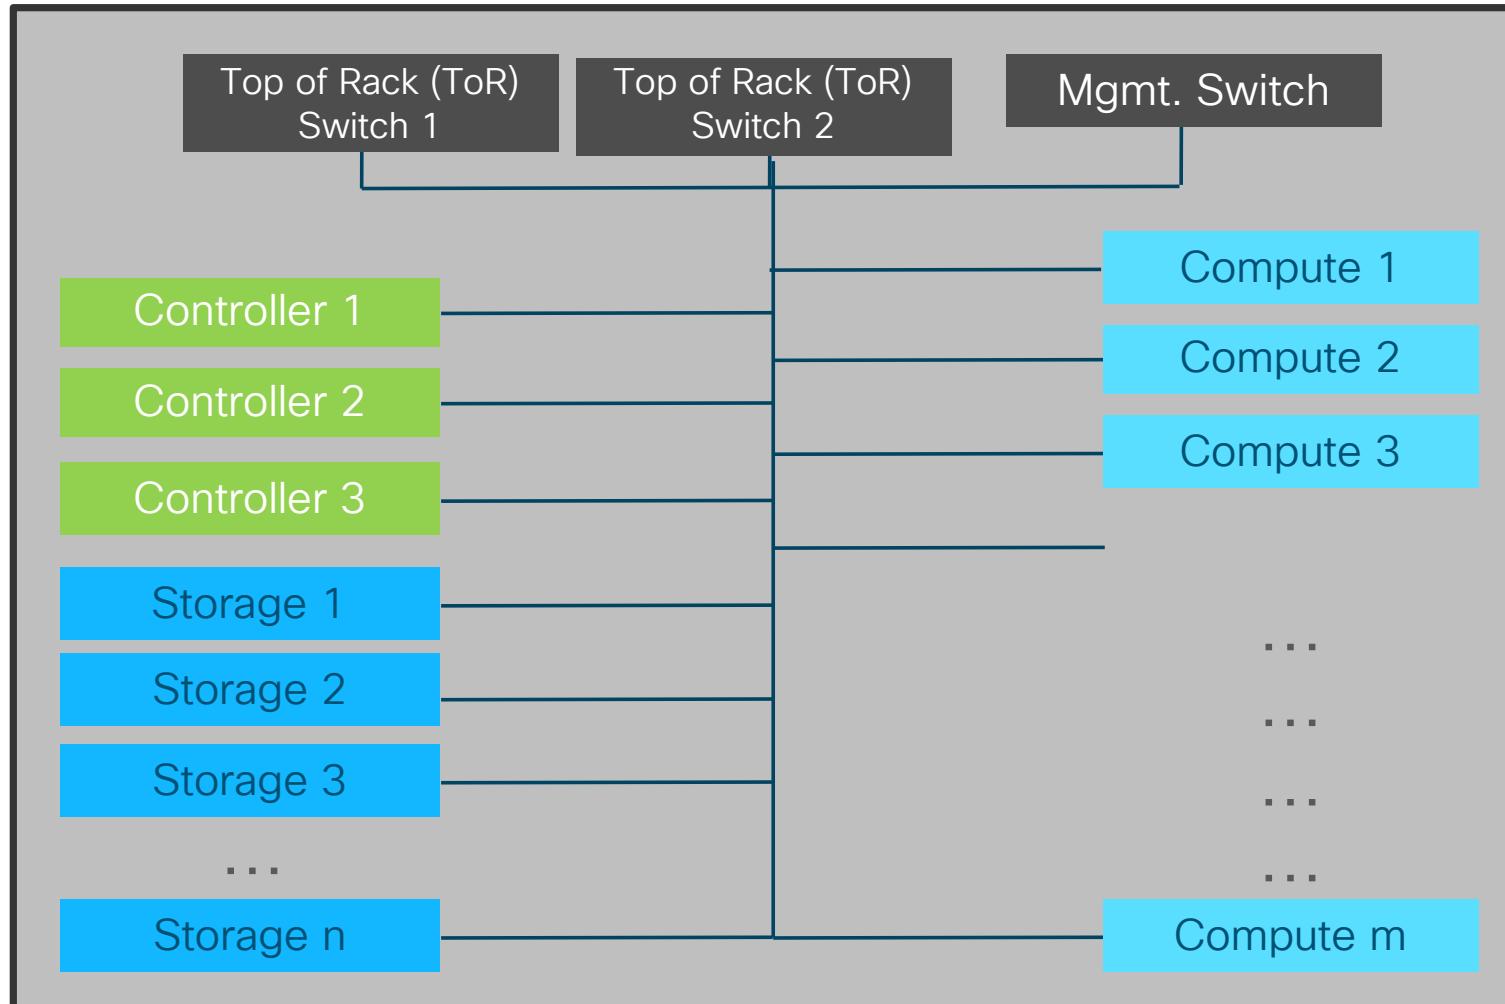
# OpenStack Logical Architecture



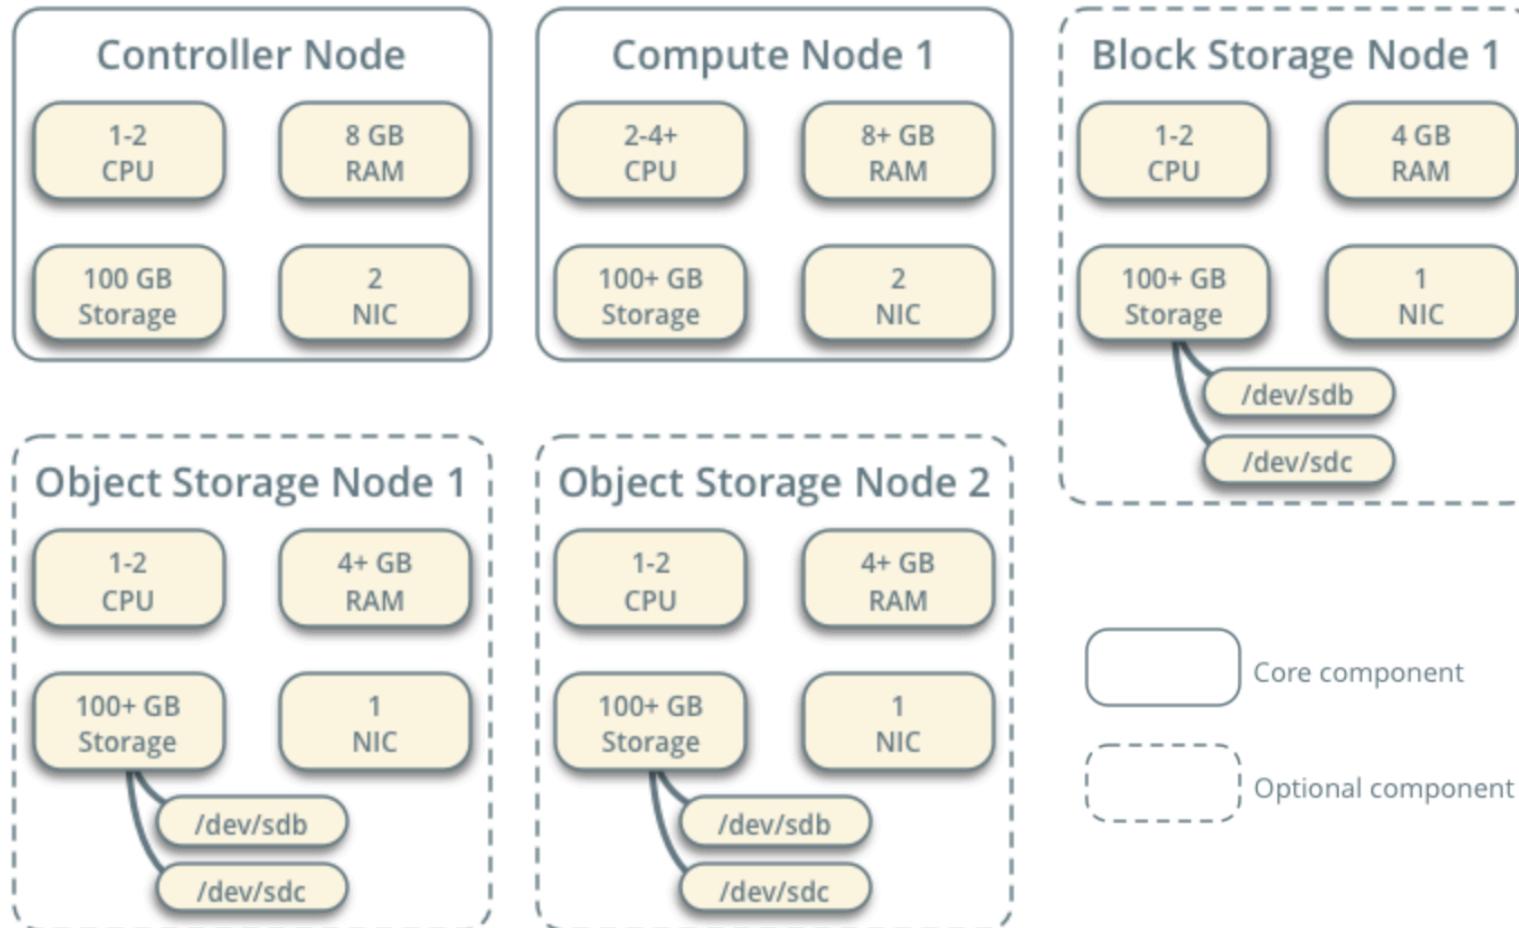
# OpenStack Conceptual Architecture



# Typical OpenStack Deployment for Production



# Minimal OpenStack Deployment for Proof-of-concept



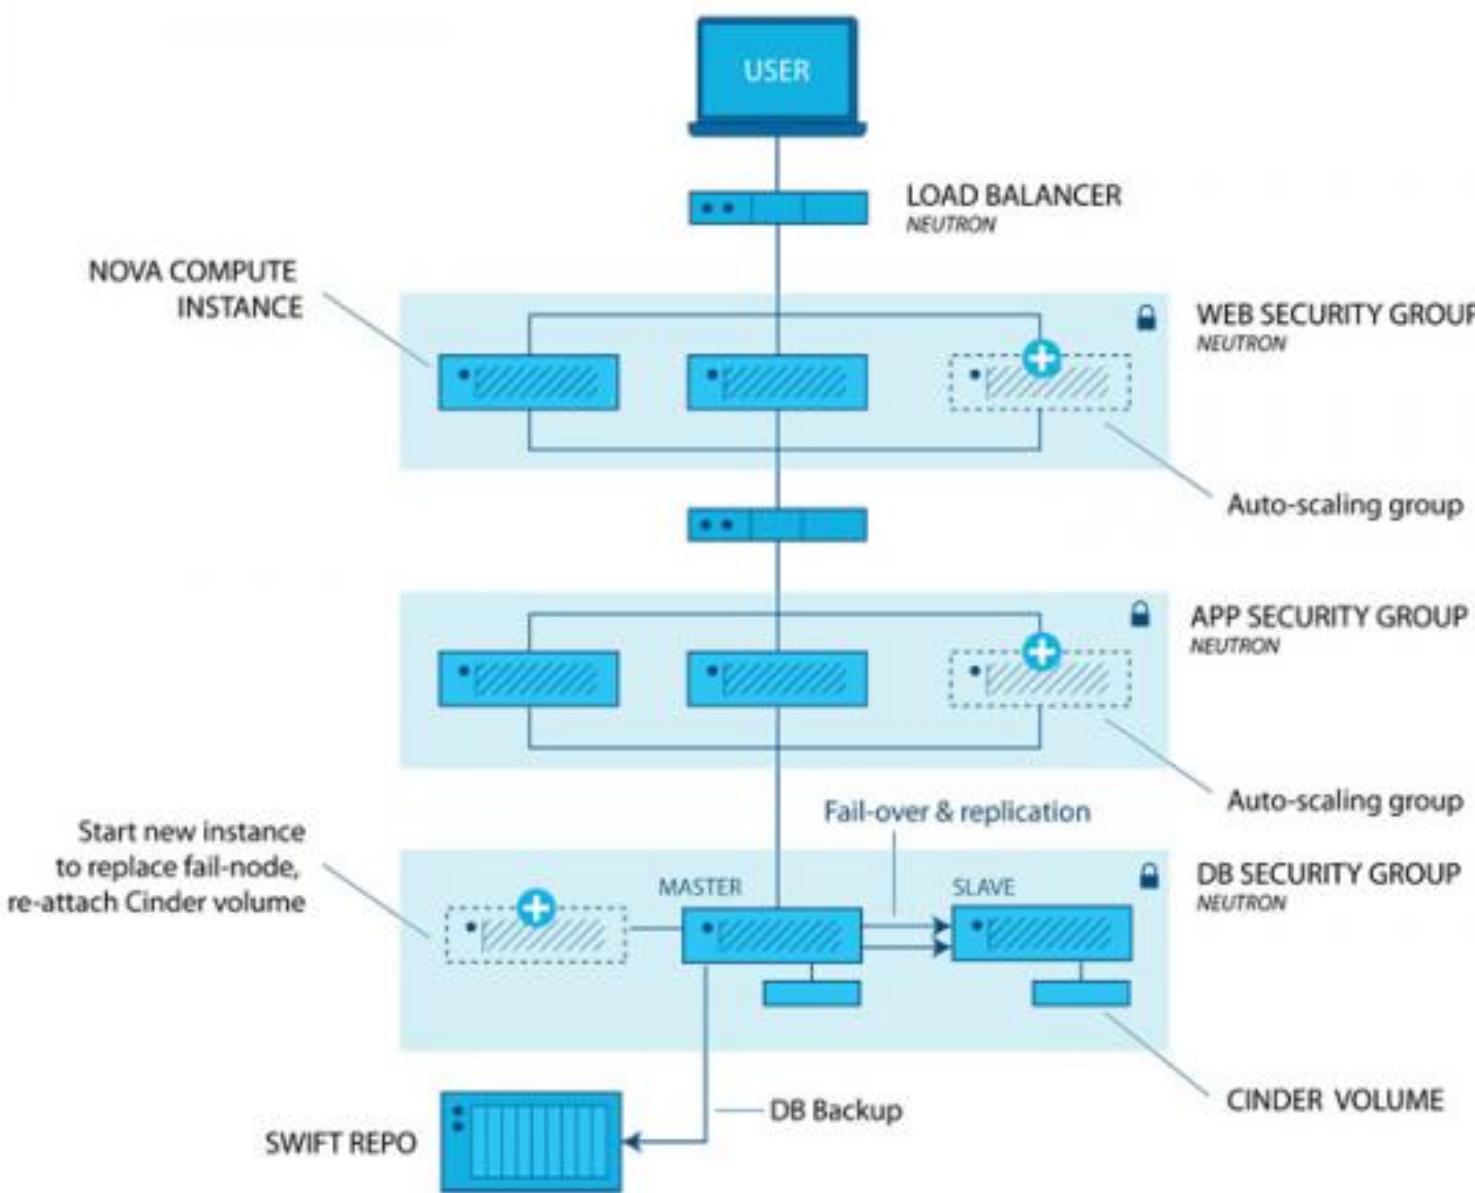
## Controller

Identity service, image service, management portions of compute and networking, networking agents and the dashboard.  
Also includes supporting services such as SQL database, message queue and NTP.

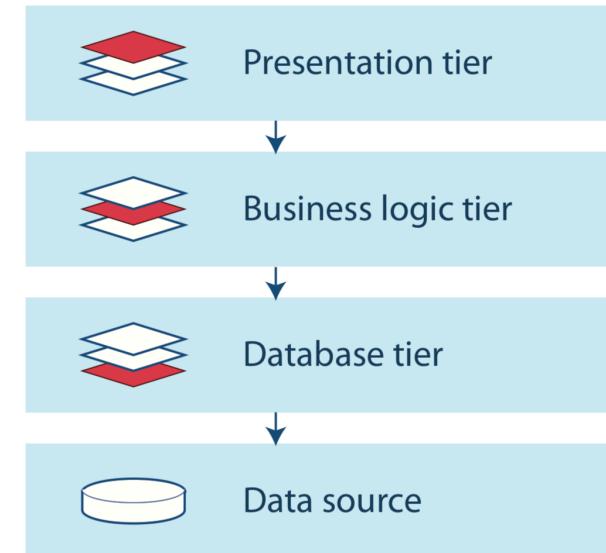
## Compute

Runs the hypervisor portion of compute that operates instances (default KVM).  
Also runs a networking service agent that connects instances to virtual networks and provides firewalling services to instances via security groups.  
Each node requires a minimum of two network interfaces.

# Use case 1: Web Applications



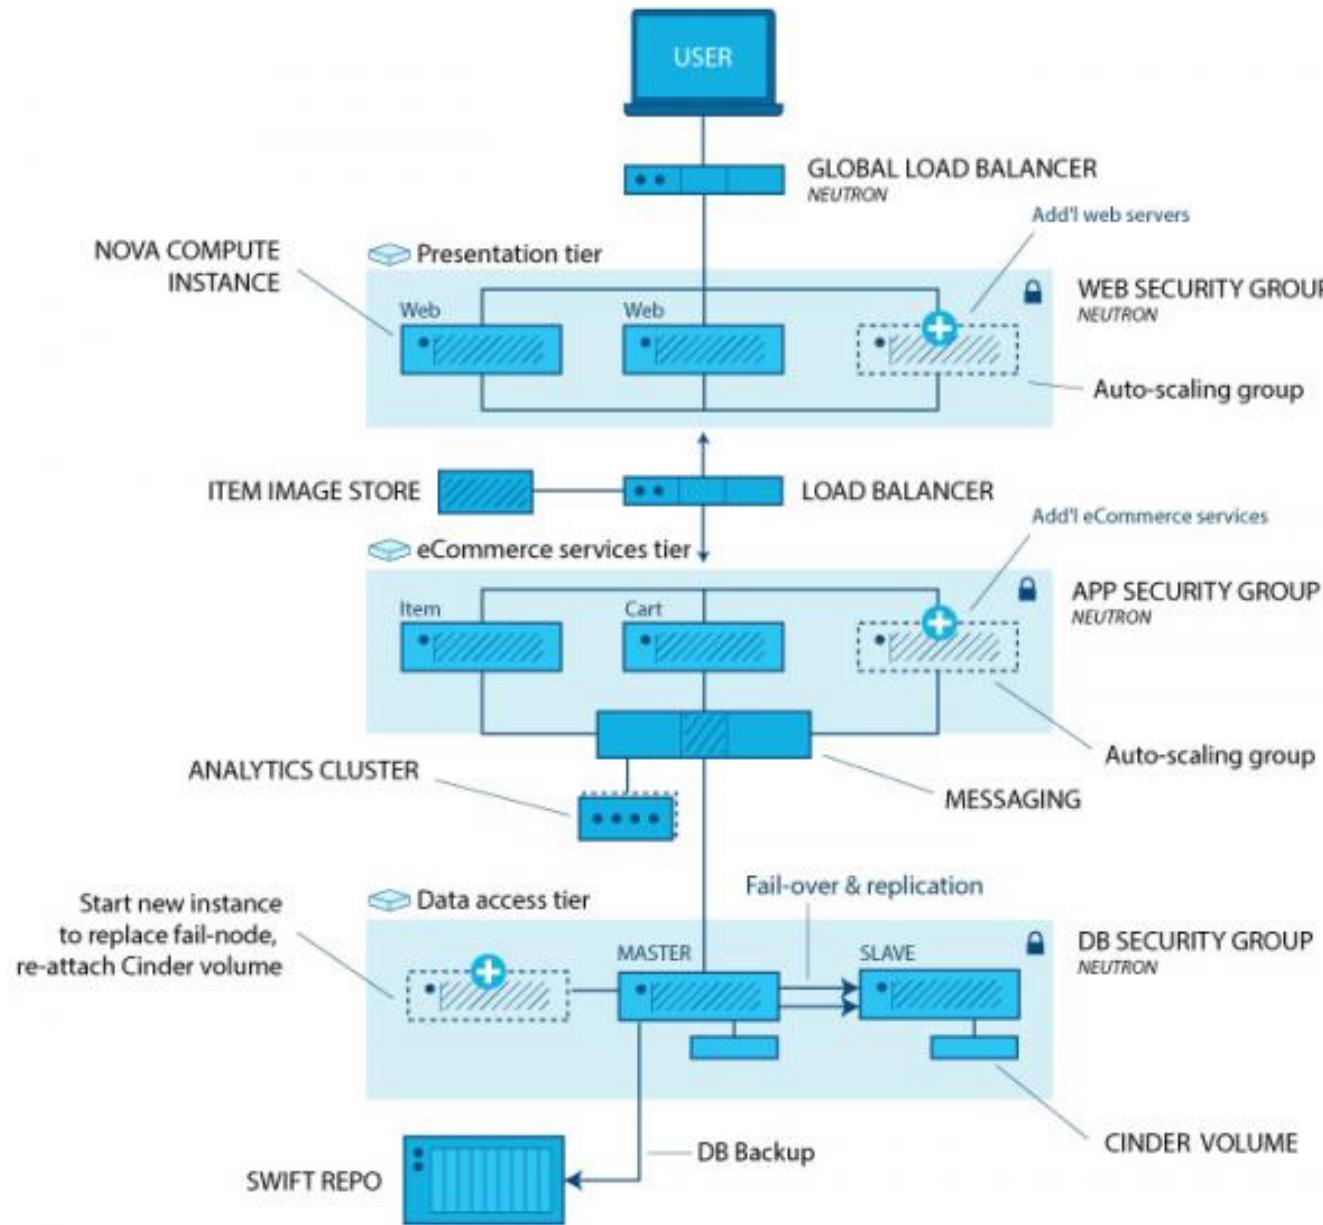
**Three-tier web application architecture overview**



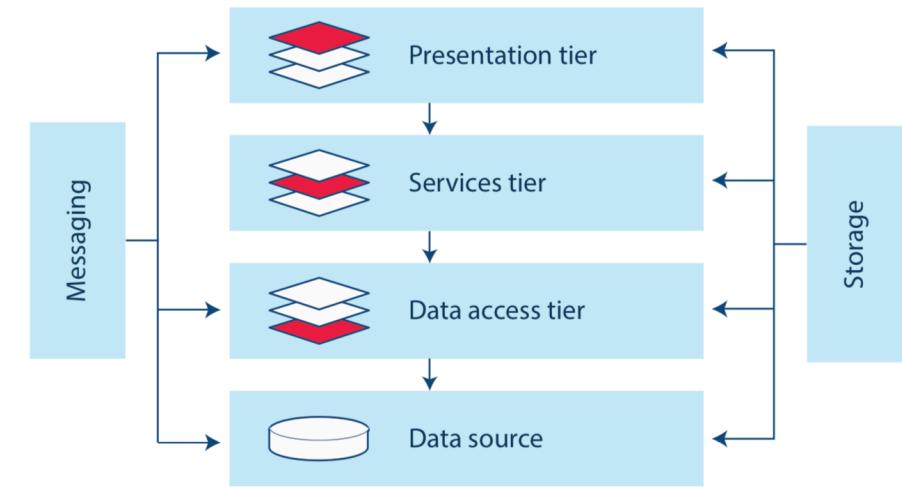
## Enterprises using it:

Workday, Betfair, Ancestry.com, JFE Steel, HMRC (Her Majesty's Revenue and Customs) and LivePerson

# Use case 2: eCommerce



**eCommerce high-level architecture**



## Enterprises using it:

Best Buy, Walmart, overstock.com, eBay, Staples, Nike, and Snapdeal

# Openstack Services for both use cases

Authentication & authorization



KEYSTONE

Interfaces for managing OpenStack



HORIZON



OS CLIENT CLI

Orchestration



HEAT

Networking resources



NEUTRON

Persistent storage resources



CINDER



SWIFT

Instance management



GLANCE

Compute resources



NOVA

Database



TROVE

Telemetry & data collection



CEILOMETER

Legend

Core service

Optional service

# How to get started with OpenStack

## Option 1: Sign up for a public cloud trial

Quickly gain access to OpenStack infrastructure via trial programs from participating OpenStack public cloud providers around the world. Experience the freedom, performance and interoperability of open source infrastructure with the [Public Cloud Global Passport Program](#).

[Try A Public Cloud](#)



## Option 2: Local Dev Environment: [devstack.org](#)

With some technical skills, DevStack is a great option to install and run an OpenStack cloud on your laptop (or even inside the VM on a cloud). DevStack is ideal for potential users who want to see what the Dashboard looks like from an admin or user perspective, and OpenStack contributors wanting to test against a complete local environment.

[Running DevStack: Quick Start](#)

[Wiki](#)

<https://www.openstack.org/software/start>

# Devstack Overview



DevStack is a series of **extensible scripts** used to quickly bring up a complete OpenStack environment based on the latest versions of everything from git master

- **Host OS:** Latest Ubuntu and RHEL releases
- **Databases:** MySQL as packaged by the host OS
- **Queues:** Rabbit as packaged by the host OS
- **Web Server:** Apache as packaged by the host OS
- **Networking:** A basic configuration approximating the original FlatDHCP mode using linuxbridge or OpenVSwitch
- **Services:** The default services configured by DevStack are Identity (**keystone**), Object Storage (**swift**), Image Service (**glance**), Block Storage (**cinder**), Compute (**nova**), Placement (**placement**), Networking (**neutron**), Dashboard (**horizon**)

# OpenStack APIs

- Use the OpenStack APIs to launch server instances, create images, create storage containers and objects, and complete other actions in the OpenStack cloud.
- After you authenticate through Identity, you can use the other OpenStack APIs to create and manage resources in your OpenStack cloud. You can launch instances from images and assign metadata to instances through the Compute API or the OpenStack command-line client.

# Methods to send API requests

- **cURL**

A command-line tool that lets you send HTTP requests and receive responses.

- **OpenStack command-line client**

The OpenStack project provides a command-line client that enables you to access APIs through easy-to-use commands

- **REST clients**

Browser-based graphical interfaces for REST in Firefox and Chrome, AdvancedRestClient, Postman, etc

- **OpenStack Python Software Development Kit (SDK)**

Use this SDK to write Python automation scripts that create and manage resources in your OpenStack cloud. The SDK implements Python bindings to the OpenStack API, which enables you to perform automation tasks in Python by making calls on Python objects rather than making REST calls directly

# Using the OpenStack APIs

1. Issue an authentication request with a payload of credentials to OpenStack Identity to get an authentication token

Credentials are usually a combination of your user name and password, and optionally, the name or ID of the project of your cloud.

2. When you send API requests, you include the token in the X-Auth-Token header. If you access multiple OpenStack services, you must get a token for each service.

A token is valid for a limited time before it expires. A token can also become invalid for other reasons. For example, if the roles for a user change, existing tokens for that user are no longer valid.

<https://developer.openstack.org/api-guide/quick-start/api-quick-start.html>

# Openstack Toolkits 1/2

- **Go**
  - [Gophercloud](#) provides a Go binding to OpenStack cloud APIs.
- **Java**
  - [OpenStack4j](#) A fluent Java OpenStack API.
  - [OpenStack Java SDK](#) is a Java binding for the OpenStack APIs.
  - [Apache jclouds](#) is a Multi-cloud SDK for Java with OpenStack support
- **JavaScript**
  - [pkgcloud](#) is a Multi-cloud library for Node.js that supports OpenStack
  - [jstack](#) is a JavaScript client library for the OpenStack API.
  - [js-openclient](#) is a very opinionated core client which can be used in either Node.js or in the browser (browser support not yet complete) to communicate with a RESTful APIs, including but not limited to any OpenStack-compatible API. (last updated 2015)
  - [node-openstack-wrapper](#) is a convenience wrapper for many of Openstack's common features with a focus on projects/tenants.
- **.NET**
  - [OpenStack.NET](#) is a .NET SDK for OpenStack.

# Openstack Toolkits 2/2

- **PHP**

- [php-opencloud/openstack](#) is a PHP SDK for OpenStack. It is actively maintained and supports latest OpenStack identity api (Keystone v3).
- SDK covers following api: BlockStorage v2, Compute v2, Identity Keystone v2 and v3, Images v2, Networking v2 including Layer3 extension and Security Groups extension, Object Storage v1, Gnocchi v1

- **Python**

- [OpenStack Shade](#) shade is a simple client library for operating OpenStack clouds that abstracts deployer differences
- The [SDK-Development/PythonOpenStackSDK](#) project is a proposed solution to offering an SDK that provides a single point of entry for consumers, and a base from which other tools can be built upon, such as command-line interfaces.
- The [OpenStackClients](#) are the native Python bindings for the OpenStack APIs. They are used to implement the command-line interfaces (which ship with the library).
- [Apache libcloud](#) is a Python library that abstracts away differences among multiple cloud provider APIs.

- **Ruby**

- [fog](#) is a Multi-cloud Ruby library with support for OpenStack
- [Misty](#) is a dedicated HTTP client for OpenStack APIs
- [Aviator](#) An elegantly designed OpenStack SDK for Ruby (hasn't been updated since 2015)