

# Projetos e APIs




Aula Teórica nº3

2020/2021



# OpenStack

- Open source cloud computing platform for all types of clouds
- Aims to be simple to implement, massively scalable, and feature rich
- Developers and cloud computing technologists from around the world create the OpenStack project
- OpenStack provides an Infrastructure-as-a-Service (IaaS) solution through a set of interrelated services
- Each service offers an application programming interface (API) that facilitates this integration
- It's possible to install some or all services




Compute

	NOVA	Compute Service
	ZUN	Containers Service
	QINLING	Functions Service




Bare Metal

	IRONIC	Bare Metal Provisioning Service
	CYBORG	Accelerators resource management





Storage

	SWIFT	Object store
	CINDER	Block Storage
	MANILA	Shared filesystems




Networking

	NEUTRON	Networking
	OCTAVIA	Load balancer
	DESIGNATE	DNS service





Orchestration

	HEAT	Orchestration
	SENLIN	Clustering service
	MISTRAL	Workflow service
	ZAQAR	Messaging Service
	BLAZAR	Resource reservation service
	AODH	Alarming Service

Workload Provisioning




	MAGNUM	Container Orchestration Engine Provisioning
	SAHARA	Big Data Processing Framework Provisioning
	TROVE	Database as a Service

Application Lifecycle





	MASAKARI	Instances High Availability Service
	MURANO	Application Catalog
	SOLUM	Software Development Lifecycle Automation
	FREEZER	Backup, Restore, and Disaster Recovery

OpenStack  
Services

## Monitoring Tools

	<b>CEILOMETER</b>	Metering & Data Collection Service
	<b>PANKO</b>	Event, Metadata Indexing Service
	<b>MONASCA</b>	Monitoring

## Optimization/policy tools

	<b>WATCHER</b>	Optimization Service
	<b>VITRAGE</b>	Root Cause Analysis service
	<b>CONGRESS</b>	Governance
	<b>RALLY</b>	Benchmark service

## Billing / Business Logic

	<b>CLOUDKITTY</b>	Billing and chargebacks
---	-------------------	-------------------------

## Multi-Region Tools

	<b>TRICIRCLE</b>	Networking Automation for Multi-Region Deployments
---	------------------	--

## Containers

	<b>KURYR</b>	OpenStack Networking integration for containers
---	--------------	---


## NFV

	<b>TACKER</b>	NFV Orchestration
---	---------------	-------------------

## API Proxies

	<b>EC2API</b>	EC2 API proxy
---	---------------	---------------

## Web Frontend

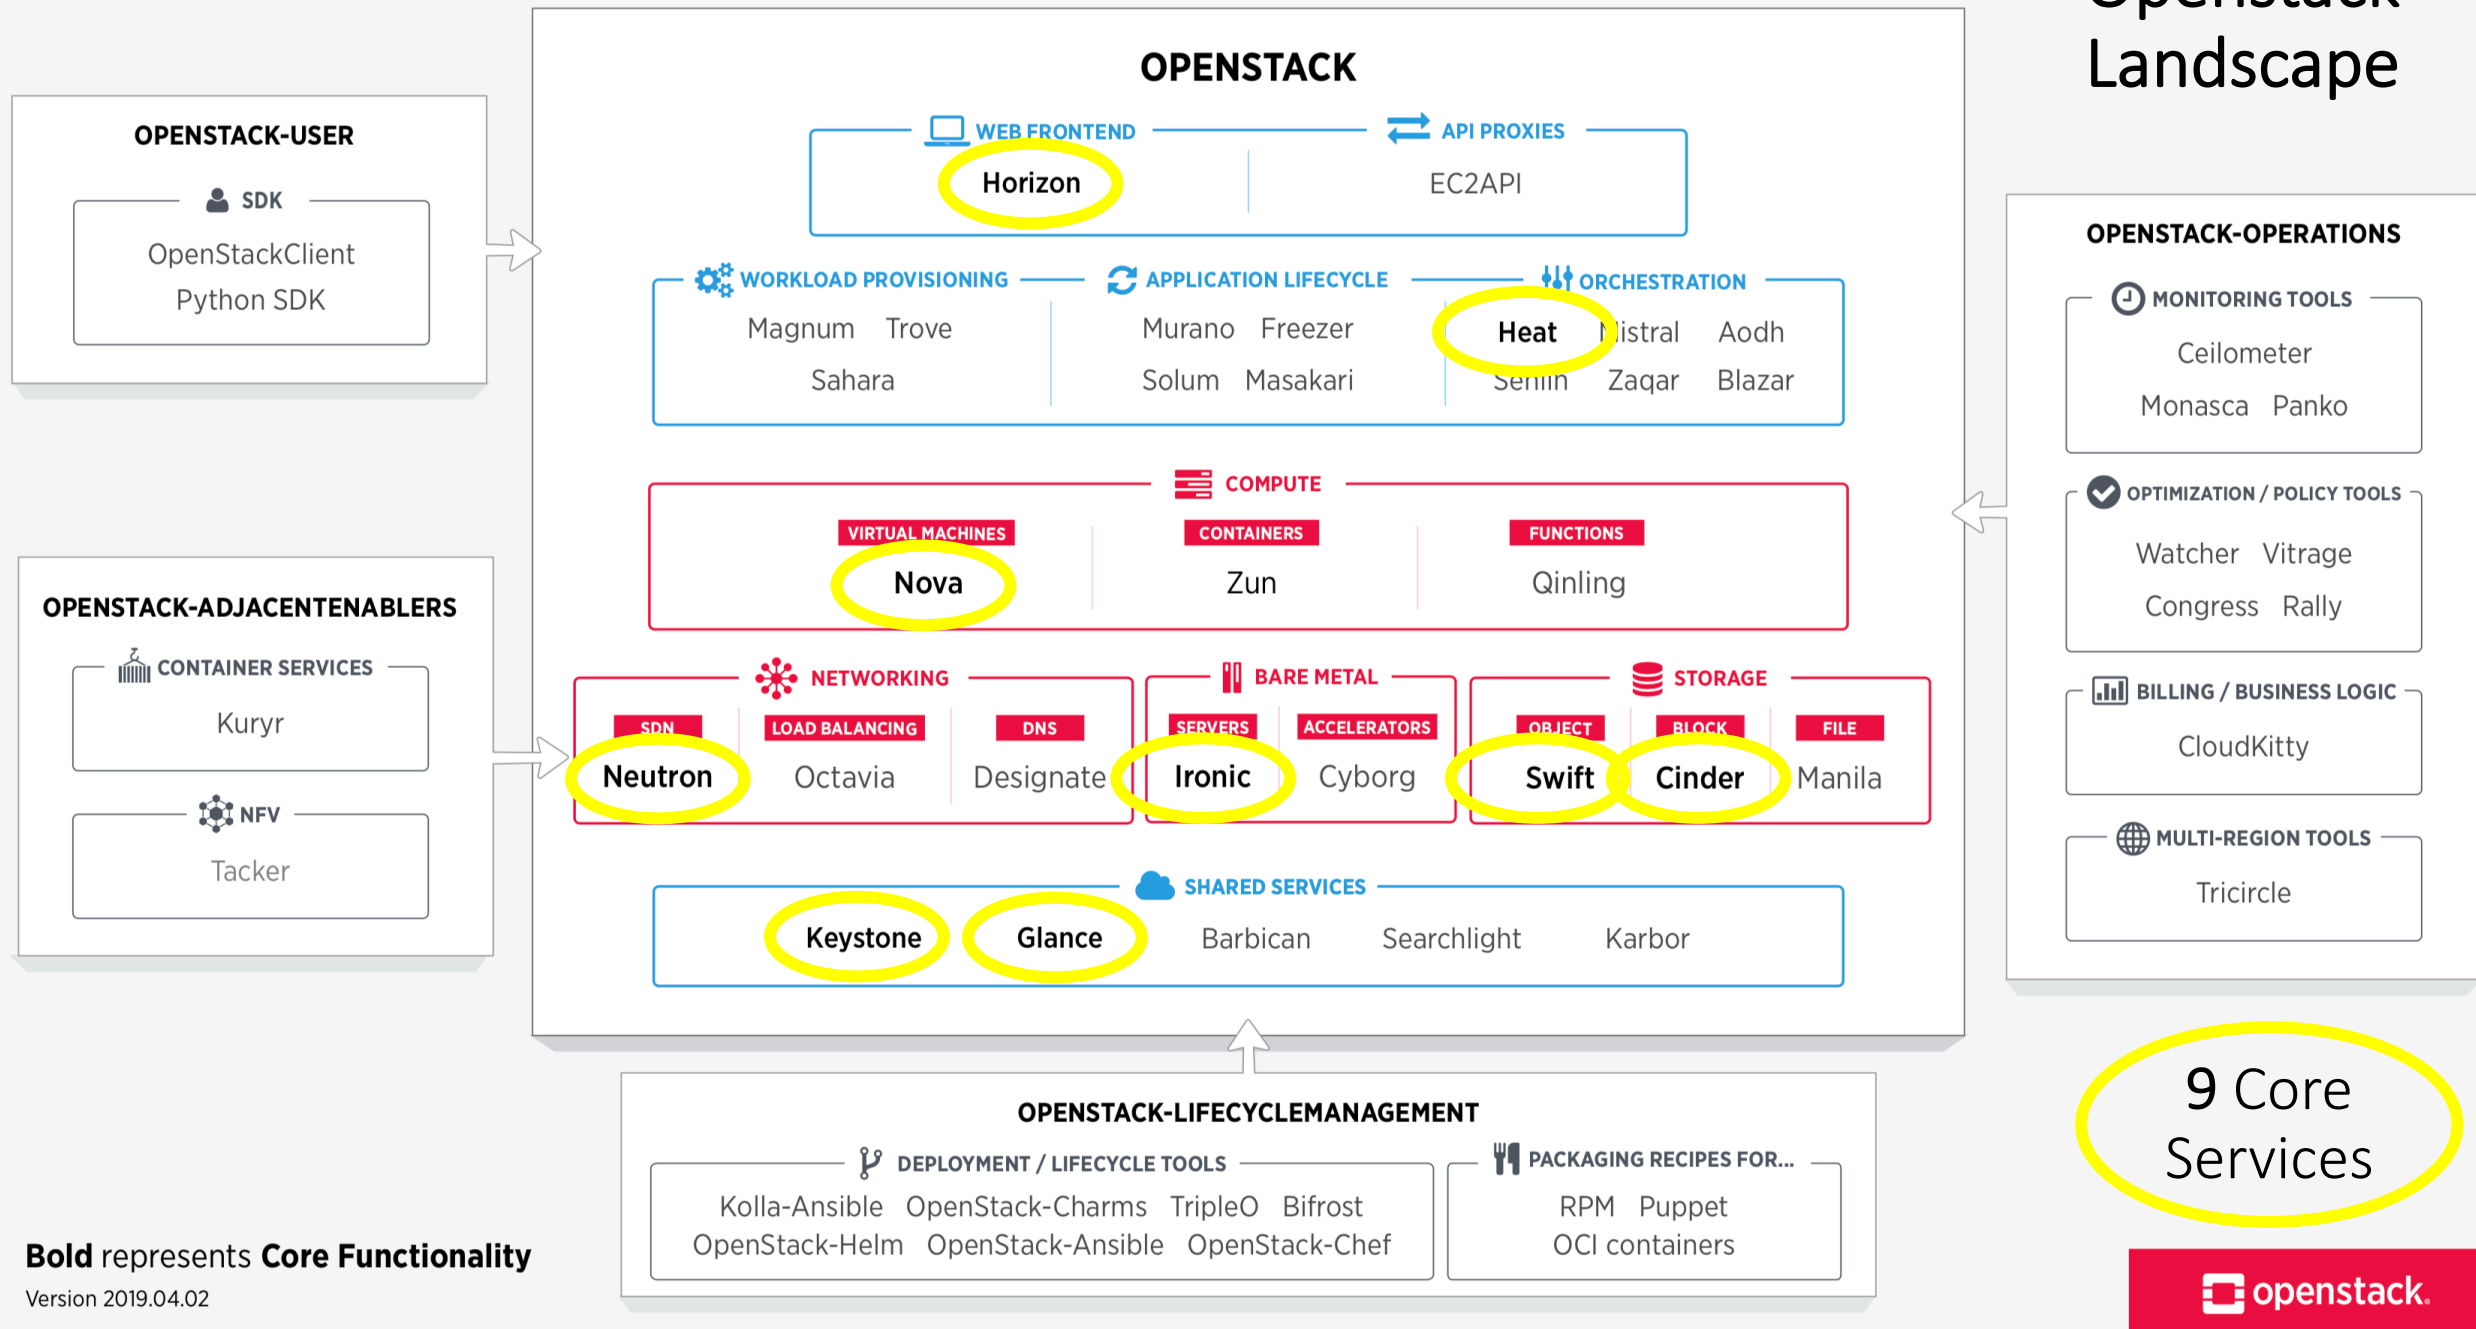
	<b>HORIZON</b>	Dashboard
---	----------------	-----------

## Swift add-ons

	<b>STORLETS</b>	Computable object storage
---	-----------------	---------------------------

OpenStack  
Services

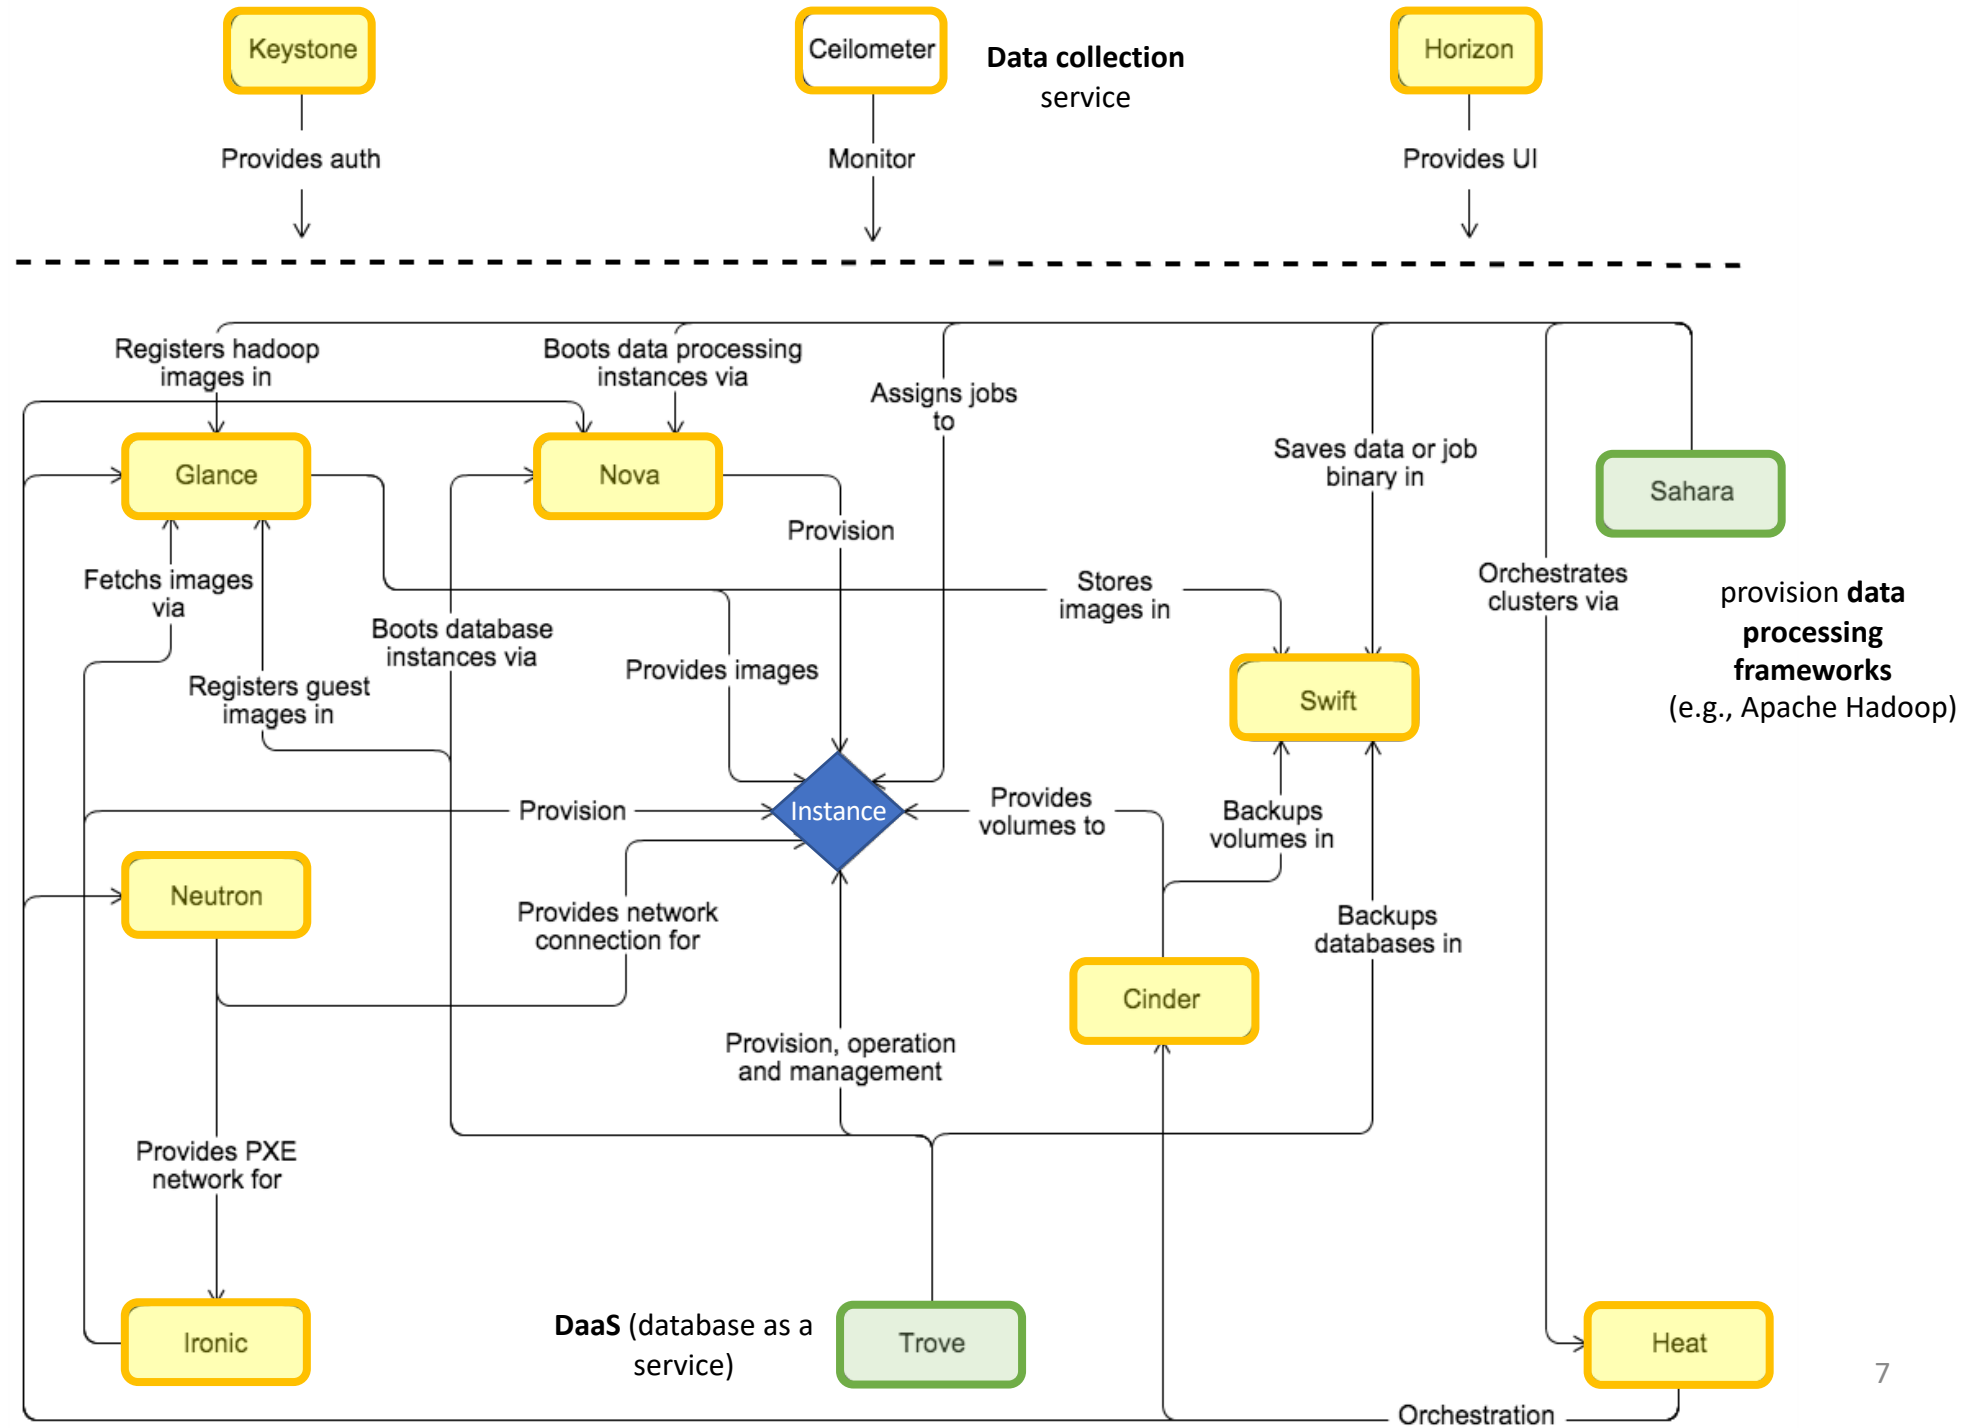
# Openstack Landscape



# OpenStack Services

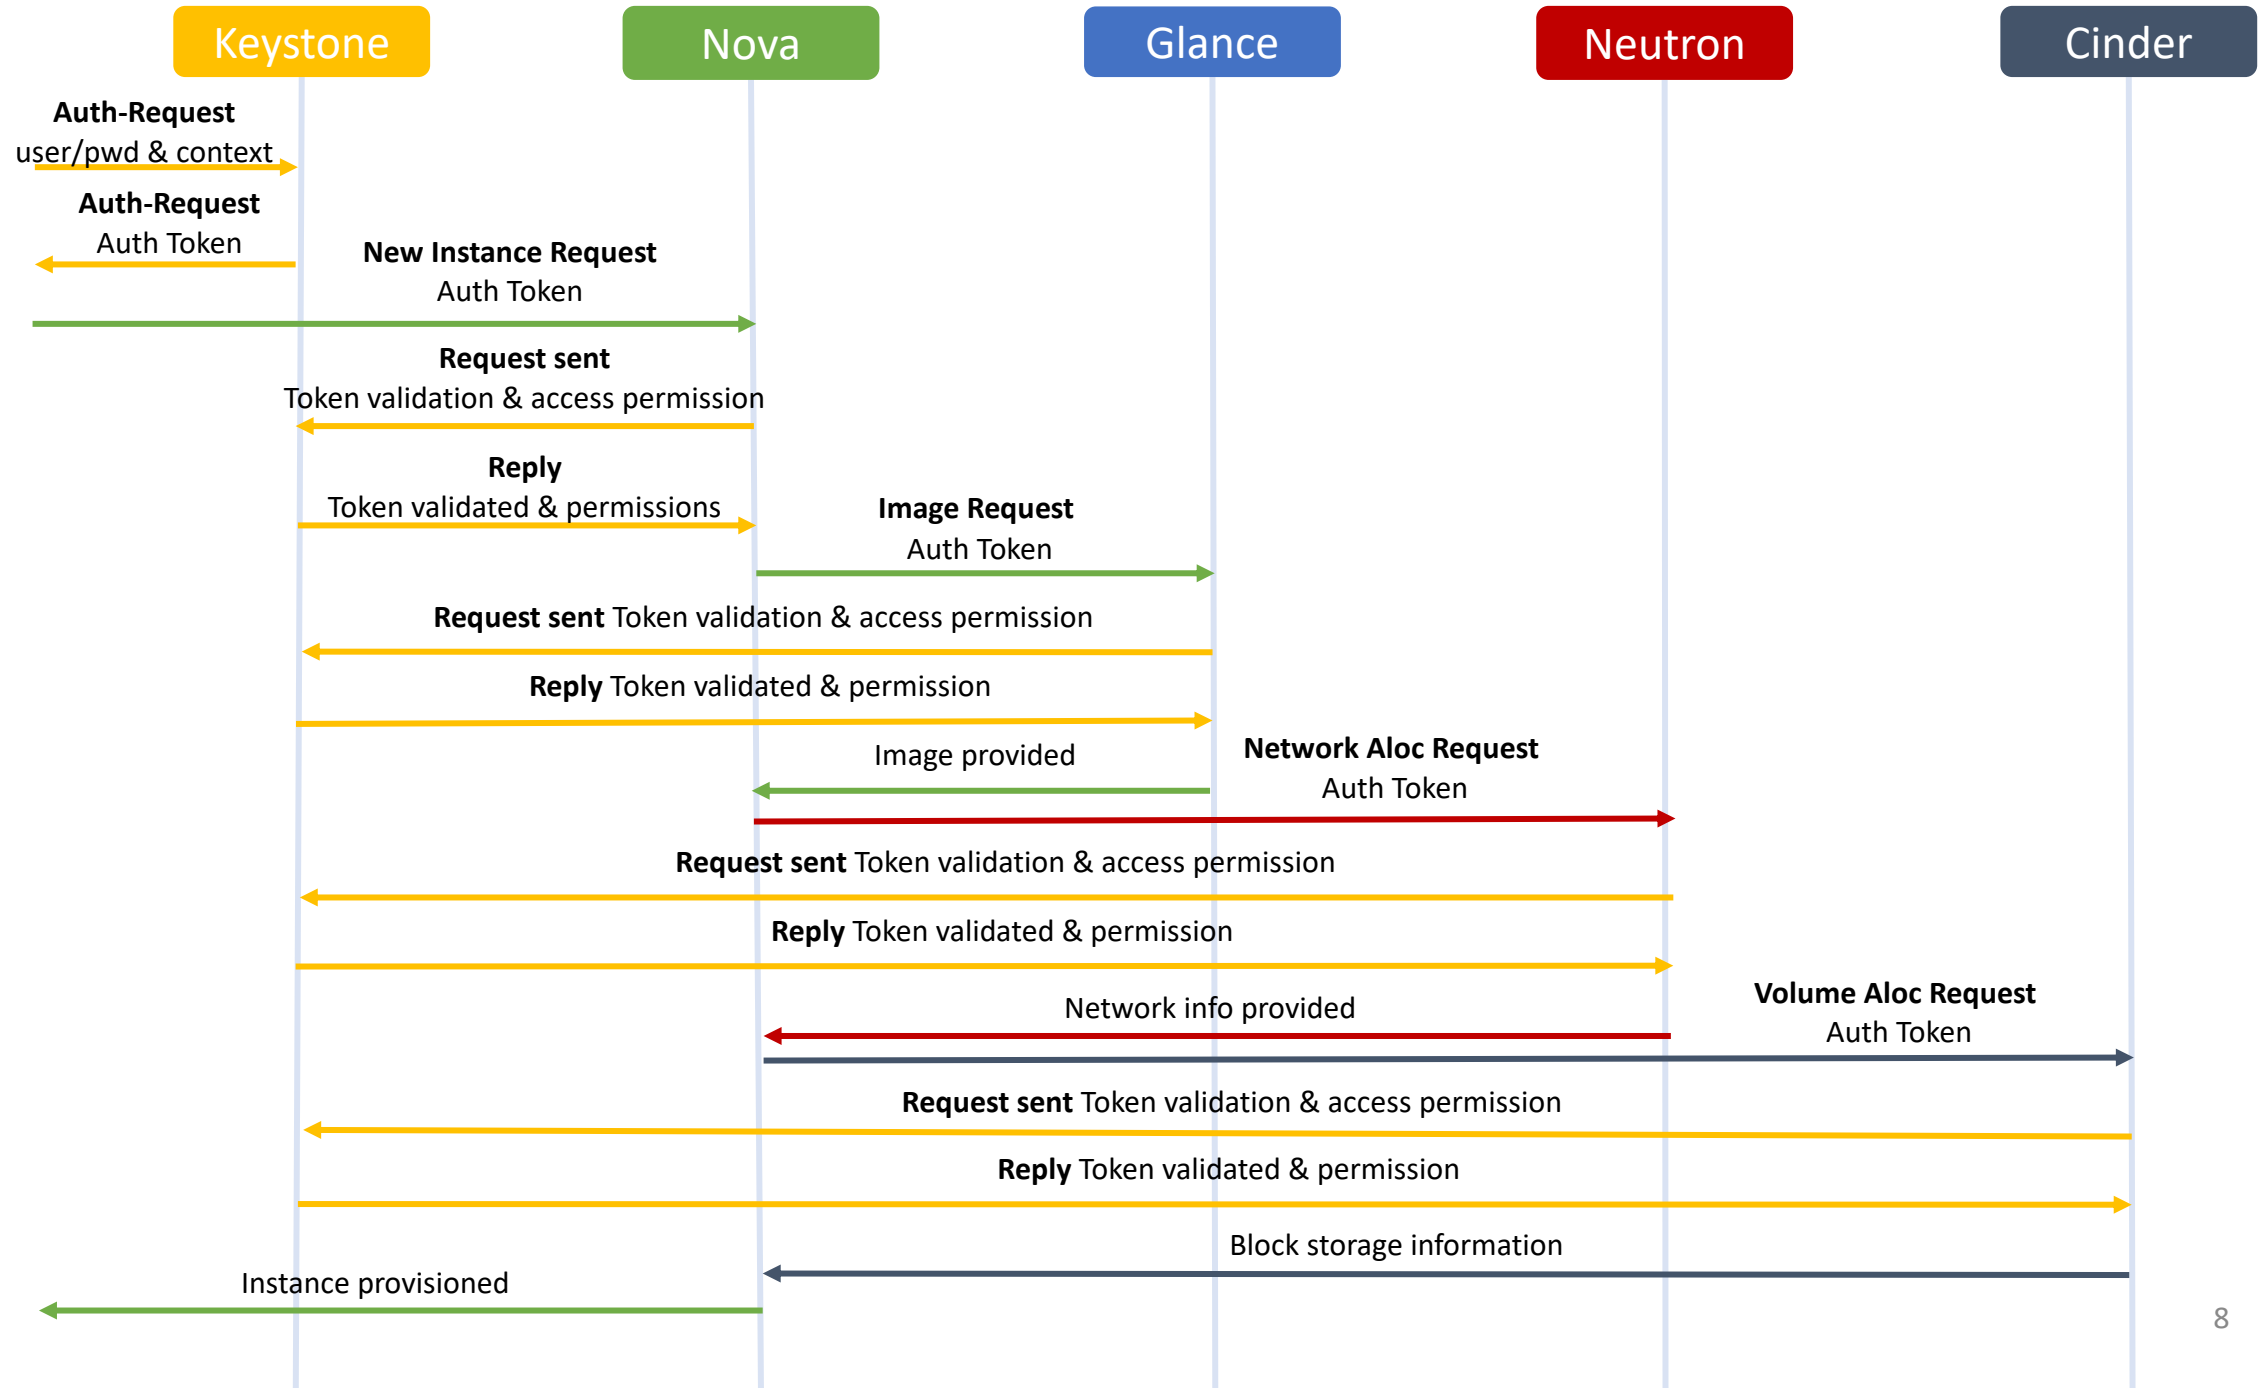
- OpenStack consists of several independent parts, named the **OpenStack services**
- Those services work together and include the **Compute, Identity, Networking, Image, Block Storage, Object Storage, Telemetry, Orchestration, and Database** services
- It's possible to install any of these projects separately and configure them stand-alone or as connected entities
- Individual **services interact with each other through public APIs**, except where privileged administrator commands are necessary
- OpenStack services are composed of **several processes**
- For communication between the processes of one service, an **AMQP message broker** is used
- The service's state is stored in a **database**

# OpenStack Conceptual Architecture



# Instance creation Step-by-Step

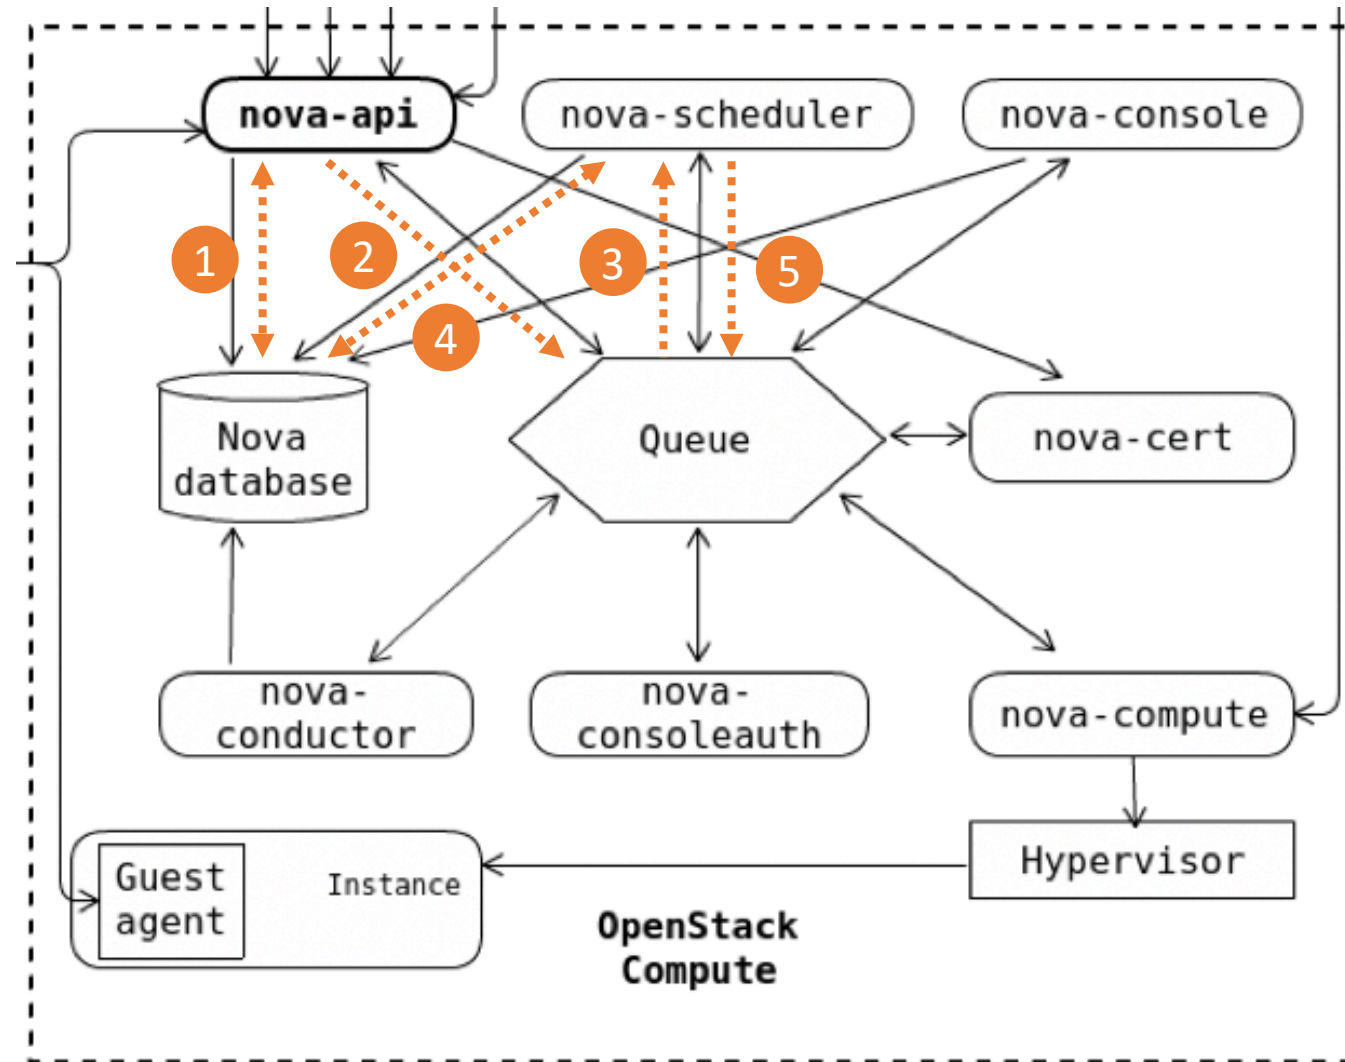
Horizon Dashboard or OpenStack CLI





# Inside Nova: the processes (1/2)

1. After getting the response from keystone, then **nova-api** checks for conflicts with **nova-database** and then it creates initial database entry for **new instance** or VM.
2. **nova-api** sends the **rpc.call** request to **nova-scheduler** expecting to get updated instance entry with **host ID** specified
3. **nova-scheduler** picks the request from the **queue**
4. **nova-scheduler** talks to **nova-database** to locate an appropriate **host ID** using **filtering and weighing mechanism**
5. **nova-scheduler** sends the **rpc.cast** request to **nova-compute** for launching an **instance** on the **appropriate host**



# Inside Nova: the processes (2/2)

6. **nova-compute** picks the request from the queue and it sends the **rpc.call** request to **nova-conductor** to get the **VM or instance info** such as host ID and flavor (RAM,CPU and Disk)

7. **nova-conductor** takes the request from queue and communicate with **nova-database**, to get the **instance information** and sends the information to the queue

8. **nova-compute** picks the instance information from the queue

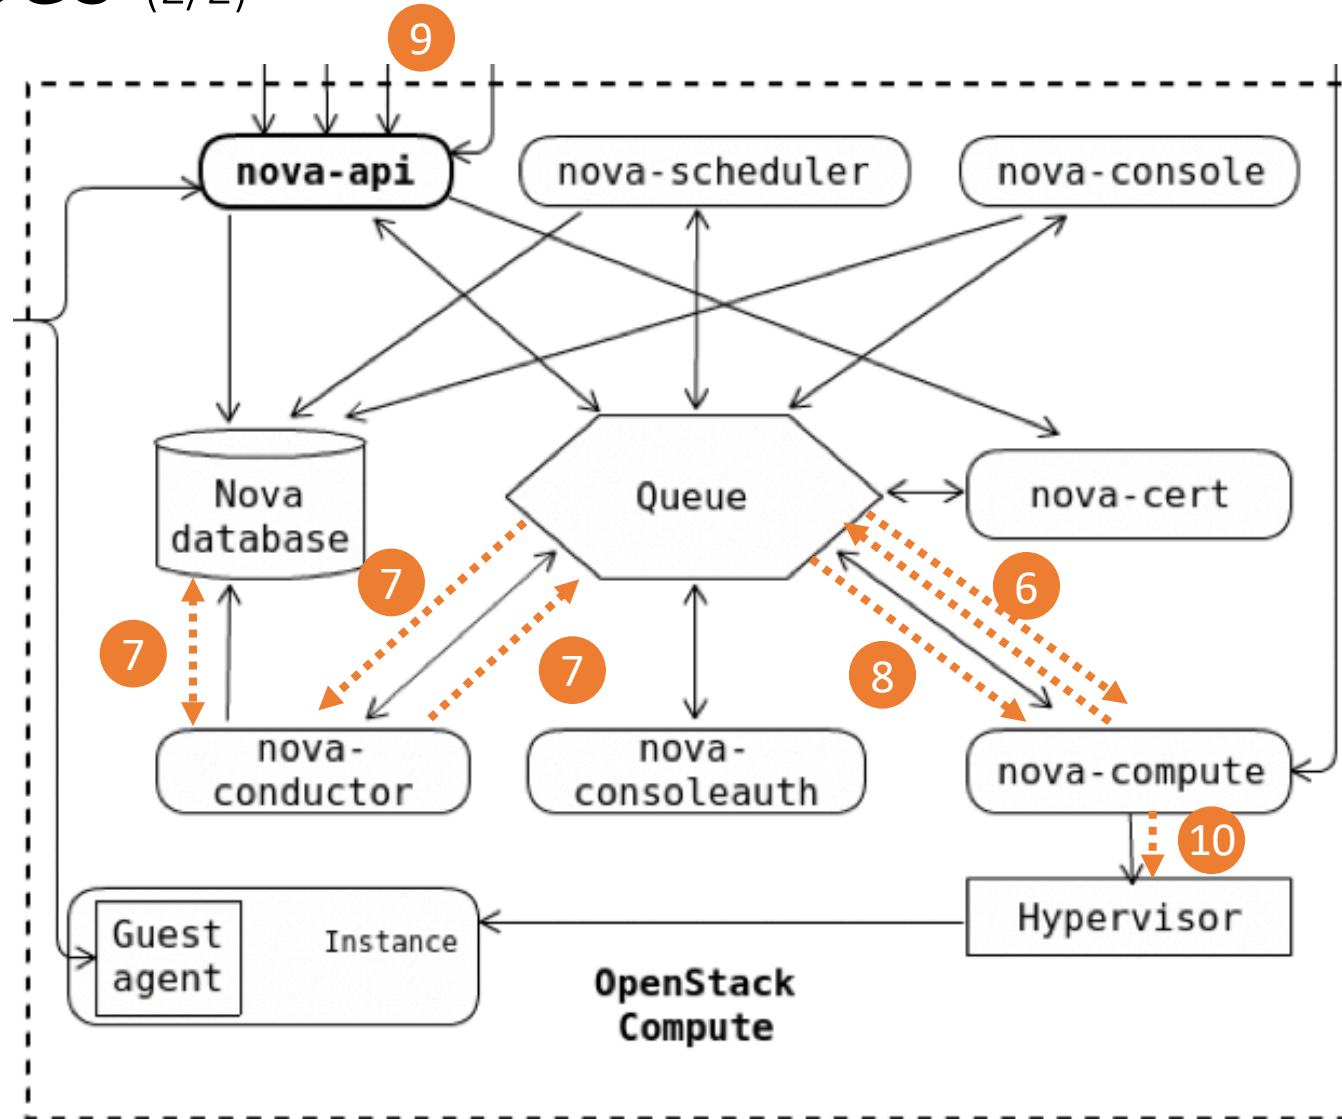
9. Through API **nova-compute** gets the image, networking and block storage information needed for the **instance**:

Glance

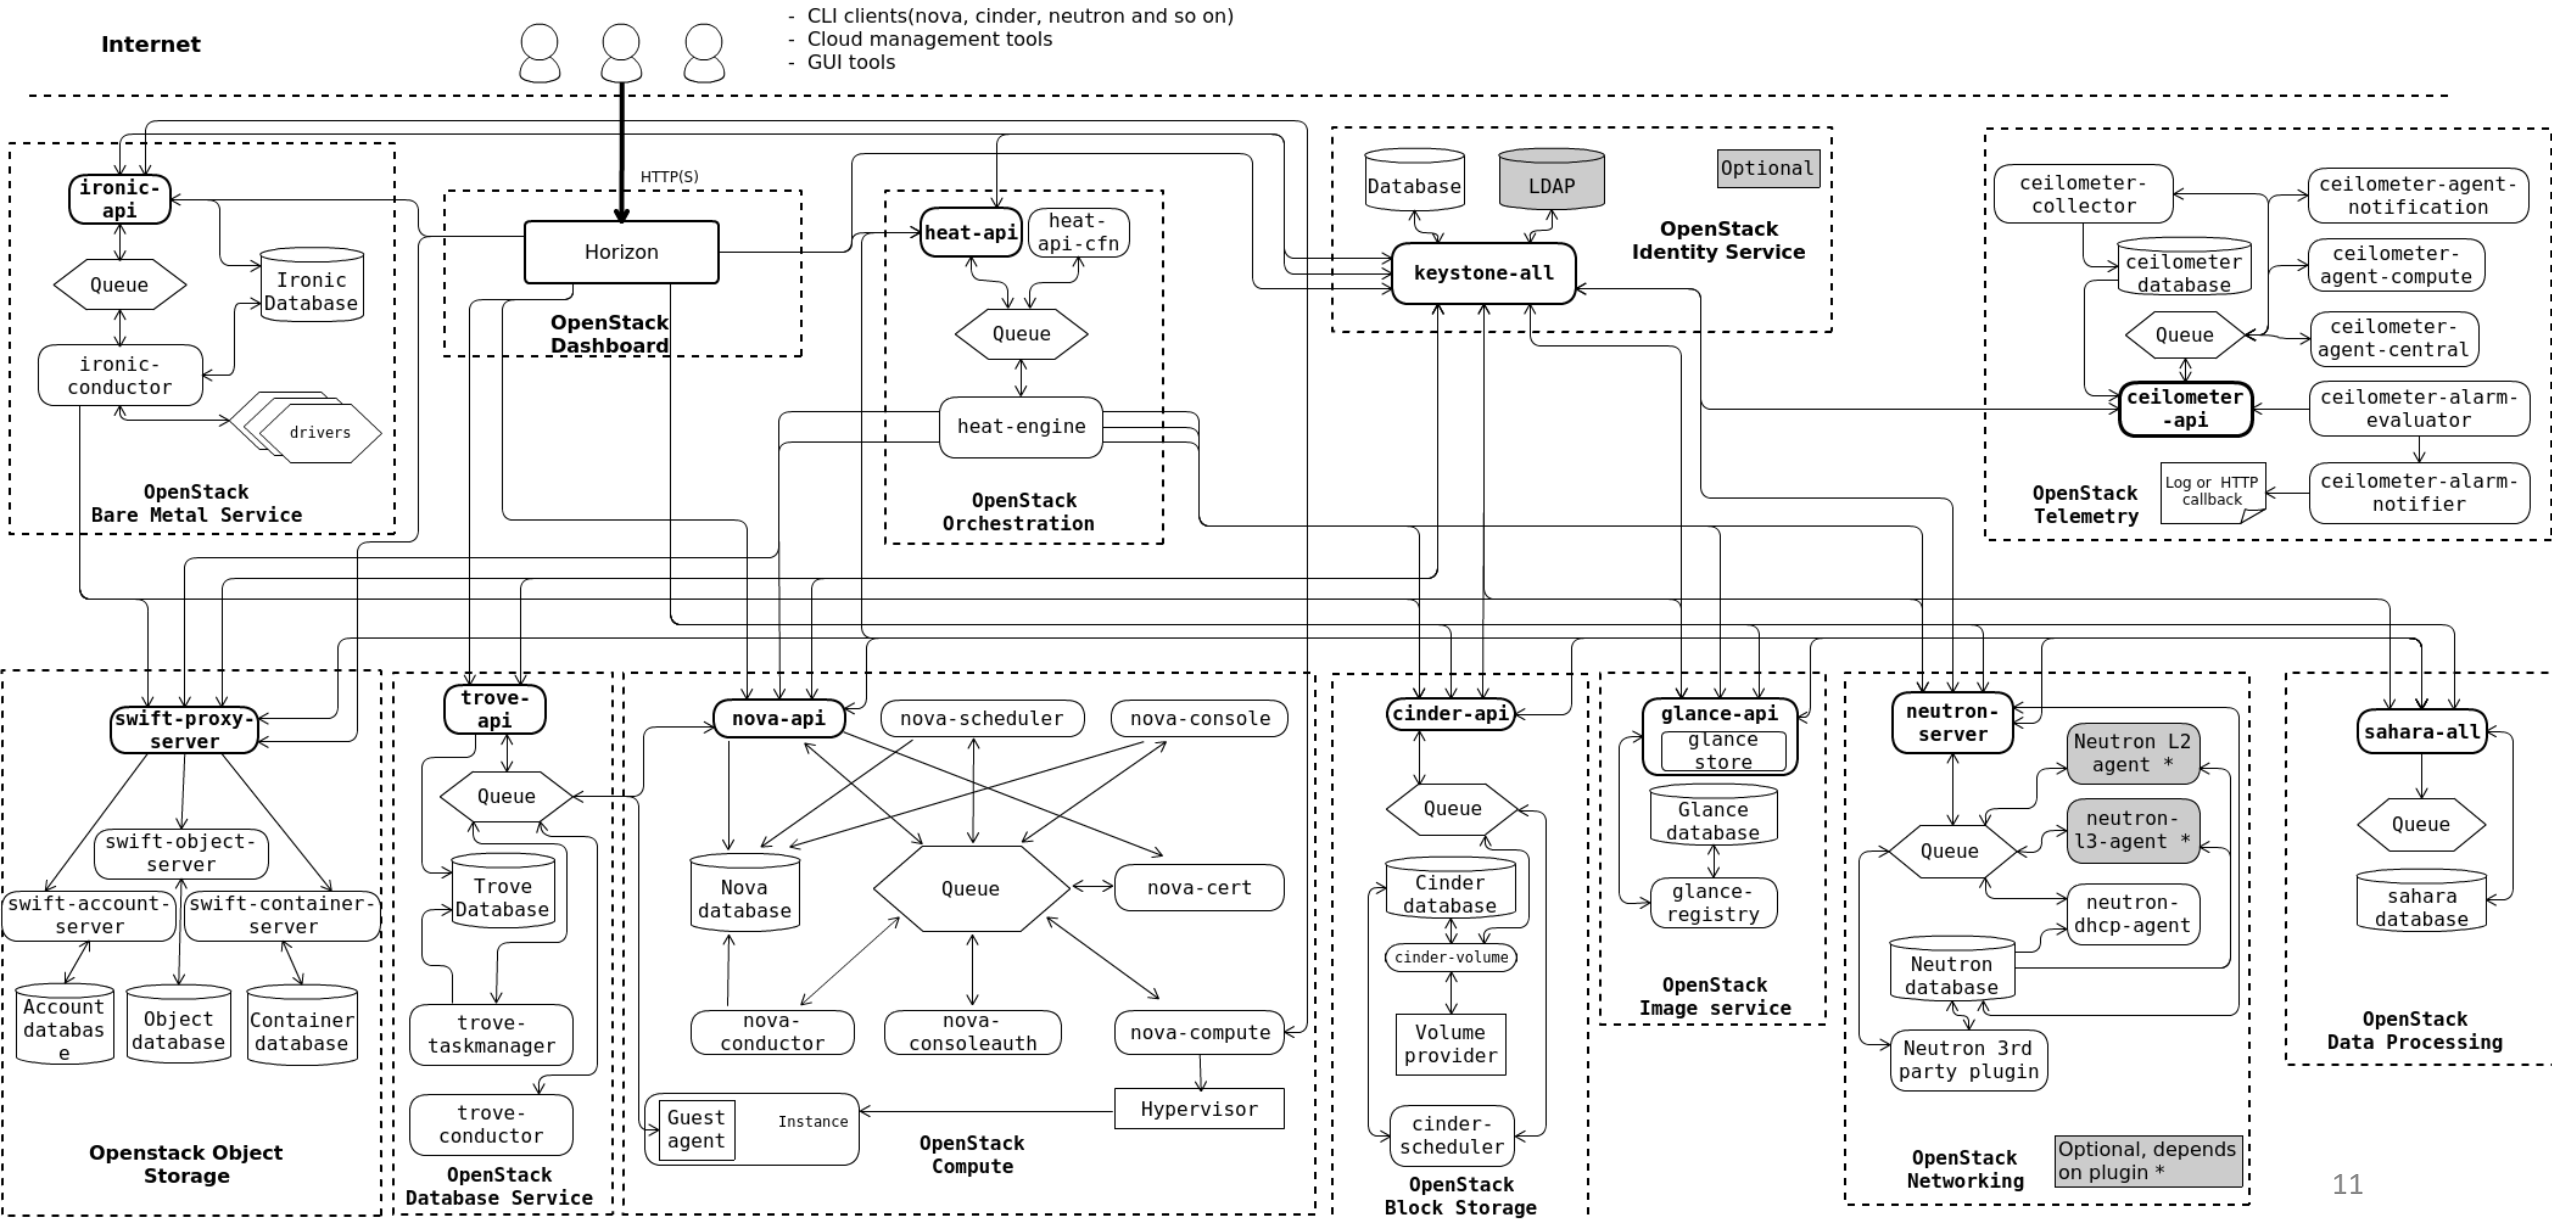
Neutron

Cinder

10. **nova-compute** generates data for the hypervisor driver and executes the request on the hypervisor using libvirt or API and then finally a VM is created on the hypervisor.



# OpenStack Logical Architecture



# Orchestration Service (Heat)

- Heat provides a **template based orchestration** for describing a cloud application
- A Heat template describes the infrastructure for a cloud application in text files which are readable and writable by humans (YAML)
- Templates specify the relationships between resources (e.g. this volume is connected to this server)
- The templates allow creation of most OpenStack resource types (such as instances, floating ips, volumes, security groups, users, etc)
- Heat primarily manages infrastructure, but the templates **integrate well with SW configuration management tools** such as Puppet and Ansible

**heat\_template\_version:** 2015-10-15

**description:** Launch a basic instance with CirrOS image using the  
``m1.tiny`` flavor, ``mykey`` key, and one network.

**parameters:**

**NetID:**

**type:** string

**description:** Network ID to use for the instance.

**resources:**

**server:**

**type:** OS::Nova::Server

**properties:**

**image:** cirros

**flavor:** m1.tiny

**key\_name:** mykey

**networks:**

– **network:** { **get\_param:** NetID }

**outputs:**

**instance\_name:**

**description:** Name of the instance.

**value:** { **get\_attr:** [ server, name ] }

**instance\_ip:**

**description:** IP address of the instance.

**value:** { **get\_attr:** [ server, first\_address ] }

# Heat – Launch an instance

1. Create the demo-  
template.yml file  
with the content in  
this slide

2. Determine available networks:

```
$ openstack network list
```

ID	Name	Subnets
4716ddfe-6e60-40e7-b2a8-42e57bf3c31c	selfservice	2112d5eb-f9d6-45fd-906e-7cabd38b7c7c
b5b6993c-ddf9-40e7-91d0-86806a42edb8	provider	310911f6-acf0-4a47-824e-3032916582ff

3. Set the NET\_ID environment variable to reflect the ID of a network. For example, using the provider network:

```
$ export NET_ID=$(openstack network list | awk '/ provider / { print $2 }')
```

4. Create a stack of one CirrOS instance on the provider network:

```
$ openstack stack create -t demo-template.yml --parameter "NetID=$NET_ID" stack
```

ID	Stack Name	Stack Status	Creation Time
Updated Time			
dbf46d1b-0b97-4d45-a0b3-9662a1eb6cf3	stack	CREATE_IN_PROGRESS	2015-10-13T15:27:20
None			

5. After a short time, verify successful creation of the stack:

```
$ openstack stack list
```

```
+-----+-----+-----+-----+
| ID                | Stack Name | Stack Status  | Creation Time      | U
pdated Time |
+-----+-----+-----+-----+
| dbf46d1b-0b97-4d45-a0b3-9662a1eb6cf3 | stack      | CREATE_COMPLETE | 2015-10-13T15:27:20 | N
one          |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

6. Show the name and IP address of the instance...

```
$ openstack stack output show --all stack
```

```
[
  {
    "output_value": "stack-server-3nzfyfofu6d4",
    "description": "Name of the instance.",
    "output_key": "instance_name"
  },
  {
    "output_value": "10.4.31.106",
    "description": "IP address of the instance.",
    "output_key": "instance_ip"
  }
]
```

7. ... and compare with the output of the OpenStack client:

```
$ openstack server list
```

ID	Name	Status	Networks
0fc2af0c-ae79-4d22-8f36-9e860c257da5	<u>stack-server-3nzyfufu6d4</u>	ACTIVE	public= <u>10.4.31.10</u>



# Methods to send API requests

- cURL

A **command-line tool** that lets you send HTTP requests and receive responses.

- OpenStack command-line client

The OpenStack project provides a **command-line client** that enables you to access APIs through easy-to-use commands

- REST clients

Browser-based graphical interfaces for REST in Firefox and Chrome, AdvancedRestClient, Postman, etc

- OpenStack Python Software Development Kit (SDK)

Use this SDK to write **Python automation scripts** that create and manage resources in your OpenStack cloud. The SDK implements Python bindings to the OpenStack API, which enables you to perform automation tasks in Python by making calls on Python objects rather than making REST calls directly

# Openstack Toolkits 1/2

- **Go**
  - [Gophercloud](#) provides a Go binding to OpenStack cloud APIs.
- **Java**
  - [OpenStack4j](#) A fluent Java OpenStack API.
  - [OpenStack Java SDK](#) is a Java binding for the OpenStack APIs.
  - [Apache jclouds](#) is a Multi-cloud SDK for Java with OpenStack support
- **JavaScript**
  - [pkgcloud](#) is a Multi-cloud library for Node.js that supports OpenStack
  - [jstack](#) is a JavaScript client library for the OpenStack API.
  - [js-openclient](#) is a very opinionated core client which can be used in either Node.js or in the browser (browser support not yet complete) to communicate with a RESTful APIs, including but not limited to any OpenStack-compatible API. (last updated 2015)
  - [node-openstack-wrapper](#) is a convenience wrapper for many of Openstack's common features with a focus on projects/tenants.
- **.NET**
  - [OpenStack.NET](#) is a .NET SDK for OpenStack.

# Openstack Toolkits 2/2

- **PHP**

- [php-opencloud/openstack](#) is a PHP SDK for OpenStack. It is actively maintained and supports latest OpenStack identity api (Keystone v3).
- SDK covers following api: BlockStorage v2, Compute v2, Identity Keystone v2 and v3, Images v2, Networking v2 including Layer3 extension and Security Groups extension, Object Storage v1, Gnocchi v1

- **Python**

- [OpenStack Shade](#) shade is a simple client library for operating OpenStack clouds that abstracts deployer differences
- The [SDK-Development/PythonOpenStackSDK](#) project is a proposed solution to offering an SDK that provides a single point of entry for consumers, and a base from which other tools can be built upon, such as command-line interfaces.
- The [OpenStackClients](#) are the native Python bindings for the OpenStack APIs. They are used to implement the command-line interfaces (which ship with the library).
- [Apache libcloud](#) is a Python library that abstracts away differences among multiple cloud provider APIs.

- **Ruby**

- [fog](#) is a Multi-cloud Ruby library with support for OpenStack
- [Misty](#) is a dedicated HTTP client for OpenStack APIs
- [Aviator](#) An elegantly designed OpenStack SDK for Ruby (hasn't been updated since 2015)

# Keystone Tokens

- Tokens are used to authenticate and authorize interactions with the various OpenStack APIs
- Tokens can express your authorization in different **scopes**
- A user may have different sets of roles, in different projects, and in different domains
- While tokens always express one identity, they may only ever express **one set of roles in one authorization scope** at a time

# Unscoped Tokens

- An unscoped token contains neither roles, a project scope, nor a domain scope
- Their primary use case is simply to prove your identity to keystone at a later time, usually to generate scoped tokens, without repeatedly presenting your original credentials

The following conditions must be met to receive an unscoped token:

- You must not specify an authorization scope in your authentication request (for example, on the command line with arguments such as **--os-project-name** or **--os-domain-id**),
- Your identity must not have a "default project" associated with it that you also have role assignments, and thus authorization, upon.

# Scoped Tokens

- **Project-scoped**
  - they express the authorization to operate in a specific tenancy (project) of the cloud and are useful to authenticate when working with most other services
- **Domain-scoped**
  - they express the authorization to operate a domain-level, above that of the user and projects contained therein (domain-level administrator)
  - They contain a limited-service catalog
  - They can also be used to work with domain-level concerns in other services, such as to configure domain-wide quotas that apply to all users or projects in a specific domain

# Identity Concepts

- **Authentication.** The process of confirming the identity of a user. To confirm an incoming request, OpenStack Identity validates a **set of credentials** that users supply. Initially, these credentials are a **username and password**, or a **username and API key**. When OpenStack Identity validates user credentials, it issues an **authentication token**. Users provide the token in subsequent requests.
- **Endpoint.** A network-accessible address, usually **a URL, through which you can access a service**.
- **Regions, Domains and Groups.** Identity services API v3 entities. A region represents a general division in an OpenStack deployment. Domains are a collection of projects and users that define administrative boundaries for managing Identity entities. Groups are a collection of users owned by a domain.
- **Project.** A **container that groups or isolates resources** or identity objects. Depending on the service operator, a project might map to a customer, account, organization, or tenant.
- **Service.** **An OpenStack service**, such as Compute (nova), Object Storage (swift), or Image service (glance), **that provides one or more endpoints through which users can access resources and perform operations**.
- **Token.** An alpha-numeric text string that enables access to OpenStack APIs and resources. A token may be revoked at any time and is **valid for a finite duration**.
- **User.** A digital representation of a person, system, or service that uses OpenStack cloud services.

# Using the OpenStack APIs

1. Issue an authentication request with a payload of credentials to OpenStack Identity to get an authentication token

Credentials are usually a combination of your username and password, and optionally, the name or ID of the project of your cloud.

2. When you send API requests, you include the token in the X-Auth-Token header.

A token is valid for a limited time before it expires. A token can also become invalid for other reasons. For example, if the roles for a user change, existing tokens for that user are no longer valid.

<https://developer.openstack.org/api-guide/quick-start/api-quick-start.html>



## Request

```
POST /identity/v3/auth/tokens HTTP/1.1
HOST: 127.0.0.1:8080
content-type: application/json
content-length: 402
```

```
{
  "auth": {
    "identity": {
      "methods": [
        "password"
      ],
      "password": {
        "user": {
          "name": "demo",
          "domain": {
            "name": "Default"
          }
        },
        "password": "devstack"
      }
    },
    "scope": {
      "project": {
        "id": "9176c2c2a1144a97bcf0de413d7a3cd9"
      }
    }
  }
}
```

*Project demo*

# Scoped Token – Example

## Response

```
date: Tue, 01 May 2018 09:50:53 GMT
server: Apache/2.4.29 (Ubuntu)
x-subject-token: gAAAAABa6Dh-eyy0aHgVnbTXAu48Y3mPA6c6pMuhuX4HhEqhYkD6jEm1YIA28kYD
RebozMgYWRzchfKVWV0pv8AA_UoAh2BfVV2THGgW3MmDiVesVS9oAZ6_HCWueMI11r1WZeEZ8pbL40Seu
HbwrSM8xekyM6xOuw1ZA_k5QDB5JiqXv-UdzA
vary: X-Auth-Token
content-type: application/json
content-length: 3397
x-openstack-request-id: req-02d627b9-2a6b-44f8-8d04-66671ad0a35d
connection: close
{
  body
}
```

Catalog that includes endpoints to access services, e.g.:

<http://10.0.2.15/compute/v2.1>

Catalog  
received in  
the  
response  
to the  
Token  
request

9176c2c2a1144a9  
7bcf0de413d7a3c  
d9 = Project\_ID

```
devstack@openstack:~/devstack$ openstack catalog list
```

Name	Type	Endpoints
keystone	identity	RegionOne admin: http://10.0.2.15/identity RegionOne public: http://10.0.2.15/identity
cinderv2	volumev2	RegionOne public: http://10.0.2.15/volume/v2/9176c2c2a1144a97bcf0de413d7a3cd9
cinderv3	volumev3	RegionOne public: http://10.0.2.15/volume/v3/9176c2c2a1144a97bcf0de413d7a3cd9
placement	placement	RegionOne public: http://10.0.2.15/placement
nova	compute	RegionOne public: http://10.0.2.15/compute/v2.1
cinder	volume	RegionOne public: http://10.0.2.15/volume/v1/9176c2c2a1144a97bcf0de413d7a3cd9
nova_legacy	compute_legacy	RegionOne public: http://10.0.2.15/compute/v2/9176c2c2a1144a97bcf0de413d7a3cd9
neutron	network	RegionOne public: http://10.0.2.15:9696/
glance	image	RegionOne public: http://10.0.2.15/image
cinder	block-storage	RegionOne public: http://10.0.2.15/volume/

## Request

# Scoped Token – Example

```
GET /compute/v2.1/servers HTTP/1.1
```

```
HOST: 127.0.0.1:8080
```

```
x-auth-token: gAAAAABa6Dh-eyy0aHgVnbTXAu48Y3mPA6c6pMuhuX4HhEqhYkD6jEm1YIA28kYDRebozMgYWRzchfKVWVOpv8AA_UoAh2BfVV2THGgW3MmDiVesVS9oAZ6_HCWueMI11r1WZeEZ8pbL40SeuHbwrSM8xekyM6x0wuw1ZA_k5QDB5JiqXv-UdzA
```

## Response

```
date: Tue, 01 May 2018 09:55:50 GMT
```

```
server: Apache/2.4.29 (Ubuntu)
```

```
content-length: 919
```

```
content-type: application/json
```

```
openstack-api-version: compute 2.1
```

```
x-openstack-nova-api-version: 2.1
```

```
vary: OpenStack-API-Version,X-OpenStack-Nova-API-Version
```

```
x-openstack-request-id: req-c6ec4efe-619a-4eaf-af4c-125ba93db35f
```

```
x-compute-request-id: req-c6ec4efe-619a-4eaf-af4c-125ba93db35f
```

```
connection: close
```

- A **self** link contains a versioned link to the resource. Use these links when the link is followed immediately.
- A **bookmark** link provides a permanent link to a resource that is appropriate for long term storage.

```
{  
  ...  
  "servers": [  
    {  
      "id": "5622e010-538e-49f4-85e2-560727253b21",  
      "links": [  
        {  
          "href": "http://127.0.0.1:8080/compute/v2.1/servers/5622e010-538e-49f4-85e2-560727253b21",  
          "rel": "self"  
        },  
        {  
          "href": "http://127.0.0.1:8080/compute/servers/5622e010-538e-49f4-85e2-560727253b21",  
          "rel": "bookmark"  
        }  
      ],  
      "name": "test_instance_2"  
    },  
  ]  
}
```

# Cinder and Nova APIs Usage Examples

```
GET /volume/v3/9176c2c2a1144a97bcf0de413d7a3cd9/volumes HTTP/1.1
HOST: 127.0.0.1:8080
x-auth-token: gAAAAABa6Cu5x90_6YWjPW9DYthkMFYVU7lfnLDkERcLnzq
UZ7a3PJ8-TzxH5GtLhzJHSI9m_u5X6JL41sRwhaF27wTSRH9pwu3T8f-wZTNn
2mhdvBSva0lJ_7ISIAE8mPorDNE04SzxvIMGhZtIzVIYKKn9EVNXhT3QmpsZR
sIB6YE95dL50x4
```

Volumes associated to **user demo**,  
**project demo** in **Cinder**

Create a new instance in **Nova**

**Response:** 202 Accepted

```
POST /compute/v2.1/servers HTTP/1.1
HOST: 127.0.0.1:8080
x-auth-token: gAAAAABa6Cu5x90_6YWjPW9DYthkMFYVU7lfnLDkERcLnzqUZ7a3PJ8-TzxH5GtLhzJH
SI9m_u5X6JL41sRwhaF27wTSRH9pwu3T8f-wZTNn2mhdvBSva0lJ_7ISIAE8mPorDNE04SzxvIMGhZtIzV
IYKKn9EVNXhT3QmpsZRsIB6YE95dL50x4
content-type: application/json
content-length: 191
{
  "server": {
    "name": "test_instance_auto",
    "imageRef": "6eda4305-d711-4686-85bc-c47856c2e2c8",
    "flavorRef": http://openstack.example.com/flavors/1
  }
}
```