

Programação Avançada (ProgA)

Chapter #1 - Processes

Patrício Domingues
ESTG/IPLeiria, 2019

✓ Program

- An executable file (.exe in Windows)

✓ Process

- A running program
- Several processes can be running the same program
 - Example: chrome.exe
- Each process
 - Program counter (PC)
 - Owner, PID
 - Security attributes
 - Address space

| Processes | Performance | App history | Startup | Users | Details | Services |
|------------------------|-------------|-------------|----------|------------|---------|----------|
| Name | 8% CPU | 82% Memory | 0% Disk | 0% Network | | |
| Git for Windows | 0% | 0,1 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0,6% | 130,9 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0,9% | 113,7 MB | 0,1 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0,4% | 82,7 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0,3% | 70,2 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0,2% | 69,9 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0% | 69,7 MB | 0,1 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0,3% | 63,1 MB | 0,1 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0% | 46,8 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0% | 37,5 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0,6% | 31,9 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0% | 22,0 MB | 0 MB/s | 0 Mbps | | |
| Google Chrome (32 bit) | 0% | 14,1 MB | 0 MB/s | 0 Mbps | | |

Processes in Linux (#1)

✓ top Número de tasks
 (~processos)

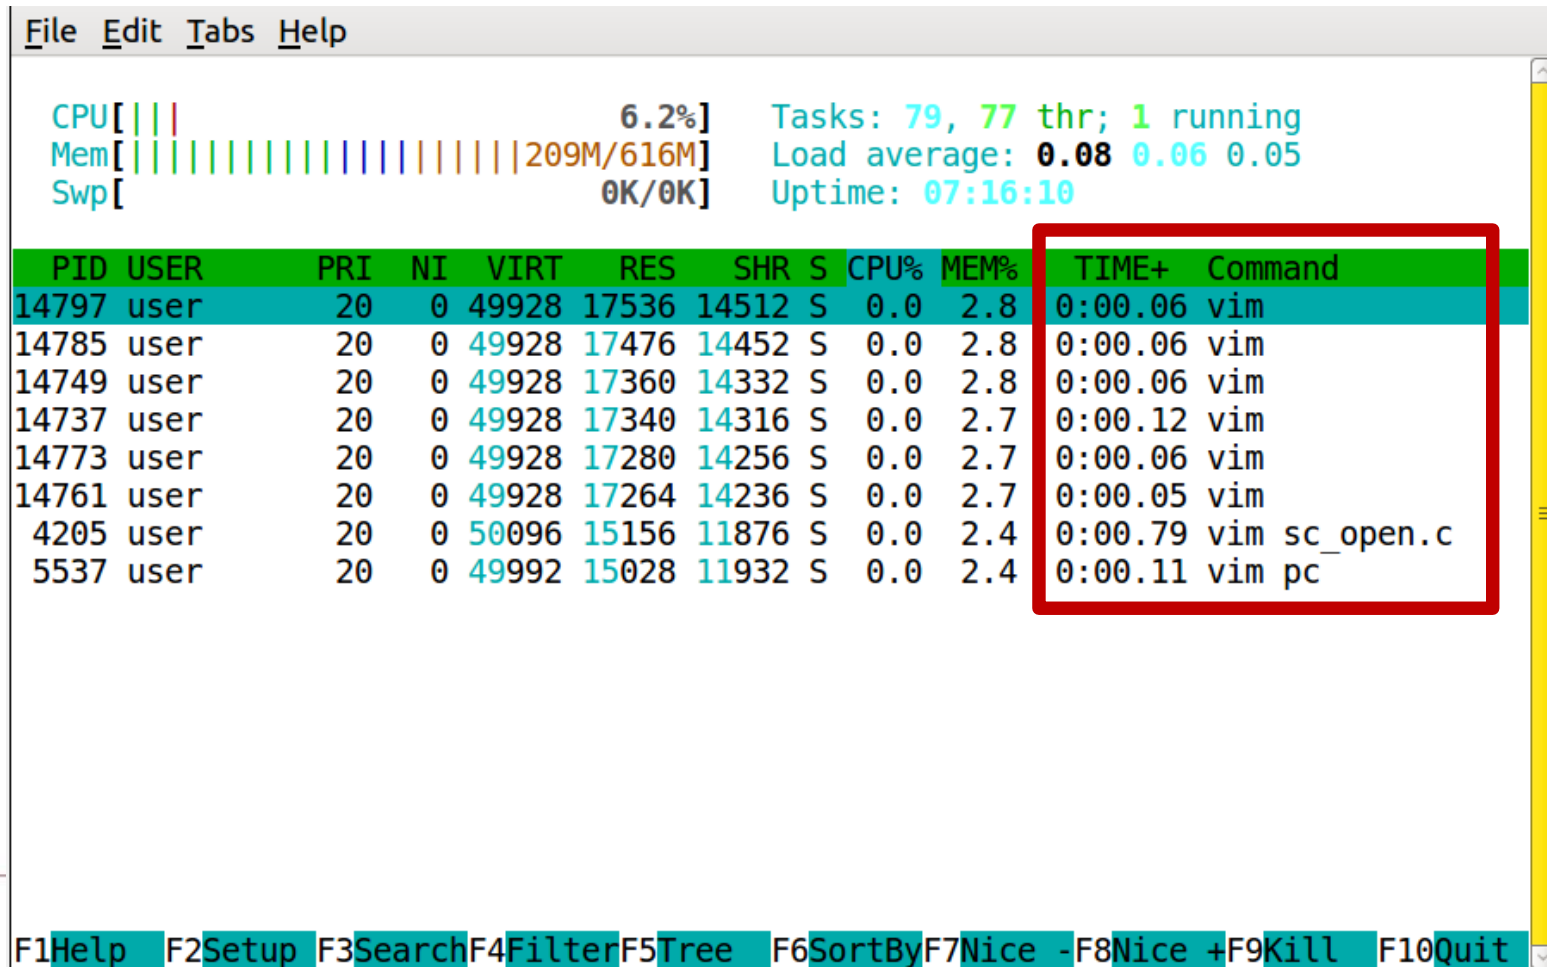
Carga média do sistema

| File Edit Tabs Help | | | | | | | | | | |
|--|------|----|----|--------|-------|-------|---|------|------|----------------------|
| top - 09:26:08 up 7:10, 1 user, load average: 0.10, 0.04, 0.05 | | | | | | | | | | |
| Tasks: 157 total, 1 running, 156 sleeping, 0 stopped, 0 zombie | | | | | | | | | | |
| %Cpu(s): 0.7 us, 0.7 sy, 0.0 ni, 98.3 id, 0.3 wa, 0.0 hi, 0.0 si, 0.0 st | | | | | | | | | | |
| KiB Mem : 630988 total, 31236 free, 194964 used, 404788 buff/cache | | | | | | | | | | |
| KiB Swap: 0 total, 0 free, 0 used. 406196 avail Mem | | | | | | | | | | |
| PID | USER | PR | NI | VIRT | RES | SHR | S | %CPU | %MEM | TIME+ COMMAND |
| 1996 | root | 20 | 0 | 253324 | 52416 | 9892 | S | 1.0 | 8.3 | 0:31.06 Xorg |
| 2550 | user | 20 | 0 | 75592 | 25740 | 19148 | S | 0.0 | 4.1 | 0:39.78 vmtoolsd |
| 2540 | user | 20 | 0 | 116076 | 25256 | 19776 | S | 0.0 | 4.0 | 0:03.64 pcmanfm |
| 2707 | user | 20 | 0 | 99088 | 24160 | 18820 | S | 1.0 | 3.8 | 0:09.60 lxterminal |
| 2454 | user | 20 | 0 | 95912 | 23544 | 18612 | S | 0.0 | 3.7 | 0:00.86 lxsession |
| 2538 | user | 20 | 0 | 107580 | 23544 | 19428 | S | 0.0 | 3.7 | 0:21.26 lxpanel |
| 2553 | user | 20 | 0 | 125216 | 22904 | 19044 | S | 0.0 | 3.6 | 0:01.10 nm-applet |
| 14797 | user | 20 | 0 | 49928 | 17536 | 14512 | S | 0.0 | 2.8 | 0:00.06 vim |
| 14785 | user | 20 | 0 | 49928 | 17476 | 14452 | S | 0.0 | 2.8 | 0:00.06 vim |
| 2535 | user | 20 | 0 | 73604 | 17428 | 13056 | S | 0.0 | 2.8 | 0:05.07 openbox |
| 14749 | user | 20 | 0 | 49928 | 17360 | 14332 | S | 0.0 | 2.8 | 0:00.06 vim |
| 14737 | user | 20 | 0 | 49928 | 17340 | 14316 | S | 0.0 | 2.7 | 0:00.12 vim |
| 14773 | user | 20 | 0 | 49928 | 17280 | 14256 | S | 0.0 | 2.7 | 0:00.06 vim |
| 14761 | user | 20 | 0 | 49928 | 17264 | 14236 | S | 0.0 | 2.7 | 0:00.05 vim |
| 2562 | user | 20 | 0 | 77116 | 16508 | 14308 | S | 0.0 | 2.6 | 0:00.24 xfce4-power+ |
| 4205 | user | 20 | 0 | 50096 | 15156 | 11876 | S | 0.0 | 2.4 | 0:00.79 vim |
| 5537 | user | 20 | 0 | 49992 | 15028 | 11932 | S | 0.0 | 2.4 | 0:00.11 vim |

Processes no Linux (#2)

✓ htop

- by default, htop is not installed in ubuntu
- `sudo apt-get install htop`



Processes in Linux (#3)

- The pseudo filesystem /proc keeps the data regarding the activity of the system
- For each process, there is a subdir in /proc
 - The name of the subdir is the PID of the process
 - NOTE: in bash, \$\$ holds the PID of the process running bash

```
user@ubuntu: /proc
File Edit Tabs Help
1249 181 2154 2620 53 73 fb scsi
13 1824 22 2625 54 74 filesystems self
1323 1837 2244 2632 541 75 fs slabinfo
14 1842 2246 2635 544 7520 interrupts softirqs
14179 1843 23 2648 55 7535 iomem stat
14264 1846 24 2653 5537 76 ioports swaps
14725 1856 240 2658 558 78 irq sys
14736 1874 2445 2663 56 8 kallsyms sysrq-trigger
14737 1877 2446 2678 57 8436 kcore sysvipc
14738 1881 2454 2707 58 9 keys thread-self
14749 1891 2502 2708 586 92 key-users timer_list
14750 1898 2505 28 59 93 kmsg timer_stats
14761 19 2506 29 598 94 kpagecgroup tty
14762 1902 2518 3 60 96 kpagecount uptime
14773 191 2523 30 61 acpi kpageflags version
14774 1966 2535 31 62 asound loadavg version_signature
14785 1978 2538 337 63 buddyinfo locks vmallocinfo
14786 199 2540 4190 64 bus mdstat vmstat
14797 1996 2550 4205 643 cgroups meminfo zoneinfo
14820 2 2553 46 65 cmdline misc
14889 20 2562 47 66 consoles modules
14922 204 2575 49 67 cpuinfo mounts
14956 206 2576 5 68 crypto mtrr
user@ubuntu:/proc$
```

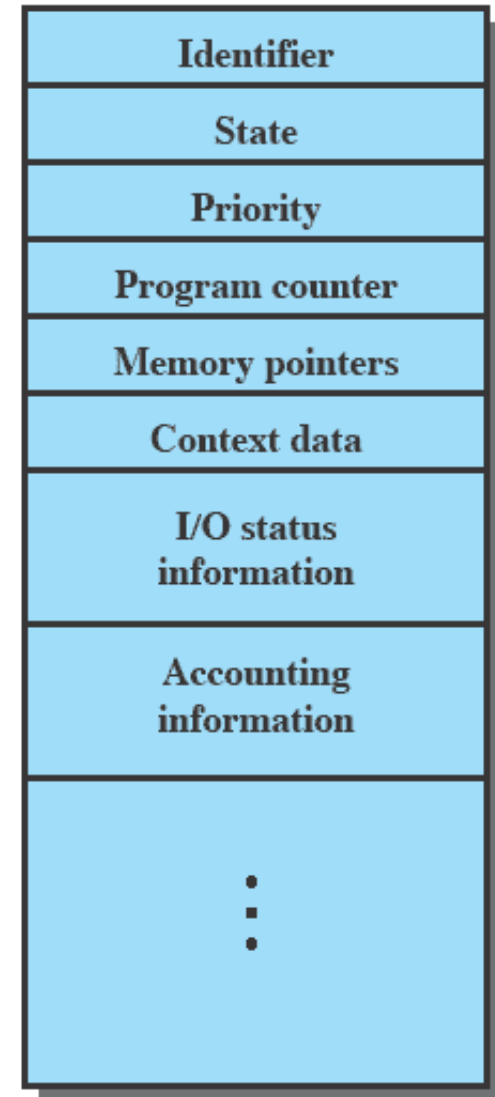
- `-n 2`: every 2 seconds
- ```
watch -n 2 cat /proc/stat
```

[illegible]



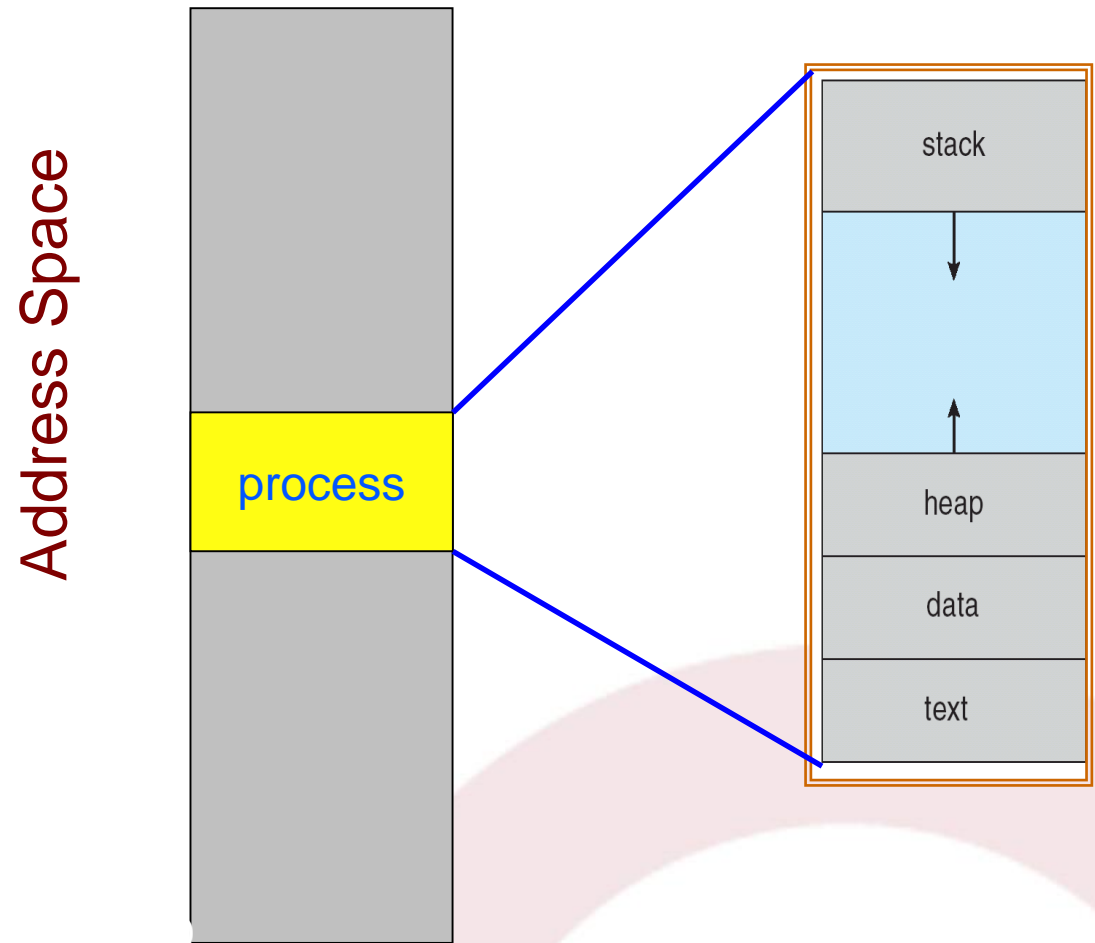
# Process Control Block - PCB

- ✓ Process Control Block – PCB
  - OS struct that keeps the data of a process
  - PCB is essential for multitasking
    - It allows the process to be interrupted and then resumed from its precise execution point



# Process in memory

- ✓ A process has several (logical) segments in memory
  - Text segment
  - Data segment
  - Heap segment
  - Stack segment





- ✓ **Memory Image**
  - (logical) representation of the process in memory
  - Several segments
- ✓ **text segment**
  - Holds the code/instruction of the process
- ✓ **data segment**
  - Holds variables in two subsections
    - .data: global and “static” variables, initialized with a non-zero value
    - .bss: global and “static” variables, non-initialized

>>

## ✓ Heap segment

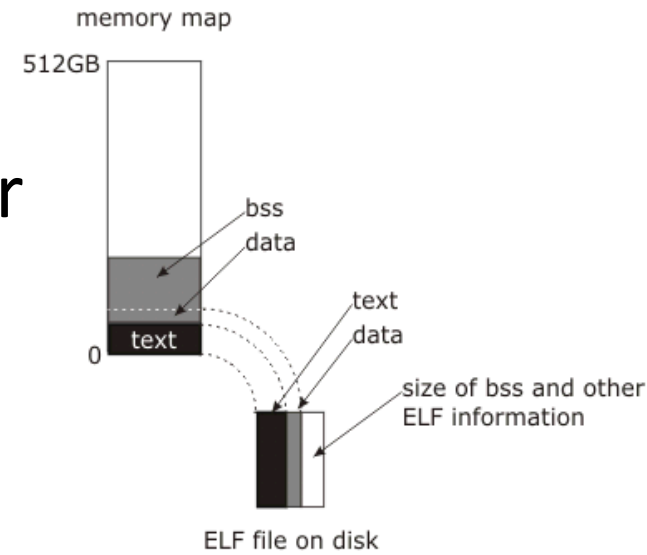
- Used for dynamic memory
  - Example: malloc, calloc, realloc,...
  - It has a variable length

## ✓ Stack segment

- Holds the so-called automatic/local variables
  - Variables that are automatically created and destroyed in functions and methods
  - It has a variable length

# Executable file

- ✓ A process executes...an executable
- ✓ BSS: Block Started by Symbol
  - Data segment that hold static and globals variables set with zero as their initial value
  - BSS does not take space in the executable file
- ✓ DATA segment
  - “static” and “globals” variables set with an initial value different from zero
  - The executable file has a DATA segment



# *size* command (1)

- **size** command
  - Unix tool that reports the length of each section and total size for a given executable file
- **Example**

```
#include <stdio.h>

int main(void){
 printf("Programa em C\n");
 return 0;
}
```
- **size a.exe**

```
text=877,data=256,bss=8,dec=1141,hex=475,filename=a.exe
```

  - $877 + 256 + 8 = 1141$  bytes (ou seja, 475 hexadecimal)
  - **Size of the executable file**
    - 7136 bytes

# size command (2)

- Example #2

```
#include <stdio.h>
char A[100000]; /* Vetor com 100000 chars = 100000 bytes */
int main(void){
 printf("Programa em C\n");
 return 0;
}
```

- `size a.exe`

`text=877,data=256,bss=100032,dec=101165,hex=18b2d,filename=a.exe`

–  $877 + 256 + 100032 = 101165$  bytes (ou seja, 18b2d hexadecimal)

– Size of the executable file

- 7145 bytes

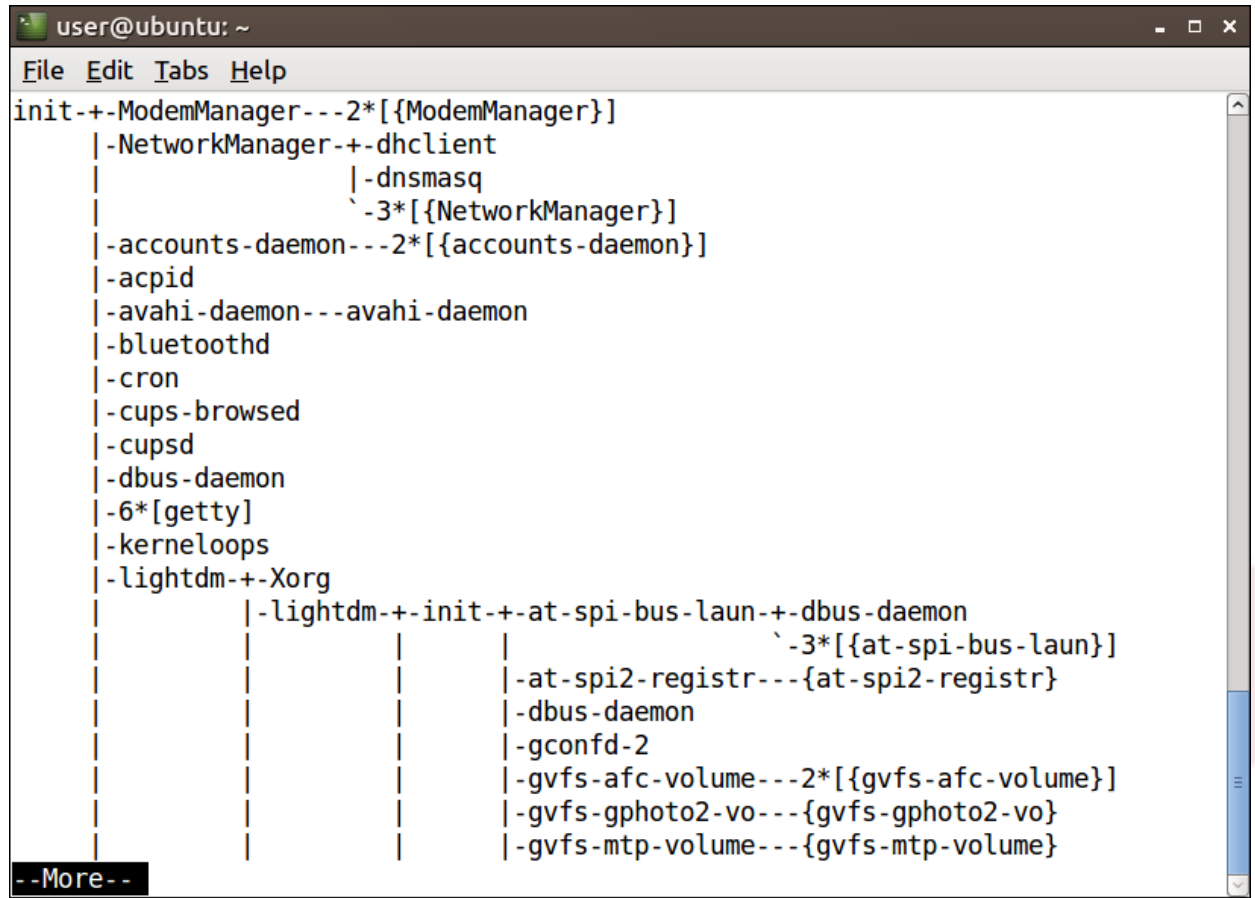
- Conclusion

- Despite the declaration of the A vector to hold 100000 chars, the executable only increased 9 bytes (7136 to 7145 bytes)
- Why? A is kept in BSS

# Processes in Unix (1)

## ✓ Processes in Unix

- All are descendant from the *init* (PID=1) process
- A process is created...by another process
  - Parent / son relationship
  - Tree of processes

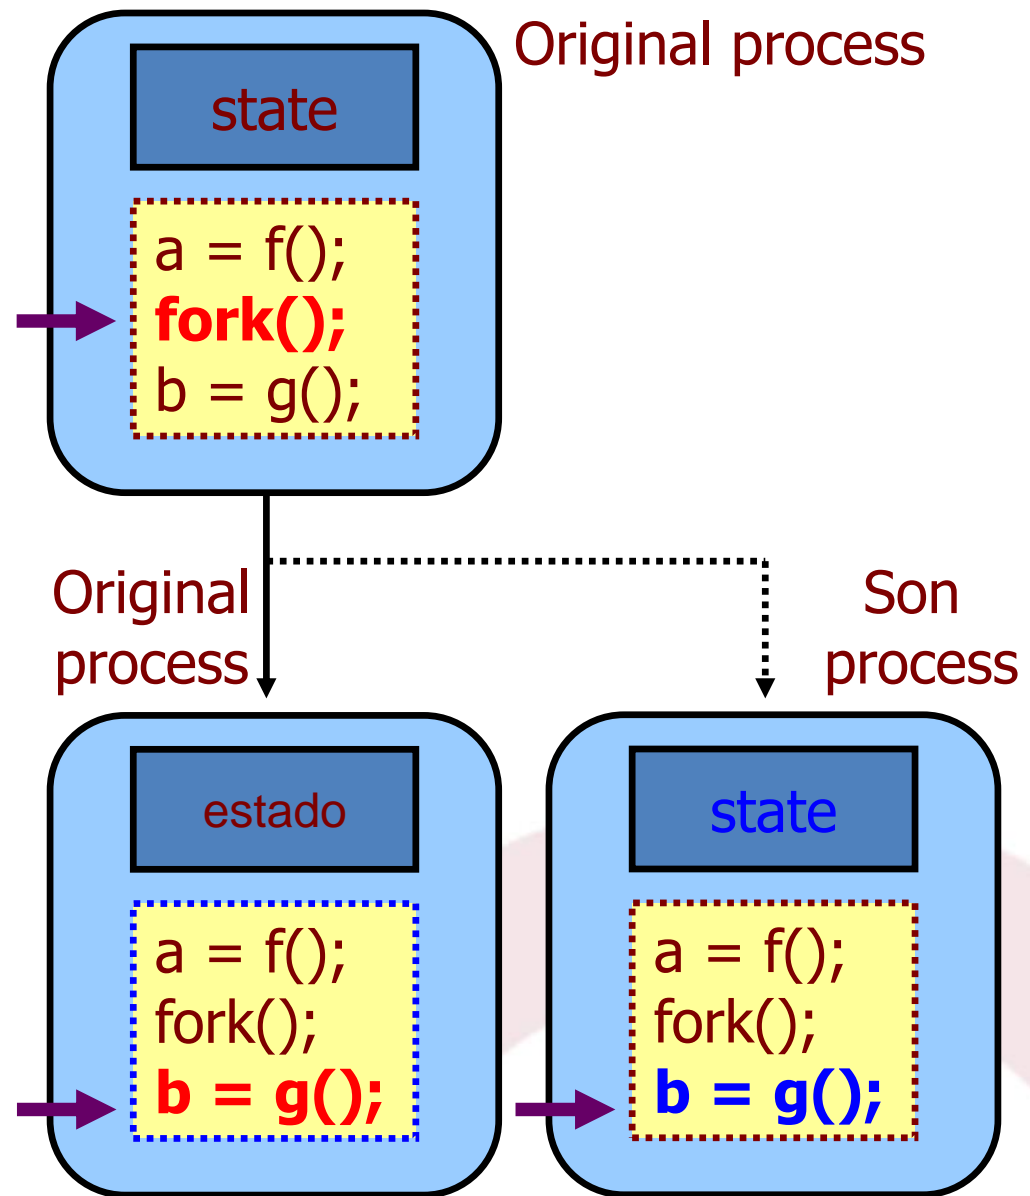


```
user@ubuntu: ~
File Edit Tabs Help
init--ModemManager---2*[{ModemManager}]
 |-NetworkManager--dhclient
 | |-dnsmasq
 | `--3*[{NetworkManager}]
 |-accounts-daemon---2*[{accounts-daemon}]
 |-acpid
 |-avahi-daemon---avahi-daemon
 |-bluetoothd
 |-cron
 |-cups-browsed
 |-cupsd
 |-dbus-daemon
 |-6*[getty]
 |-kerneloops
 |-lightdm--Xorg
 | |-lightdm--init--at-spi-bus-laun--dbus-daemon
 | | `--3*[{at-spi-bus-laun}]
 | | |-at-spi2-registr---{at-spi2-registr}
 | | |-dbus-daemon
 | | |-gconfd-2
 | | |-gvfs-afc-volume---2*[{gvfs-afc-volume}]
 | | |-gvfs-gphoto2-vo---{gvfs-gphoto2-vo}
 | | |-gvfs-mtp-volume---{gvfs-mtp-volume}
 --More--
```



# Process model in UNIX

- `fork()`
  - system call to create a process
  - It is called by the parent process
- The son process inherits all characteristics of the parent process
  - Variables, program counter, open files, allocated memory, etc.
  - The son is a snapshot of the parent
- After “fork”, each process executes separately
  - The change of a variable in one process does **not** reflect on the other one



# fork system call

- ✓ `pid_t fork(void);`
- ✓ The fork system call returns an integer:
  - 0 to the newly created son process
  - $> 0$  to the calling parent process
    - The return value corresponds to the PID of the newly created process
- ✓ It can also returns  $-1$  if an error has occurred
  - `errno` holds the error code
  - `strerror(errno)` returns the error string

# Example – fork system call

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/types.h>
```

```
int main()
{
 pid_t id;
```

```
 id = fork(); /* returns 0: son process; > 0 to the parent */
 if (id == 0)
```

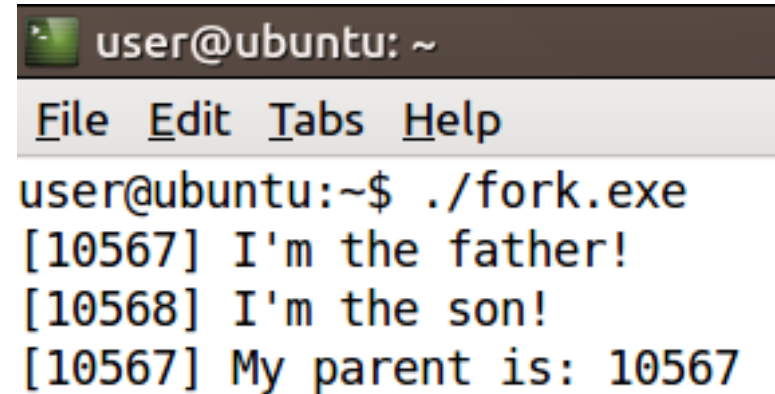
```
 { /* code only executed by the son process */
 printf("[%d] I'm the son!\n", getpid());
 printf("[%d] My parent is: %d\n", getppid(), getppid());
 }
```

```
 else if (id > 0)
```

```
 { /* code only executed by the parent process */
 printf("[%d] I'm the father!\n", getpid());
 wait(NULL);
 }
```

```
 return 0;
```

```
}
```



```
user@ubuntu: ~
File Edit Tabs Help
user@ubuntu:~$./fork.exe
[10567] I'm the father!
[10568] I'm the son!
[10567] My parent is: 10567
```

# Example – `fork` in a for loop

- ✓ How many processes are created by the following code?

```
#include <...>
int main(void){
 int i;
 for(i=0;i<3;i++){
 if(fork() == 0){
 printf("PID=%u\n", getpid());
 fflush(stdout);
 }
 }
 return 0;
}
```

Only newly created processes print their PID  
Answer: 7  
 $2^n - 1$ , with  $n=3$

# Memory cost of a fork

- ✓ The fork system call creates a new process by cloning the current one
  - Does this mean that all the memory assigned to the cloned process is copied/duplicated?
    - NO, otherwise, a fork will be computationally expensive...
  - Unix uses the “**copy-on-write**” mechanism

Copy-on-write mechanism >>

# Copy-on-Write (#1)

- ✓ The OS organizes memory in fixed-size blocks called *pages* (“*paged memory*”)
  - Usually, a page has 4 KiB
- ✓ To avoid memory duplication, after a fork, the father and son process share the same memory pages
  - So there is no duplication of memory content

**Continue >>**



# Copy-on-Write (#2)

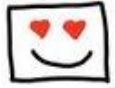
- ✓ The shared pages are marked as **readonly**
- ✓ When one of the process tries to write a shared page (e.g, to change the value of a variable: `i++`)
  - The OS flags it and creates a copy of the page
  - Thereafter, both father and son processes have their own page
    - They still share the other pages
- ✓ Only written pages (“dirty”) are duplicated
  - This scheme is also known as *copy-on-demand*

# copy on write

JULIA EVANS  
@b0rk

On Linux, you start new processes using the `fork()` or `clone()` system call

calling fork gives you a child process that's a copy of you



parent



child

the cloned process has EXACTLY the same memory

- same heap
- same stack
- same memory maps

if the parent has 36B of memory, the child will too

copying all that memory every time we fork would be **slow** and a **waste of RAM**



often processes call `exec` right after `fork` which means they don't use the parent process's memory basically at all!

so Linux lets them share physical RAM and only copies the memory when one of them tries to **write**.



process

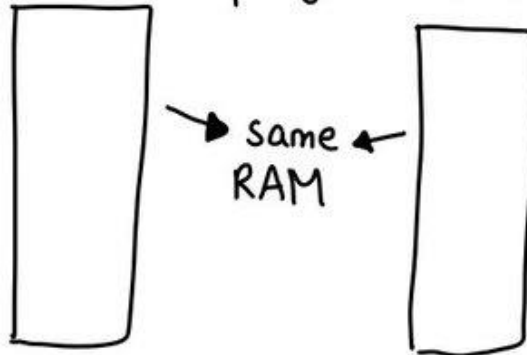
I'd like to change that memory

ok I'll make you your own copy!



Linux

Linux does this by giving both the processes identical page tables



but marks every page as **read only**

when a process tries to write to a shared memory address

- ① there's a **page fault**
- ② Linux makes a copy of the page & updates the page table
- ③ the process continues, blissfully ignorant



It's just like I have my own copy

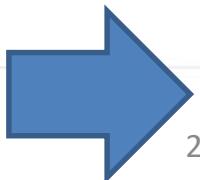
# Process ID (PID) – #1

- ✓ Process ID (PID)
  - Integer identifier of a process
- ✓ The Linux kernel allocates process IDs to processes in a strictly linear fashion.
  - If pid 37 is the highest number currently allocated, pid 38 will be allocated next, even if the process last
- ✓ For compatibility with old UNIX, the max value for PID is 32768 (16-bit signed int)
- ✓ This value can be changed
  - `/proc/sys/kernel/pid_max`

# Process ID (PID) - #2

- ✓ Within a C program, the PID of the calling process is returned with `getpid()`
  - `pid_t getpid(void);`
- ✓ The PID of the parent process is available through `getppid()`
  - `pid_t getppid(void);`

```
printf ("My pid=%jd\n", getpid ());
printf ("Parent's pid=%jd\n", getppid ());
```



# What is `intmax_t`?

- ✓ `intmax_t`
  - Signed integer capable of representing any value of any signed integer type
- ✓ Defined by C99 and C11
  - It requires `#include <inttypes.h>`
- ✓ The format string for `printf` is `%jd` (decimal)
  - There is also the format string for hexadecimal representation
    - `%jx`

# Running applications

## ✓ But...

- If all new processes execute the code of their parents, how can new applications be run?
  - The **fork** system call creates a clone of the parent process

## ✓ How do we run an application?

- `vim`, `ps`, `ls`, `find`, `firefox`,...

## ✓ Answer

- The “`exec`” family of system calls
  - These syscalls replace the image of the calling process



# The *exec* family of system calls

## ✓ “exec” system calls

```
int execl(const char *path, const char *arg, ...);
int execlp(const char *file, const char *arg, ...);
int execlenv(const char *path, const char *arg, ..., char *const envp[]);
int execv(const char *path, char *const argv[]);
int execvp(const char *file, char *const argv[]);
```

## ✓ Usage of an exec system call is

- exec...(application\_to\_be\_run)

## ✓ “exec”

- Replaces the image of the current process by another one from a given executable
  - Functions with “p” are “PATH”-aware
  - Functions with “v” get their parameters from a vector of strings
  - Functions with “l” get their parameters from a list, where itens are separated by “,” and the list ends with NULL

## ✓ Example: **execl("/bin/ps", "ps", "aux",NULL);**

# Example - running “ls” (1)

- ✓ Running “ls -a” resorting to “execvp”
- ✓ Question
  - ✓ Why the “This cannot happen!” message?

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
int main(){
 if (execvp("ls", "ls", "-a", NULL) == -1)
 perror("Error executing ls: ");
 else
 printf("This cannot happen!\n");
 return 0;
}
```

# Example – running “ls” (2)

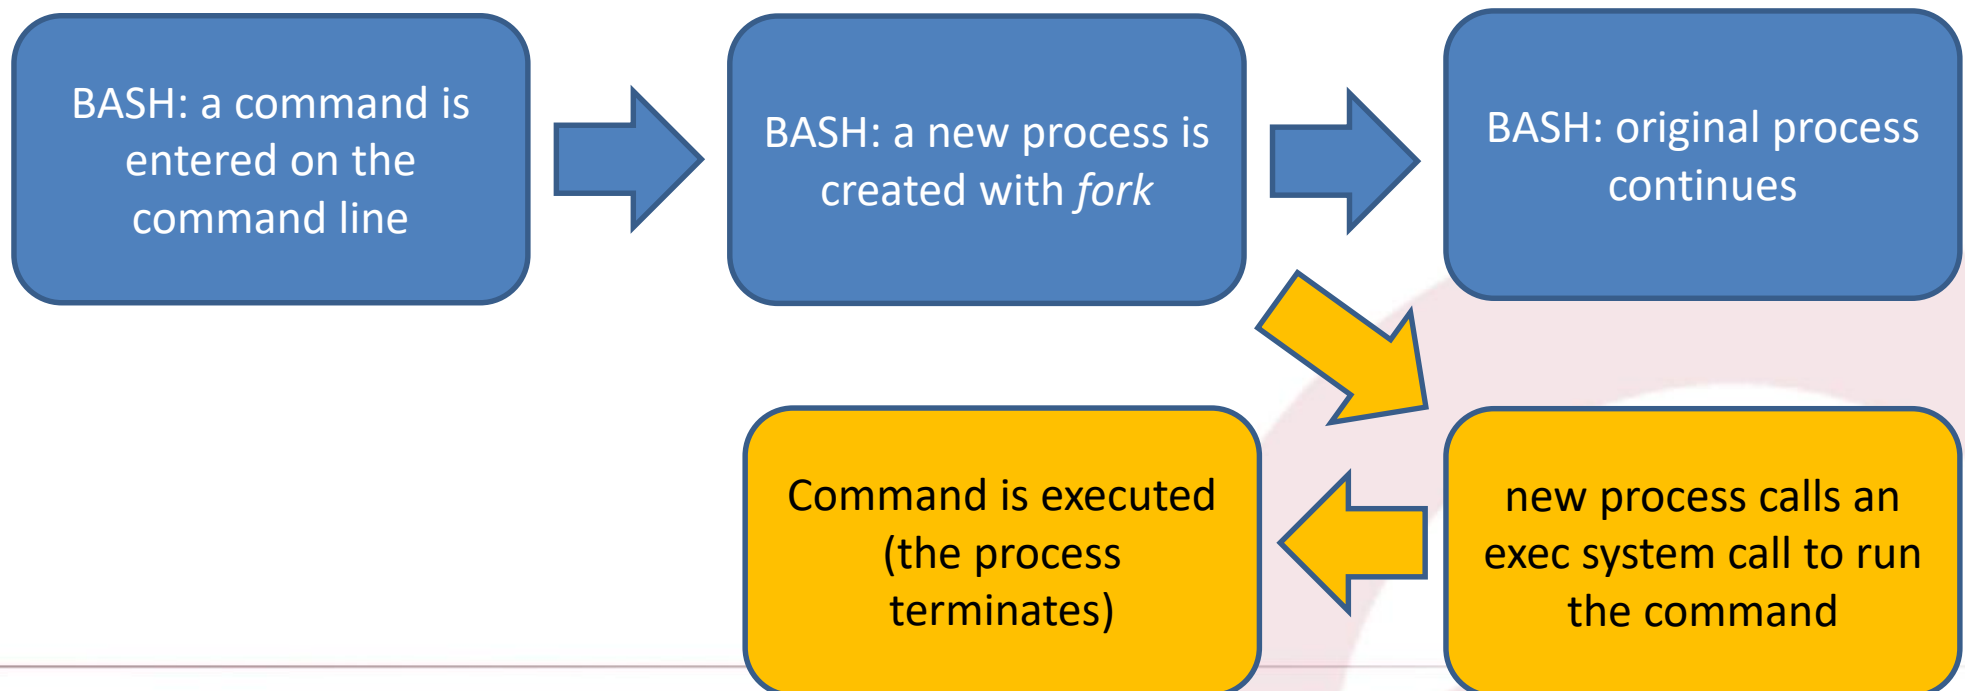
- ✓ Process launches “ls” via `execvp`
  - The exec system calls **NEVER** return when the execution is successful
  - The calling process image is replaced by the image of the executable called via “exec”
    - The calling process runs the executable
      - “ls” in our example
      - Therefore, `printf(“This cannot happen!”)` is removed from memory (as well as all the code of the calling process)
      - The code is replaced by the code of “ls -a”

# Executing a command

## ✓ Executing a command

### – Example with bash

- Applies to other shells (sh, zsh, etc.)
- 1<sup>st</sup> – fork
- 2<sup>nd</sup> – system call from the exec family



## ✓ *wait* and *waitpid*

```
pid_t wait(int *status);
```

```
pid_t waitpid(pid_t pid, int *status, int options);
```

- system calls used to synchronize a parent process with its children
- Wait for state changes on their children processes
  - Child is stopped (SIGSTOP) or terminates
  - Child is resumed by a signal (SIGCONT)
- Example
  - `wait(&status);` -- waits until a children process terminates
  - `waitpid(-1, &status, 0);` -- same as above

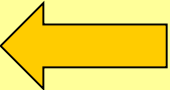
# Zombie process

- ✓ A child that terminates, but has not been waited for becomes a "zombie"
- ✓ The kernel maintains a minimal set of information about the zombie process
  - PID, termination status, resource usage information
- ✓ As long as a zombie is not removed from the system via `wait`, it will consume a slot in the kernel process table
- ✓ If a parent process terminates, then its "zombie" children (if any) are adopted by *init*, which automatically performs a `wait` to remove the zombies



# Creating zombies...

```
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
void worker() {
 printf("[%d] Hi, I'm a worker process! Going to die...\n",
 getpid());
}
int main()
{ int i;
 for (i=0; i<5; i++) {
 if (fork() == 0) {
 worker();
 exit(0);
 }
 }
 system("ps aux | grep -i zombie");
 printf("[%d] Big father is sleeping!\n", getpid());
 sleep(10);
 return 0;
}
```



Question: how many  
processes are created?

# Creating zombies...

## ✓ Results

```
user@ubuntu: ~/SO
File Edit Tabs Help
user@ubuntu:~/SO$./zombie.exe
[17512] Hi, I'm a worker process! Going to terminate...
[17513] Hi, I'm a worker process! Going to terminate...
[17514] Hi, I'm a worker process! Going to terminate...
[17511] Hi, I'm a worker process! Going to terminate...
[17510] Hi, I'm a worker process! Going to terminate...
user 17467 0.2 0.4 12348 4616 pts/16 S+ 12:59 0:00 vim zombie.c
user 17509 0.0 0.0 2024 276 pts/5 S+ 13:02 0:00 ./zombie.exe
user 17510 0.0 0.0 0 0 pts/5 Z+ 13:02 0:00 [zombie.exe] <defunct>
user 17511 0.0 0.0 0 0 pts/5 Z+ 13:02 0:00 [zombie.exe] <defunct>
user 17512 0.0 0.0 0 0 pts/5 Z+ 13:02 0:00 [zombie.exe] <defunct>
user 17513 0.0 0.0 0 0 pts/5 Z+ 13:02 0:00 [zombie.exe] <defunct>
user 17514 0.0 0.0 0 0 pts/5 Z+ 13:02 0:00 [zombie.exe] <defunct>
user 17515 0.0 0.0 2268 552 pts/5 S+ 13:02 0:00 sh -c ps aux | grep -i zombie
user 17517 0.0 0.0 4680 832 pts/5 S+ 13:02 0:00 grep -i zombie
[17509] Big father is sleeping!
user@ubuntu:~/SO$
```

# The system function

- ✓ In the zombie code, we have the following line of code

```
system("ps aux | grep -i zombie");
```

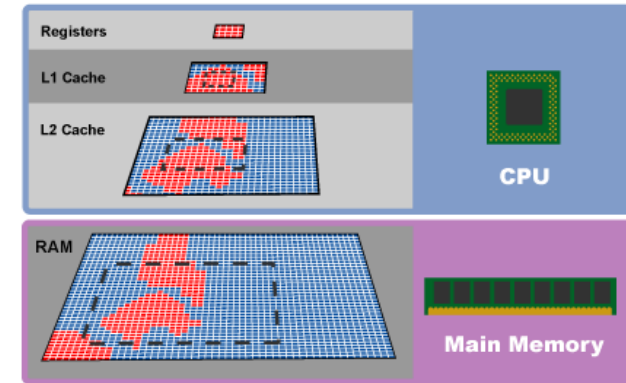
- ✓ system launches a shell that executes the command given as string
  - It executes **/bin/sh -c command\_line** and waits for its termination
  - The **shell** process is the one that actually executes the command line
- It is costly, since it has to i) **fork** a process and then ii) **exec** its image to execute the Shell

# CPU affinity (1)

- ✓ On a multicore system, the scheduler needs to decide which processes runs on each CPU
- ✓ Once a process is running on a CPU, it should remain there
  - Avoid the “cold cache” effect of moving a process to another CPU/core
- ✓ How to enforce CPU affinity at the programming level?

# CPU affinity (2)

- ✓ On a multicore system, the OS scheduler needs to decide which processes runs on each CPU
- ✓ Once a process is running on a CPU, the OS scheduler tries to keep it there
  - Avoid the “cold cache” effect of moving a process to another CPU/core
    - When a process moves to another CPU/core, the cache(s) of the CPU/core do not have content of the process



<http://bit.ly/256cAlr>

- ✓ CPU affinity of a process can be controlled programmatically
  - Hard affinity
- ✓ `int sched_setaffinity(pid_t pid, size_t setsize, const cpu_set_t *set);`
- ✓ `int sched_getaffinity(pid_t pid, size_t setsize, cpu_set_t *set);`
- ✓ `void CPU_SET(unsigned long cpu, cpu_set_t *set);`
- ✓ `void CPU_CLR(unsigned long cpu, cpu_set_t *set);`
- ✓ `int CPU_ISSET(unsigned long cpu, cpu_set_t *set);`
- ✓ `void CPU_ZERO(cpu_set_t *set);`

## ✓ Example

```
#define _GNU_SOURCE
#include <sched.h>
#include <stdio.h>
cpu_set_t set;
int ret, i;
CPU_ZERO (&set);
ret = sched_getaffinity(0, sizeof (cpu_set_t), &set);
if (ret == -1) {
 perror ("sched_getaffinity");
}
for (i=0; i < CPU_SETSIZE; i++) {
 int cpu;
 cpu = CPU_ISSET(i, &set);
 printf ("cpu=%i is %s\n", i, cpu?"set":"unset");
}
```



# Termination of a process

- ✓ Reason for a process to terminate
  - Regular termination
    - exit, return of main function, etc.
  - Process has exceeded maximum CPU time (e.g., “ulimit” from bash)
  - Not enough memory
  - I/O failure
  - Invalid instruction (e.g., “divide by zero”)
  - OS action
    - Deadlock or OOM (Out of Memory Killer)
  - User action
    - Kill -9 PID or killall -9 process\_name
  - ...

# the **exit** function

- ✓ The **exit** function terminates the calling process
  - `void exit(int status);`
- ✓ It returns the int **status** to the operating system
- ✓ The **exit** function can call other functions before terminating the process
  - It calls the function previously registered with the following functions
    - `atexit`
    - `on_exit`

# The `_exit` function

- ✓ What about to terminate the current process without calling the function registered with **`atexit(3)`** and **`on_exit(3)`**?
- ✓ Function **`_exit(2)`**
  - Note the leading underscore in the name



# ulimits (bash)

✓ The bash shell has an internal set of limits

– `ulimit`

- internal command

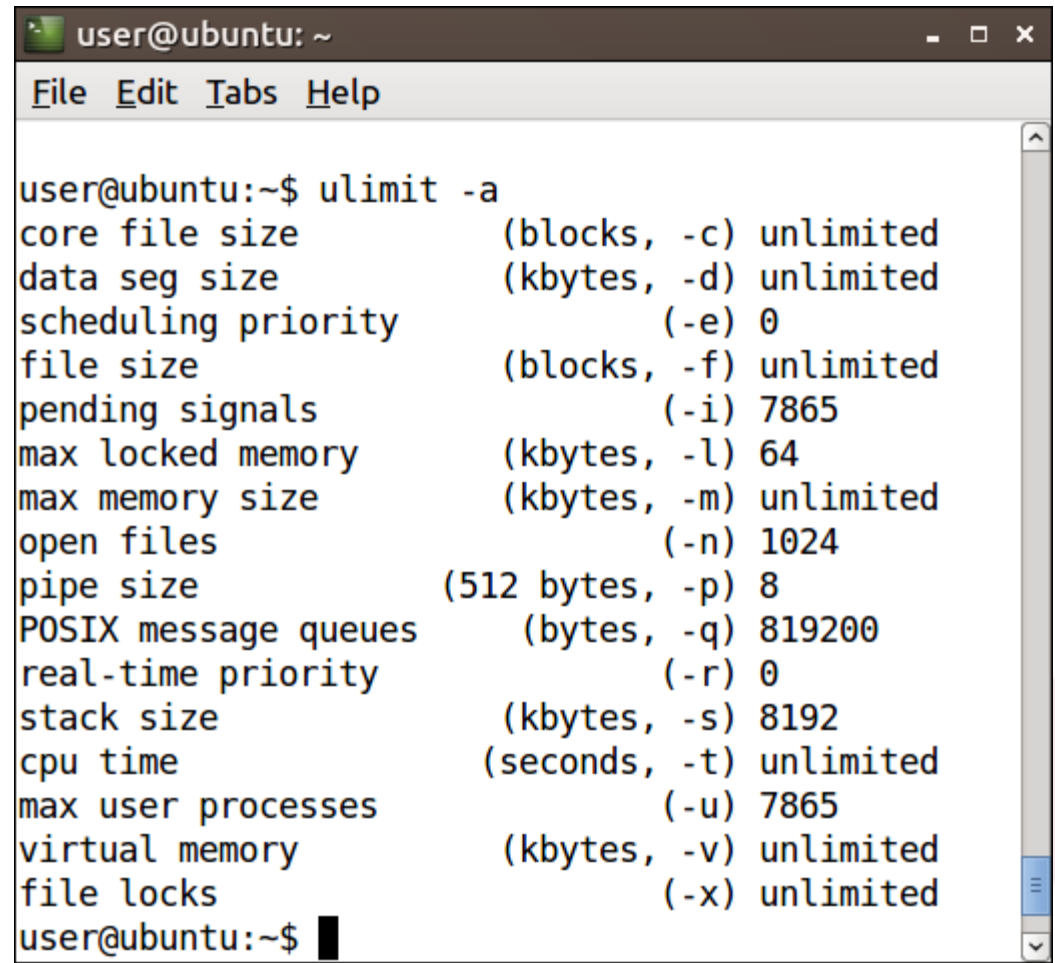
- Processed by bash, there is no `ulimit` executable

- There is no man for `ulimit`

- `help ulimit`

– `ulimit -a`

- List all limits for the current session



```
user@ubuntu: ~
File Edit Tabs Help

user@ubuntu:~$ ulimit -a
core file size (blocks, -c) unlimited
data seg size (kbytes, -d) unlimited
scheduling priority (-e) 0
file size (blocks, -f) unlimited
pending signals (-i) 7865
max locked memory (kbytes, -l) 64
max memory size (kbytes, -m) unlimited
open files (-n) 1024
pipe size (512 bytes, -p) 8
POSIX message queues (bytes, -q) 819200
real-time priority (-r) 0
stack size (kbytes, -s) 8192
cpu time (seconds, -t) unlimited
max user processes (-u) 7865
virtual memory (kbytes, -v) unlimited
file locks (-x) unlimited
user@ubuntu:~$
```

- Man pages
  - man 2 fork
  - man 2 exec
  - man 3 system
  - man 3 exit
  - man bash
  - help ulimit
- *Chapter 5 – Process management, “Linux System Programming”, Robert Love, 2013*
- printf format in C99 and C11  
<http://bit.ly/1OvGzGI>

