

Standard I/O

High-level



Patrício R. Domingues

Departamento de Eng Informática

ESTG/IPLeiria

Buffered I/O (1)

- Compared to CPU and RAM, persistent storage is slow
 - Example: access disk is around 10 ms for HDD and .1 ms for SSD (approximate values)
 - Much slower than RAM (~80 ns)
- Persistent storage is organized in blocks
 - A disk (SSD, HDD, etc.) is a set of blocks
 - We can only access disk to read/write **blocks**
 - Blocks are typically 512 bytes, 1024, 2048, 4096 and 8192 bytes
 - A disk does not read/write a byte. It reads/writes a block.

Performance *byte* vs. *block*

- Effect of block size on performance
 - Copy of 2 MiB (HDD)
 - 1 byte at the time
 - 1024 bytes at the time
 - 1024 bytes at the time is much faster than 1 byte at the time

Block size	Real time	User time	System time
1 byte	18.707 seconds	1.118 seconds	17.549 seconds
1,024 bytes	0.025 seconds	0.002 seconds	0.023 seconds
1,130 bytes	0.035 seconds	0.002 seconds	0.027 seconds

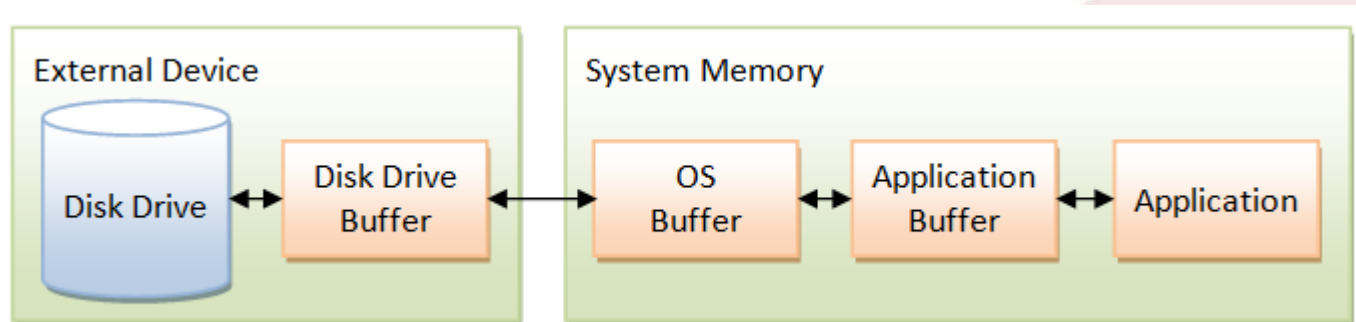
- Source: “Buffered I/O”, Chapter 3 - *Linux System Programming*, Robert Love, 2nd Edition, O’Reilly, 2013

Buffered I/O (2)

- At the storage level, read and write operations are done in blocks
 - The operating system and the device driver already **buffer** writes and reads
 - Kernel buffers, etc.
 - But programmers work with fields, lines and characters
 - Programmers work at the semantic level
 - It would be unproductive to force programmers to work with block...
- Solution
 - User buffered I/O
 - Standard I/O in C
 - `#include <stdio.h>`

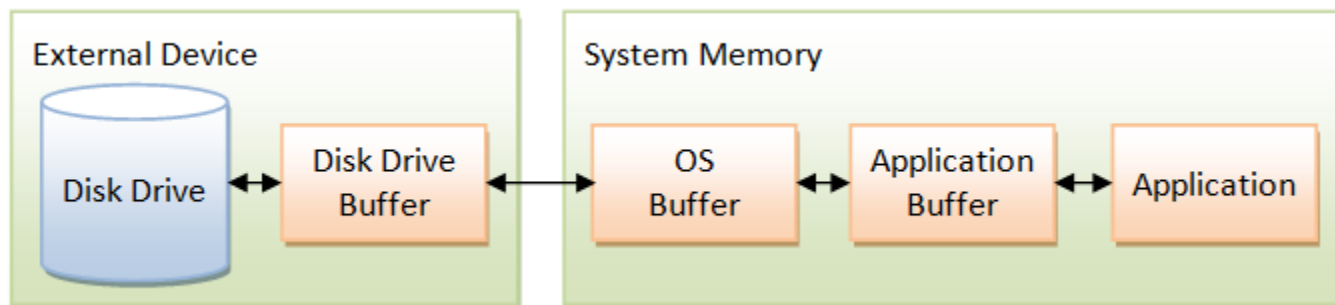
Buffered I/O (3)

- In C, user buffered I/O is available through STDIO
 - STDIO → standard Input/Output
 - FILE *, fopen, etc.
- Basics of user buffered I/O
 - The write and read operations use buffers that are set in **user space**
 - Buffer: zone of memory (e.g., chunk of 1024 bytes)



Buffered I/O (4)

- The functions of user buffered I/O (`fread`, `fwrite`, `fprintf`, etc.) transparently deal with writes and reads to/from persistent storage
- They use *write* and *read* buffers in user-space memory
- Whenever a buffer is *full* (write) or *empty* (read), appropriate actions are taken
 - Flush to disk (write); reload of buffer (read)



- Standard I/O routines operate on file through so-called *file pointers*
- File pointer type: FILE* (defined in <stdio.h>)
 - Example: FILE *f;
- Documentation often refers to an open FILE* as a *stream*
 - *Input stream*
 - *Output stream*
 - *Input/Output stream*

The *fopen* function (1)

- `FILE *fopen (const char *path, const char *mode) ;`
 - *path*: path for the file to be opened
 - *mode*: how the file will be used (**string**)
 - r: reading
 - r+: reading and writing (file position at the start of the file)
 - w: writing. If file exists, it is truncated to zero
 - w+: writing and reading. If file exists, it is truncated to zero. File position at the start of the file.
 - a: writing in append mode. File created if it does not exist. File position at the end of the file.
 - a+: writing and reading. All writes will append to the file.

The *fopen* function (2)

- *mode*: how the file will be used
 - Some OSes (e.g., Microsoft Windows) distinguish between text and binary files
 - Next slides
 - For this OS, the mode parameter needs to be complemented with **b** for binary files.
 - e.g. **rb**, **rb+**, etc.

- Differences **text vs. binary** files (e.g. Windows)
 - Newlines in **text mode**
 - `\n` → `\r\n` when file is written to disk;
 - `\r\n` → `\n` when file is read from disk in TEXT mode.
 - No conversion in BINARY mode.
 - `\r` is carriage return (hex code: 0x0D)
 - `\n` is line feed (hex code: 0x0A)
- Representation of numbers
 - Text mode: string (e.g., "64502" → five characters)
 - Binary: integer (e.g., 64502 → four bytes: 0xFBF6)

DOS/Windows vs. Unix

- Example
 - DOS/Windows format
 - `0x0d0A: \r\n`

```
linha 1
linha 2
linha 3
```



```
00000000: 6c69 6e68 6120 310d 0a6c 696e 6861 2032  linha 1..linha 2
00000010: 0d0a 6c69 6e68 6120 330d 0a        ..linha 3..
```

- Unix format
 - `0x0A: \n`

```
linha 1
linha 2
linha 3
```



```
00000000: 6c69 6e68 6120 310a 6c69 6e68 6120 320a  linha 1..linha 2.
00000010: 6c69 6e68 6120 330a                linha 3.
```

The *fopen* function (3)

- `FILE *fopen (const char *path,
const char *mode) ;`
 - *Returns NULL on error*
 - *Returns a pointer to the file if everything went ok*
 - *errno keeps the error*

Closing streams

- `int fclose (FILE *stream);`
 - Any buffered and not-yet-written data are first flushed. On success, `fclose()` returns 0.
 - On failure, it returns **EOF** and sets **errno** appropriately.
- `int fcloseall(void);`
 - Flush all streams of process and then close them all streams
 - *All streams* means that `stdin`, `stdout` and `stderr` are also closed
 - **Warning:** Linux-specific

Reading from a FILE*

- There are several functions in STDIO for reading from a FILE*
 - Reading one character at a time: **fgetc** and **ungetc**
 - Reading an entire line: **fgets**
 - Reading binary data: **fread**
 - Reading formatted data: **fscanf**

fgetc & ungetc >>

fgetc & ungetc

- Reading a character
 - `int fgetc (FILE *stream);`
- Returns
 - EOF at end of file, otherwise it returns the read character
- *Pushing back* a character to the bufer
 - `int ungetc(int c, FILE *stream);`
 - Puts back character **c** to stream

fgets >>

Reading a line - *fgets*

```
char *fgets (char *str, int size, FILE *stream);
```

- This function reads up to **one less than *size* bytes** from stream and stores the results in *str* .
- A null character (`'\0'`) is stored in the buffer after the last byte read in.
- Reading stops after an EOF or a newline character is reached.
- If a newline is read, the `\n` is stored in *str* .
- On success, *str* is returned;
- on failure, NULL is returned.

getline function >>

getline function

```
ssize_t getline(char **lineptr, size_t *n, FILE *stream);
```

- POSIX function
- It reads a line (up to "\n") from a stream and...
 - Copies to **lineptr* if **lineptr* is not NULL (and **n* is set to the size of the buffer pointed by **lineptr*). If the buffer pointed by **lineptr* is not large enough it gets reallocated
 - It needs to be an allocated buffer
 - Allocates a buffer and copies the line to this allocated buffer if **lineptr* is NULL. The address of the buffer is returned in **lineptr*
- **n* is returned with the number of bytes written to **lineptr*
- *getline* returns the number of chars read or -1 on ERROR or EOF

getdelim function

```
ssize_t getdelim(char **lineptr, size_t *n,  
                  int delim, FILE *stream);
```

- Similar to *getline*
- Only difference is parameter “delim”
 - It identifies the end-of-line terminator
 - *getline* always uses '\n' as end-of-line terminator
 - *getdelim* allows to specify a different one

- Reading chunks of binary data
 - Example: set of C structures
- ```
size_t fread (void *buf, size_t size, size_t nr, FILE *stream);
```
- Read up to *nr* elements, each of *size* bytes from *stream* into the buffer pointed by *buf*
  - It returns the number of elements read (not the number of bytes!)
    - if the number of elements read is less than *nr*, then...
      - The function might have reached EOF . Test it with `feof()`.
      - There might be an error. Test it with `ferror()`.

- Example

```
char buf[128];
size_t nr;
nr = fread (buf, sizeof(buf), 1, stream);
if (nr == 0) {
 /* error */
}
```

# Reading formatted data

- `fscanf` for reading formatted data
- `int fscanf (FILE *stream, const char *format, ...);`
  - *Returns*
    - *the number of matched an assigned inputs*
    - *EOF on error*
- *There is a family of scanf functions*
  - `scanf`: *reads from stdin*
  - `sscanf`: *reads from a string*

# Writing to a stream

- Writing to a stream
  - Three main operations
    - A character
    - A string
    - A chunk of binary data

Writing a character to a stream >>

# Writing a char to a stream

- `int fputc (int c, FILE *stream);`
  - `fputc()` writes the byte specified by `c` (cast to an unsigned char) to the stream pointed by `stream`.
    - `stream` must be open for writing
  - On success, it returns `c`
  - On error, it returns `EOF`, and `errno` is set appropriately.
- Example

```
if (fputc ('p', stream) == EOF) {
/* error */
...}
```

# Writing a string to a stream (1)

- `int fputs (const char *str, FILE *stream);`
  - `fputs()` writes the `'\0'` terminated string *str* to *stream*.
    - *stream* must be open for writing
  - On error, it returns `EOF`, and `errno` is set appropriately.
- **Example**

```
if (fputs ("texto", stream) == EOF) {
 /* error */
 ...}
```

# Writing a string to a stream (2)

- Another function to write a string to a stream is...

```
int fprintf(FILE *stream, const char *format, ...);
```

- Example

- Remember, `stderr` is a stream!

```
char *str = "this is a test!";
```

```
fprintf(stderr, "%s", str);
```



# Writing binary data

- `size_t fwrite (void *buf, size_t size, size_t nr, FILE *stream);`
  - write to `stream` up to `nr` elements from buffer `buf`
  - each element has `size` bytes
  - `fwrite` returns the number of elements successfully written to `stream`
    - A return value less than `nr` denotes error.

- Binary data are dependent on many issues
  - Differences in variable sizes, alignment, padding, and byte order
  - Binary data written with one application may not be readable by a different application
- Solution
  - Standardized formats
    - Examples
      - GZIP file format: RFC 1952 (<https://tools.ietf.org/html/rfc1952>)
      - PGM file format: <http://netpbm.sourceforge.net/doc/pgm.html>
      - PNG specification - ISO/IEC 15948:2003 (E): <https://www.w3.org/TR/PNG/>
      - Etc.

# Set the position of a stream

- Standard I/O functions to deal with the position of a stream
  - `int fseek (FILE *stream, long offset, int whence);`
    - whence: `SEEK_SET`, `SEEK_CUR`, `SEEK_END`
  - `void rewind (FILE *stream);`
    - Reset the file position
    - same as `fseek(stream, 0, SEEK_SET)`

# Get the position of a stream

- What's the current position of a stream?
- `long ftell (FILE *stream);`
  - Returns the current stream position
  - On error, it returns -1 and `errno` is set appropriately

# Flushing a stream (1)

- Flushing a stream
  - Have the buffer written through the `write()` system call
- Function `int fflush (FILE *stream);`
  - Any unwritten data in `stream` is flushed to the OS
  - If `stream` is `NULL`, all open input streams in the process are flushed
- Returns:
  - 0 on success
  - EOF on error and sets `errno` appropriately

# Flushing a stream (2)

- Flushing a stream
  - The standard I/O buffer is written through the `write()` system call to the OS buffer
  - It does not mean that the content is actually written to file
    - Standard I/O buffers are in user-space
    - OS buffers are in kernel space
- Using standard I/O functions, we minimize the number of (costly) system calls
  - A system call (e.g., *write*) is issued only when the disk or some other storage has to be accessed.

# Flushing a stream (3)

- To provide for effective write to persistent storage, the programmer needs to...
  - *fflush* the stream
  - *fsync* the associated file descriptor
    - The file descriptor of a stream can be obtained with:
      - `int fileno (FILE *stream);`

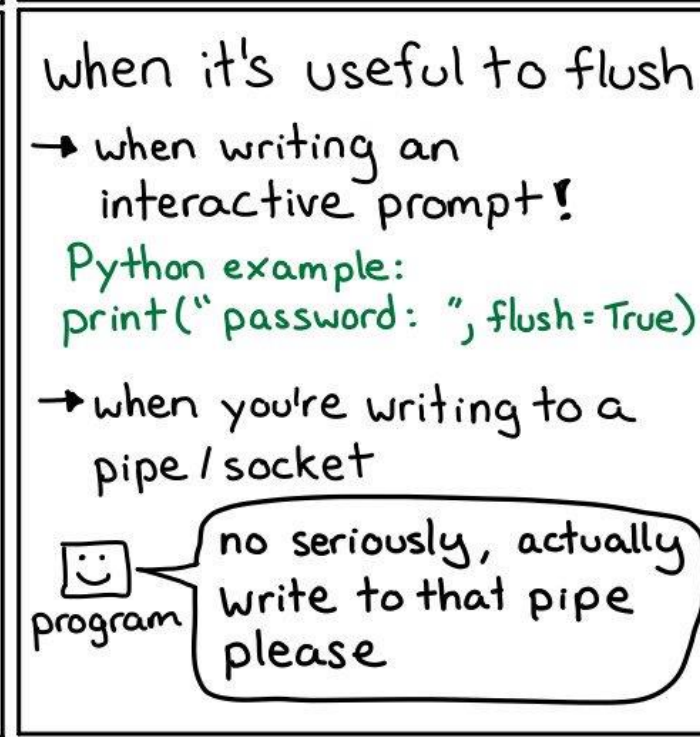
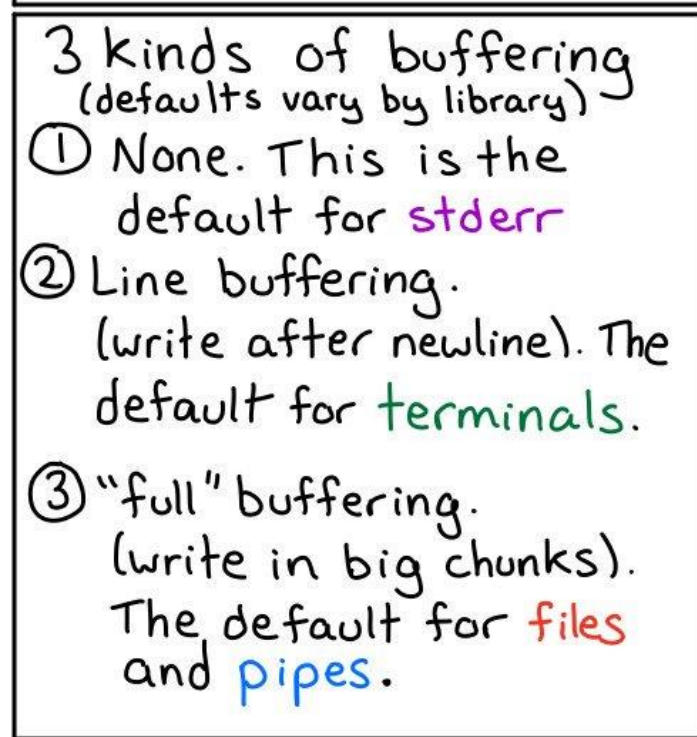
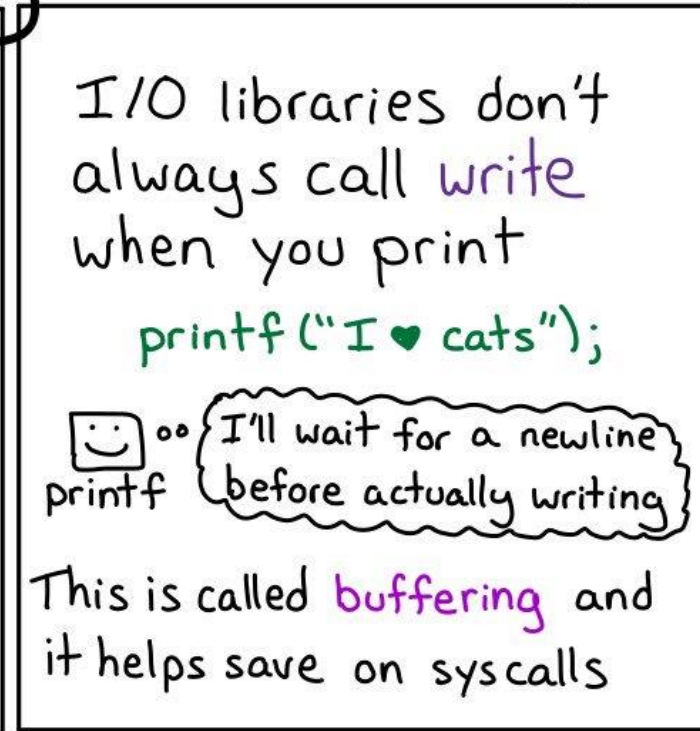

- Function `ferror` tests if the error indicator is set on *stream*
  - `int ferror (FILE *stream);`
  - Returns non-zero on error, 0 otherwise
- Function `feof` tests if file position of stream is at end-of-file
  - `int feof (FILE *stream);`
- Function `clearerr` clears error status of *stream*
  - `void clearerr (FILE *stream);`



- Types of buffering
  - *Unbuffered*
    - No user buffer is performed. Data are sent directly to the kernel. STDERR is, by default, unbuffered. Otherwise, this mode is rarely used.
  - *Line-buffered*
    - With each newline (`\n`), the buffer is submitted to the kernel. This is the default for stdout.
  - *Block-buffered or full-buffering*
    - A block is a fixed number of bytes.

# Controlling buffering (2)

- Function `setvbuf`
  - To be called right after opening `stream`
  - `int setvbuf (FILE *stream, char *buf, int mode, size_t size);`
- Mode can be
  - `_IONBF`: unbuffered
  - `_IOLBF`: line-buffered
  - `_IOFBF`: Block-buffered (*full-buffered*)
- `buf` may point to a buffer of `size` bytes that standard I/O will use as the buffer for `stream`.
- If `buf` is `NULL`, a buffer of `size` bytes is allocated automatically



# remove function

```
int remove(const char *pathname);
```

- Remove is a function available in stdio
- It removes an entry in the filesystem
  - file
  - directory
- In practice, it calls the appropriate system call
  - **unlink** for a file
  - **rmdir** for a directory

- “*Buffered I/O*”, Chapter 3 - *Linux System Programming*, Robert Love, 2<sup>nd</sup> Edition, O’Reilly, 2013
- Man pages for standard I/O functions
  - man 3 printf
  - man 3 fopen
  - ...