

Signals

(POSIX systems)

Patrício R. Domingues

Departamento de Eng Informática

ESTG/IPLeiria

Pointer to functions (1)

- The C language has the concept of *pointer to a function*
 - The pointer holds the address of a function
 - The function can be called through the pointer
 - Call syntax is similar to a regular function call
- What's the datatype of a pointer to a function?
 - Datatype is function
 - A function has a signature defined by the following elements
 - Datatype of the return value
 - Datatype of the parameters (if any)

Pointer to functions (2)

- Examples of signatures of functions
 - `int F1(int a, int b);`
 - `int F2(int y, int z);`
 - F1 e F2 have the same signature (the name of the parameters is irrelevant)
 - `double F3 (int a, int b);`
 - F3 has a signature different from F1 (and obviously from F2)
 - `int *F4(int a, int b);`
 - F4 has a signature diferente from F1, F2 and F3

Pointer to functions (3)

- Declaration of a pointer to a function
 - Pointer points to the signature
- Example
 - `int (*PtrF1)(int , int);`
 - PtrF1 can point to functions with signature “**int ... (int, int);**”
 - `double (*PtrF2)(double ,char *);`
 - PtrF2 can point to functions with signature “**double ... (double,char*);**”

Pointer to functions (4)

- How to get the address of a function?
 - simple
 - The name of the function holds the address of the function
 - Thus, the pointer only needs to be assigned the name of the function
 - This can also be done through “&FunctionName”

- Example

```
int F1(int a, int b){  
    return a + b;  
}
```

```
int (*PtrF1)(int , int) = NULL;  /*PtrF1 declaration*/
```

```
int Result;
```

```
PtrF1 = F1;           /* PtrF1 points to the F1 function */
```

```
Result = PtrF1(10,30);  /* Call of F1 through PtrF1 */
```

```
Result = (*PtrF1)(10,30); /* Same as previous line of code */
```

What's a *signal*? (1)

- Signal is an asynchronous notification delivered to a process

- Notification

- Related to an event


- Example:

- » control + C within the shell has the effect of delivering a signal (SIGINT) to the running process
 - » Process tries to access an invalid memory address. SO delivers the SIGSEGV (segmentation violation) to the process.

- Asynchronous

- The signal can happen at any time

- Therefore, the notification can be delivered at any time...



<https://bit.ly/2wLTTca>

What's a *signal*? (2)

- In UNIX, a signal is represented by an integer number
- Each signal has also a symbolic name
 - Easier for us humans to remember
 - Examples: SIGINT, SIGHUP, SIGSEGV, SIGPIPE, etc.
- There are around 40 different signals
 - man 7 signal
 - kill -l signal
 - List the available signals

signals

drawings.juns.ca

If you've ever used
⚡ kill ⚡
you've used signals



the Linux kernel sends
your process signals in
lots of situations



you can send signals
yourself with the **kill**
system call or command

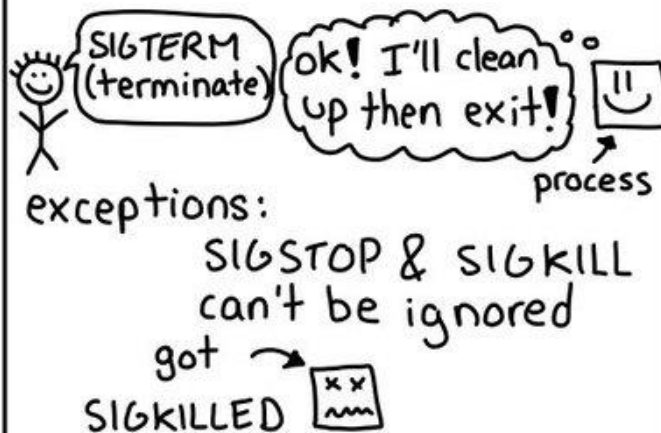
SIGINT Ctrl-C } various
SIGTERM kill } levels of
SIGKILL kill -9 } "die"

SIGHUP kill -HUP
↑
often interpreted as
"reload config", eg by nginx

Every signal has a default
action, one of:

- ☺ ignore
- ☒ kill process
- ☒ ☹ kill process AND
make core dump file
- ⏸ stop process
- ⏪ resume process

Your program can set
custom handlers for
almost any signal



exceptions:

SIGSTOP & SIGKILL
can't be ignored

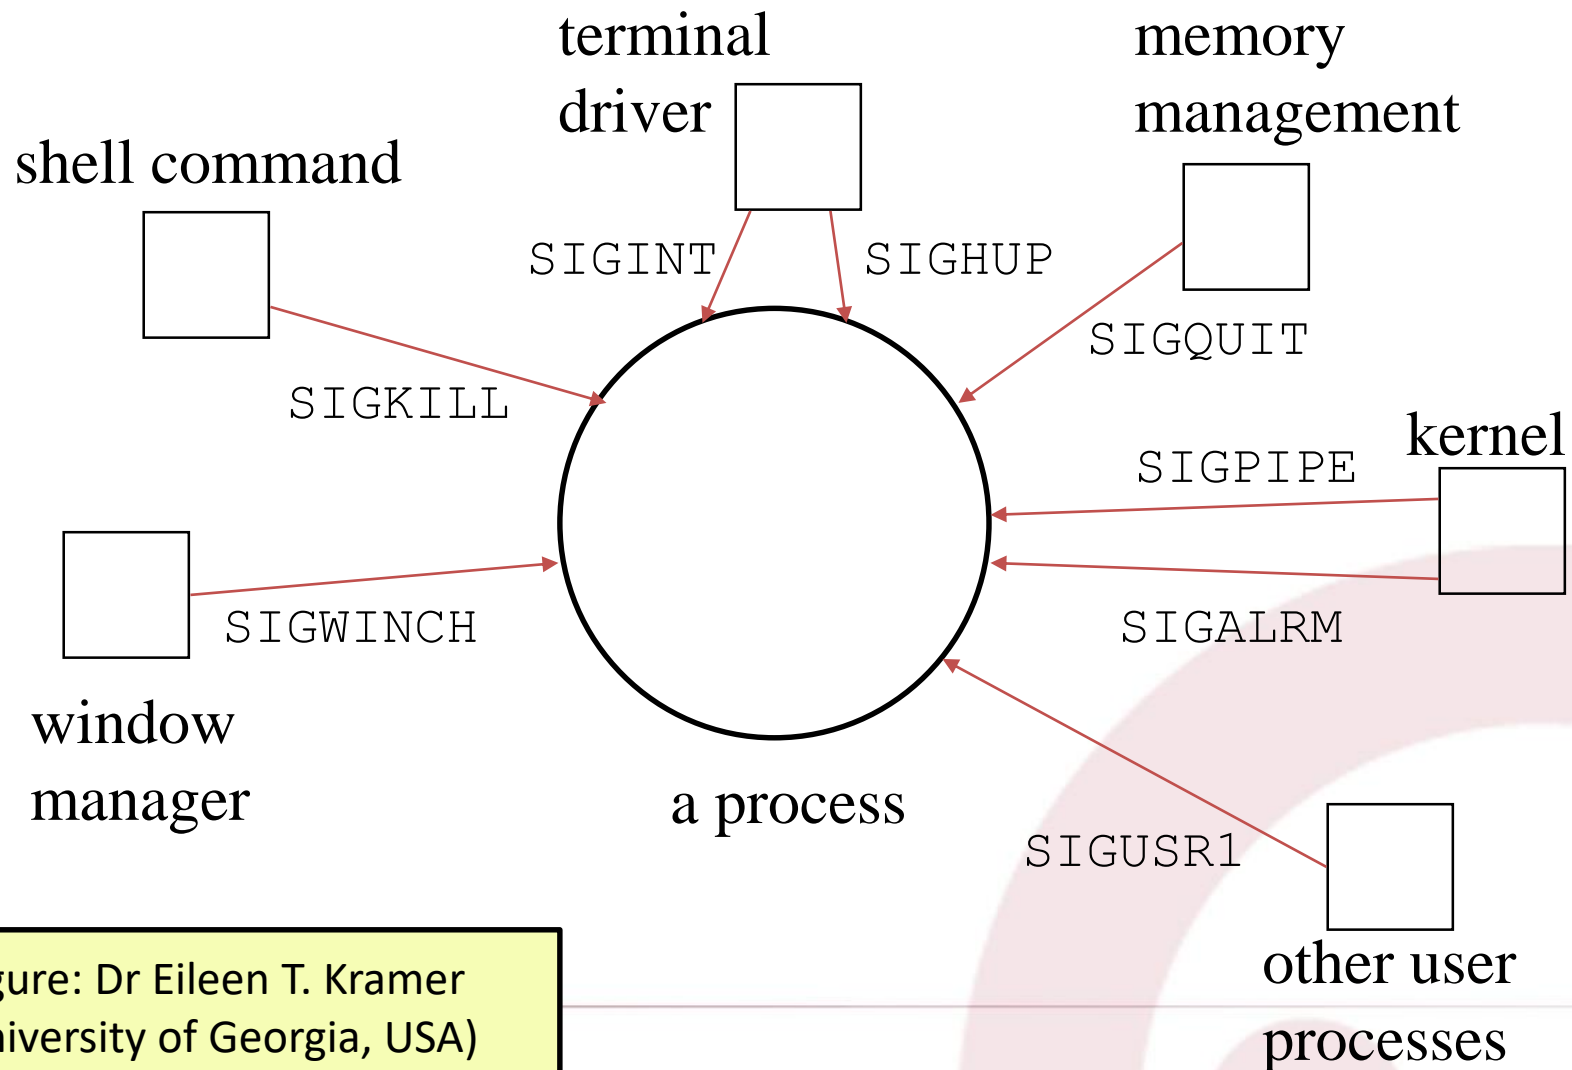
got
SIGKILLED ☒

signals can be hard
to handle correctly since
they can happen at
ANY time



Source of signals

- Which elements can send a signal to a process?



The *kill* command (1)

- kill command (tries to) deliver a signal to a destination process
 - Syntax
 - `kill [options]PID1 PID2 PID3...`
 - Send *signal* to processes with PID1, PID2, PID3, ...
 - Why the name *kill*?
 - By default, when a process receives a signal, it terminates
 - A process can be configured to have a different behavior
 - *Handling* the signal
 - More on the command kill
 - `man kill`

The *kill* command (2)

- Examples

- `kill -SIGINT 1234`

- Send the signal SIGINT to process with PID 1234

- `kill -9 1234`

- Send the signal 9 (**SIGKILL**) to processes 1234 777

- `kill -SIGKILL 1234`

- Same as above

- `kill -kill $$`

- Kills the shell
 - \$\$ is the PID of the shell

SIGKILL and SIGSTOP are the only signal that cannot be captured (more on this later)

Sending a signal

- A regular user can only send a signal to processes that she/he owns
 - `kill -KILL 1`
`bash: kill: (1) - Operation not permitted`
 - Process with PID 1 is a system process (process *init*)
- An administrator can send a signal to any process

The *killall* command

- Killall is a variant of the kill command
- `killall [options] name`
 - It sends the signal to all processes whose name is *name*
- `killall -INT bash`
 - *(tries) to send the SIGINT signal to all processes whose name is bash*
 - For a regular user, only the bash processes that belong to the user receive the signal

Kill

JULIA EVANS
@bork

kill doesn't just kill programs



you can send ANY signal to a program with kill!

kill - SIGNAL PID

↑
name or number

which signal kill sends

	<u>name</u>	num
kill	=> SIGTERM	15
kill -q	=> SIGKILL	9
kill -KILL		
kill -HUP	=> SIGHUP	1
kill -STOP	=> SIGSTOP	19

kill -l
lists all signals

1 HUP	2 INT	3 QUIT	4 ILL
5 TRAP	6 ABRT	7 BUS	8 FPE
9 KILL	10 USR1	11 SEGV	12 USR2
13 PIPE	14 ALRM	15 TERM	16 STKFLT
17 CHLD	18 CONT	19 STOP	20 TSTP
21 TTIN	22 TTOU	23 URG	24 XCPU
25 XFSZ	26 VTALRM	27 PROF	28 WINCH
29 POLL	30 PWR	31 SYS	

killall - SIGNAL NAME

signals all processes called NAME
for example

\$ killall firefox

useful flags:

-w wait for all signaled processes to die

-i ask before signalling

pgrep

prints PIDs of matching running programs

pgrep fire matches firefox
fire bird
NOT bash firefox.sh

To search the whole command line (eg bash firefox.sh)

use **pgrep -f**

pkill

same as pgrep, but signals PIDs found. ex:

pkill -q -f firefox



I use pkill more than killall these days

The kill() function (#1)

- How do we send a signal programmatically?
 - kill function
 - `int kill(pid_t pid, int sig);`
 - return -1 on error (setting *errno*), 0 on success
 - Interpretation of pid
 - *pid* > 0: send signal to process with PID *pid*
 - *pid* == 0: send signal to all processes whose group ID equals the sender
 - » Parent process sending a signal to all children
 - Interpretation of sig
 - Sig > 0: send signal sig
 - Sig==0: do not actually send a signal, but acts if has done so
 - » Use-case: checking if a given process exist

The kill() function (#2)

- `man 2 kill`
 - Need to specify the section 2 of the manual
 - Section 2 of man describes system calls
 - Kill, read, write, open...
 - `man 2 intro`

INTRO(2)

Linux Programmer's Manual

INTRO(2)

NAME

`intro` - introduction to system calls

DESCRIPTION

Section 2 of the manual describes the Linux system calls. A system call is an entry point into the Linux kernel. Usually, system calls are not invoked directly: instead, most system calls have corresponding C library wrapper functions which perform the steps required (e.g., trapping to kernel mode) in order to invoke the system call. Thus, making a system call looks the same as invoking a normal library function.

For a list of the Linux system calls, see `syscalls(2)`.

Signals sent by functions

- Some system functions send a signal
 - abort() sends the SIGABRT signal to the calling process
 - `void abort(void);`
 - alarm(): schedule the delivery of the SIGALRM signal to the calling process
 - Appropriate for setting timeouts
 - `unsigned int alarm(unsigned int seconds);`

Responding to a *signal*

- For given signal, a process can be configured to...
 - Execute the default action
 - The default action for many (not all!) signals is to terminate the process
 - SIGKILL and SIGSTOP always results in their default actions
 - SIGKILL: terminates the process
 - SIGSTOP: stops the process (same as ctrl+Z on the shell)
 - Ignore the signal
 - Not available for SIGKILL and SIGSTOP
 - Launch a signal *handler*
 - *Signal handler*: function called whenever a signal is received

Setting a *signal handler*

- The `sigaction` library call
 - Installs a signal handler for a given signal

```
#include <signal.h>
```

```
int sigaction(int signum, const struct  
    sigaction *act, struct sigaction *oldact);
```

- `signum`: signal number to be configured
- `act`: pointer to `sigaction` struct that contains the configuration to use
- `oldact`: `sigaction` struct filled with the previous configuration
- Question
 - Why const `struct sigaction *act` **VS.** `struct sigaction *oldact`?

struct sigaction (#1)

- The *struct sigaction*

```
struct sigaction {  
    void (*sa_handler)(int);  
    void (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer)(void);    /* << obsolete */  
};
```

sa_handler:

- pointer to the function that will handle the signal. The function should have one integer parameter and does not return anything
 - The integer parameter corresponds to the signal number that triggered the handler
- SIG_DFL to restore the default behavior
- SIG_IGN to set the process to ignore the signal

struct sigaction (#2)

- The *struct sigaction*

```
struct sigaction {  
    void (*sa_handler)(int);  
    void (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer)(void);    /* << obsolete */  
};
```

sa_sigaction:

- If *sa_flags* is *SA_SIGINFO*, *sa_sigaction* specifies the signal handling function
- This has improved functionalities for the signal handler function, namely a *siginfo_t* structure (see next slide)
- *sa_sigaction* is not compatible with *sa_handler*. Use one or the other, but not both!

- The *struct siginfo_t*

```
struct siginfo_t {  
    int si_signo;      /* Signal number */  
    int si_errno;      /* An errno value */  
    int si_code;       /* Signal code */  
    pid_t si_pid;      /* Sending process ID */  
    uid_t si_uid;      /* Real user ID of sending process */  
    int   si_status;    /* Exit value or signal */  
    clock_t si_utime;   /* User time consumed */  
    clock_t si_stime;   /* System time consumed */  
    (...)  
}
```

- The `struct siginfo_t` returns a large number of data regarding the signal and the context of the call

struct sigaction (#4)

- The *struct sigaction*

```
struct sigaction {  
    void (*sa_handler)(int);  
    void (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer)(void);    /* << obsolete */  
};
```

sa_mask:

- Set of signal to be blocked during the call of a signal handler
 - Avoid race conditions
 - During the handler execution, the signal being processed by the handler is blocked by default
 - Regarding `sigset_t`, see **man 3 sigsetops**

- The *struct sigaction*

```
struct sigaction {  
    void (*sa_handler)(int);  
    void (*sa_sigaction)(int, siginfo_t *, void *);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_restorer)(void);    /* << obsolete */  
};
```

sa_flags:

- specifies a set of flags which modify the behavior of the signal.
- It can comprise several flags (through bitwise OR)
- Example
 - SA_NOCLDSTOP | SA_ONSTACK | ...
 - SA_SIGINFO to use the field `sa_sigaction` as handler
 - | is the bitwise OR operator
 - See **man 3 sigaction**

Example with *sigaction*

```
int main(void)    {
    struct sigaction act;
    act.sa_handler = process_signal;
    sigemptyset( &act.sa_mask );
    act.sa_flags = 0;
    sigaction( SIGINT, &act, 0 );
    while(1)
    {
        printf("waiting one second - press CTRL+C :)\n");
        sleep(1);
    }
    return 0;
}

void process_signal(int signum){
    printf("Capturing %d (SIGINT=%d)\n",signum, SIGINT);
}
```

The program will not terminate with
CTRL+C (SIGINT)
However, it will terminate with CTRL+\
(SIGQUIT)

- SIGINT can be send to a foreground process with
CTRL+C

Example with `siginfo_t` (1)

```
/* Setting sigaction() to use siginfo_t. */
static void hdl(int sig, siginfo_t *siginfo, void *context) {
    printf ("Sending PID: %ld, UID: %ld\n",
        (long) siginfo->si_pid, (long) siginfo->si_uid);
}

int main (int argc, char *argv[]) {
    struct sigaction act;
    memset (&act, '\0', sizeof(act));
    /* Use sa_sigaction field: handle has
two additional parameters */
    act.sa_sigaction = &hdl;
```


Example with `siginfo_t` (2)

```
/* SA_SIGINFO flag tells sigaction() to use the sa_sigaction  
field, not sa_handler. */  
  
act.sa_flags = SA_SIGINFO;  
  
if (sigaction(SIGTERM, &act, NULL) < 0) {  
    perror ("sigaction");  
    return 1;  
}  
  
while (1) {  
    sleep(10);  
}  
  
return 0;  
}
```

- Limitations
 - A signal can be received at anytime
 - Whenever the signal is received, the process interrupts its regular execution and *jumps* to the signal handler
 - The signal handler has no context except for global variables
 - This is one of the few use cases of global variables
 - In the SA_SIGINFO mode of sigaction, the third parameter of the signal handler is a buffer address given by the user.

Sets of signal (`sigset_t`)

- Sets of signals are...
 - used by functions to define which signal types are to be processed
 - represented by the `sigset_t` datatype
- There is several functions for creating, querying and changing sets of signals

Functions to deal with `sigset_t` >>



Functions handling `sigset_t`

```
#include <signal.h>
```

```
int sigemptyset(sigset_t *set );
```

```
int sigfillset(sigset_t *set );
```

```
int sigaddset(sigset_t *set, int signo);
```

```
int sigdelset(sigset_t *set, int signo);
```

```
int sigismember(const sigset_t *set, int  
signo);
```



Function `sigprocmask()` (1)

- The function `sigprocmask()` is used to fetch and/or change the signal mask of the calling thread
- The signal mask is the set of signals whose delivery is currently blocked for the caller
- **`#include <signal.h>`**
**`int sigprocmask(int how,`
**`const sigset_t *set,`
`sigset_t *oldset);`****
- **how** – indicates how mask is modified

- **#include <signal.h>**
int sigprocmask(int how,
const sigset_t *set,
sigset_t *oldset);
- **how** – indicates how mask is modified
 - **SIG_BLOCK**: set of blocked signals is the union of the current set and the **set** argument.
 - **SIG_UNBLOCK**: the signals in **set** are removed from the current set of blocked signals.
 - **SIG_SETMASK**: set of blocked signals is set to the argument **set**.

System calls and signals

- When a system call (e.g. `read()`) is interrupted by a signal...
 - 1) The signal handler is called
 - 2) The signal handler terminates and thus returns the control back to the system call
 - On UNIX, slow system calls do not resume.
 - Whenever a signal is received within a system call, the system call returns an error and sets `errno` to `EINTR`
- What is a *slow system call*?
 - System calls that perform I/O operations on devices that can block the caller *forever*
 - sockets (networks), pipes
 - The `pause()` system call
 - Blocks the process until it receives a signal
- Use `SA_RESTART` to recover slow system calls automatically

Signals in Linux (#1)

- Source: “*Signals*”, Chapter 10 - *Linux System Programming*, Robert Love, 2nd Edition, O’Reilly, 2013

Signal	Description	Default action
SIGABRT	Sent by <code>abort()</code>	Terminate with core dump
SIGALRM	Sent by <code>alarm()</code>	Terminate
SIGBUS	Hardware or alignment error	Terminate with core dump
SIGCHLD	Child has terminated	Ignored
SIGCONT	Process has continued after being stopped	Ignored
SIGFPE	Arithmetic exception	Terminate with core dump
SIGHUP	Process’s controlling terminal was closed (most frequently, the user logged out)	Terminate
SIGILL	Process tried to execute an illegal instruction	Terminate with core dump
SIGINT	User generated the interrupt character (Ctrl-C)	Terminate
SIGIO	Asynchronous I/O event	Terminate ^a
SIGKILL	Uncatchable process termination	Terminate
SIGPIPE	Process wrote to a pipe but there are no readers	Terminate
SIGPROF	Profiling timer expired	Terminate

Signals in Linux (#2)

- Source: “*Signals*”, Chapter 10 - *Linux System Programming*, Robert Love, 2nd Edition, O’Reilly, 2013

Signal	Description	Default action
SIGSEGV	Memory access violation	Terminate with core dump
SIGSTKFLT	Coprocessor stack fault	Terminate ^b
SIGSTOP	Suspends execution of the process	Stop
SIGSYS	Process tried to execute an invalid system call	Terminate with core dump
SIGTERM	Catchable process termination	Terminate
SIGTRAP	Break point encountered	Terminate with core dump
SIGTSTP	User generated the suspend character (Ctrl-Z)	Stop
SIGTTIN	Background process read from controlling terminal	Stop
SIGTTOU	Background process wrote to controlling terminal	Stop
SIGURG	Urgent I/O pending	Ignored
SIGUSR1	Process-defined signal	Terminate
SIGUSR2	Process-defined signal	Terminate
SIGVTALRM	Generated by <code>setitimer()</code> when called with the <code>ITIMER_VIRTUAL</code> flag	Terminate
SIGWINCH	Size of controlling terminal window changed	Ignored
SIGXCPU	Processor resource limits were exceeded	Terminate with core dump
SIGXFSZ	File resource limits were exceeded	Terminate with core dump

^a The behavior on other Unix systems, such as BSD, is to ignore this signal.

^b The Linux kernel no longer generates this signal; it remains only for backward compatibility.

Signal name in Linux


- The system-defined `sys_siglist[...]` vector of strings holds names of each signal
 - `extern const char *const sys_siglist[];`
- The name of signal is also available with:
 - `char *strsignal(int sig);`
- Example:

```
#include <signal.h>
for(i=0;i<20;i++) {
    printf("signal %d => '%s'\n",i,sys_siglist[i]);
}
```



IPL

escola superior de tecnologia e gestão
instituto politécnico de leiria



- “*Signals*”, Chapter 10 - *Linux System Programming*, Robert Love, 2nd Edition, O’Reilly, 2013
- Sigaction
 - <http://www.linuxprogrammingblog.com/code-examples/sigaction>