



Sistemas Operativos fork, exec & exit



Patrício Domingues
ESTG/IPLeiria

May 2019



“top” / linux

142 processes

uptime

Average load of the system

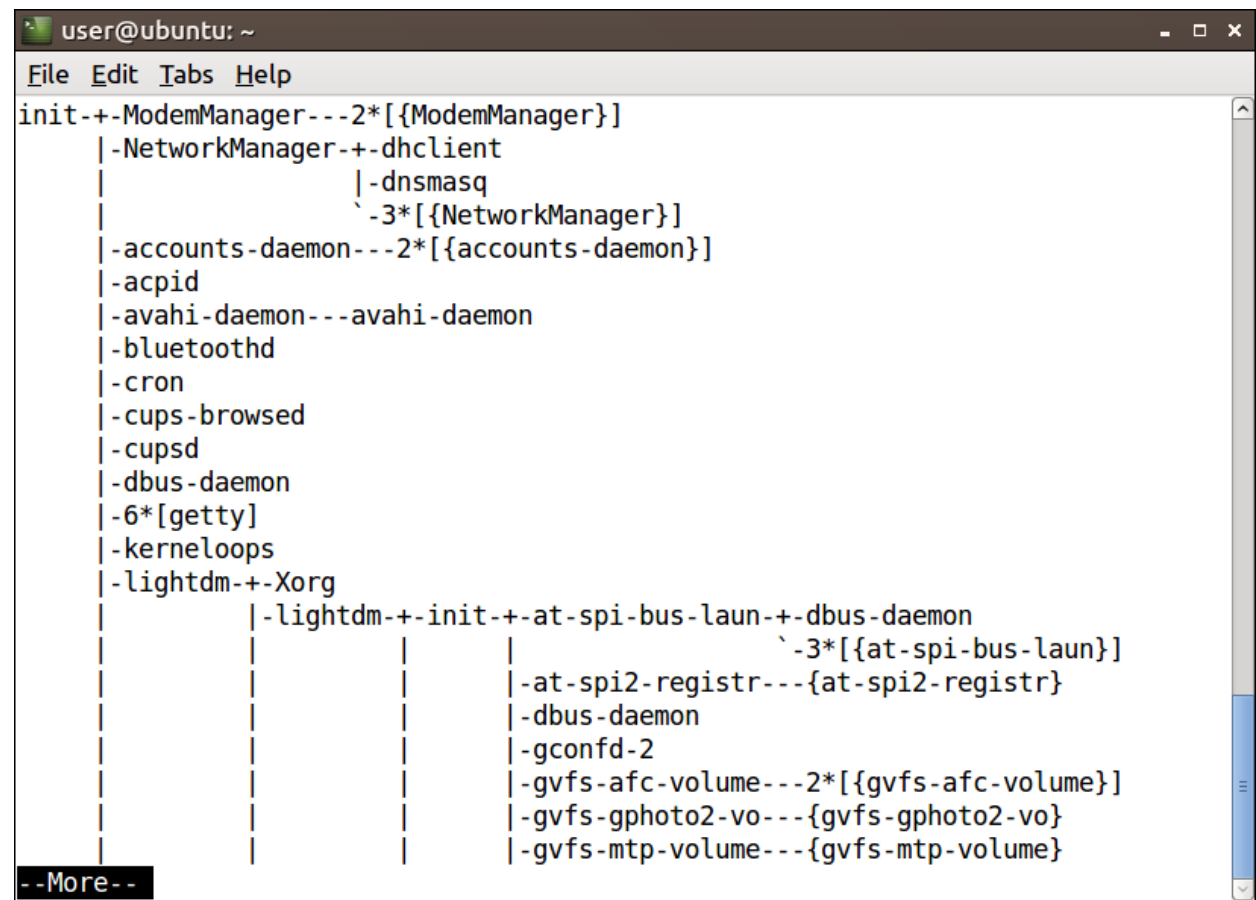
```
user@ubuntu: /bin
File Edit Tabs Help
top - 15:50:52 up 1 day, 9:44, 11 users, load average: 0.70, 0.45, 0.50
Tasks: 142 total, 1 running, 141 sleeping, 0 stopped, 0 zombie
%Cpu(s): 3.5 us, 0.7 sy, 0.0 ni, 95.8 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 1026012 total, 916316 used, 109696 free, 121464 buffers
KiB Swap: 0 total, 0 used, 0 free. 518944 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1450	root	20	0	216612	41560	10060	S	3.9	4.1	6:23.45	Xorg
2531	user	20	0	213720	24964	18468	S	0.3	2.4	5:33.43	vmtoolsd
2518	user	20	0	220292	21220	13288	S	0.0	2.1	1:20.05	pcmanfm
2510	user	20	0	208644	18860	11116	S	0.0	1.8	0:42.46	lxpanel
2463	user	20	0	123992	16992	9664	S	0.3	1.7	0:15.91	ibus-ui-gtk3
7731	user	20	0	187512	15712	10864	S	1.6	1.5	0:14.52	lxterminal
2505	user	20	0	34120	11848	7240	S	0.0	1.2	0:26.95	openbox
2441	user	20	0	48472	10488	5148	S	0.0	1.0	0:02.08	lxsession
1247	root	20	0	52116	9796	4676	S	0.0	1.0	0:23.12	NetworkManager
2429	user	20	0	45764	8640	2896	S	0.3	0.8	1:54.88	ibus-daemon
9728	user	20	0	49172	8580	5668	S	0.0	0.8	0:00.18	vim
9750	user	20	0	49172	8580	5668	S	0.0	0.8	0:00.18	vim
9788	user	20	0	49172	8580	5668	S	0.0	0.8	0:00.17	vim
9821	user	20	0	49172	8580	5668	S	0.0	0.8	0:00.16	vim
9840	user	20	0	49172	8580	5668	S	0.0	0.8	0:00.27	vim
9864	user	20	0	49172	8580	5668	S	0.0	0.8	0:00.18	vim
9730	user	20	0	49172	8576	5668	S	0.0	0.8	0:00.14	vim
9769	user	20	0	49172	8576	5668	S	0.0	0.8	0:00.15	vim
9802	user	20	0	49172	8576	5668	S	0.0	0.8	0:00.15	vim

Processes (1)

✓ Processes in Unix

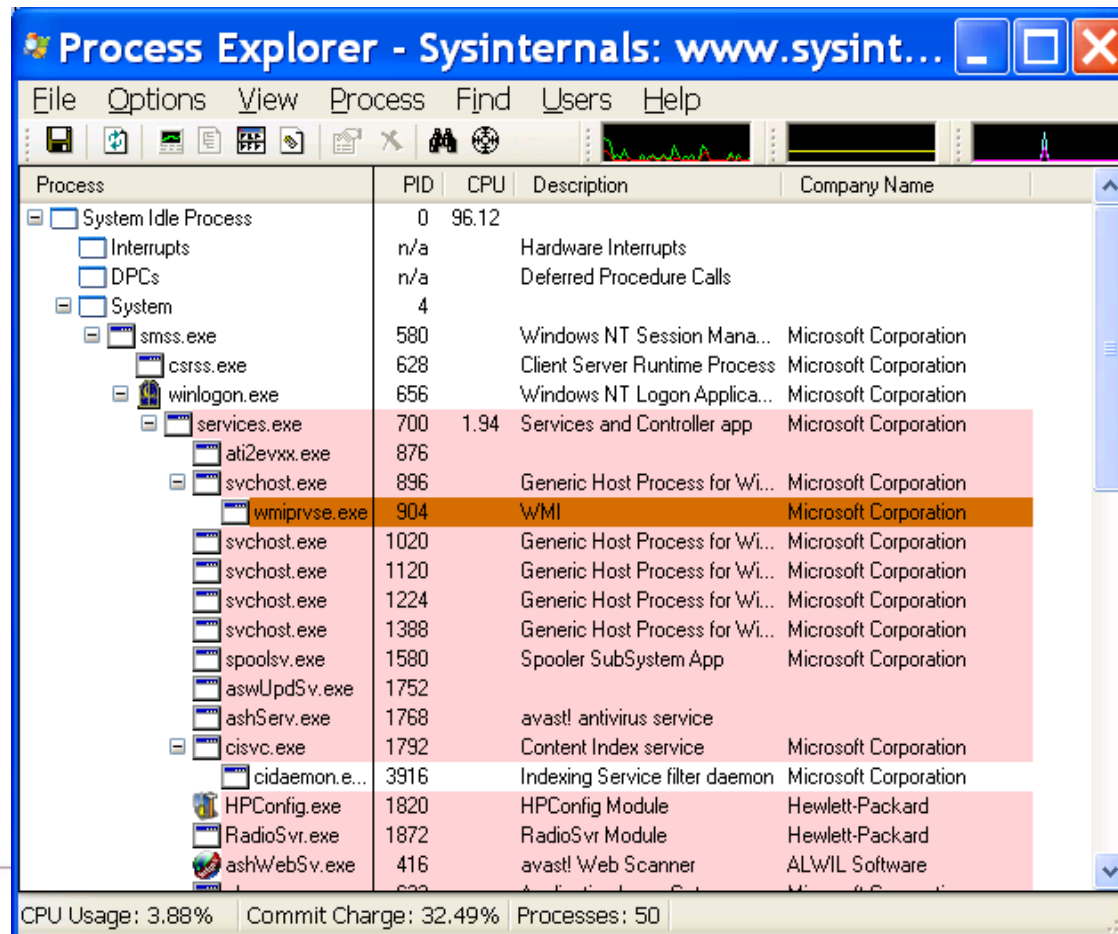
- All are descendant from the *init* (PID=1) process
- A process is created...by another process
 - Parent / son relationship
 - Tree of processes



```
user@ubuntu: ~
File Edit Tabs Help
init--ModemManager---2*[{ModemManager}]
|-NetworkManager--dhclient
|                  |-dnsmasq
|                  \-3*[{NetworkManager}]
|-accounts-daemon---2*[{accounts-daemon}]
|-acpid
|-avahi-daemon---avahi-daemon
|-bluetoothd
|-cron
|-cups-browsed
|-cupsd
|-dbus-daemon
|-6*[getty]
|-kerneloops
|-lightdm--Xorg
|          |-lightdm--init--at-spi-bus-laun--dbus-daemon
|          |                               \-3*[{at-spi-bus-laun}]
|          |   |-at-spi2-registr---{at-spi2-registr}
|          |   |-dbus-daemon
|          |   |-gconfd-2
|          |   |-gvfs-afc-volume---2*[{gvfs-afc-volume}]
|          |   |-gvfs-gphoto2-vo---{gvfs-gphoto2-vo}
|          |   \-gvfs-mtp-volume---{gvfs-mtp-volume}
--More--
```


Processes (2)

- ✓ There is also a process hierarchy in Microsoft Windows
 - “process explorer” (sysinternals.com)



Process model in UNIX

- `fork()`
 - system call to create a process
 - It is called by the parent process
- The son process inherits all characteristics of the parent process
 - Variables, program counter, open files, allocated memory, etc.
 - The son is a snapshot of the parent
- After “fork”, each process executes separately
 - The change of a variable in one process does **not** reflect on the other one



- ✓ `pid_t fork(void);`
- ✓ The fork system call returns an integer:
 - 0 to the newly created son process
 - > 0 to the calling parent process
 - The return value corresponds to the PID of the newly created process
- ✓ It can also returns -1 if an error has occurred

Example – fork system call

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
#include <sys/types.h>
```

```
int main()
{
    pid_t id;
```

```
    id = fork(); /* returns 0: son process; > 0 to the parent */
    if (id == 0)
```

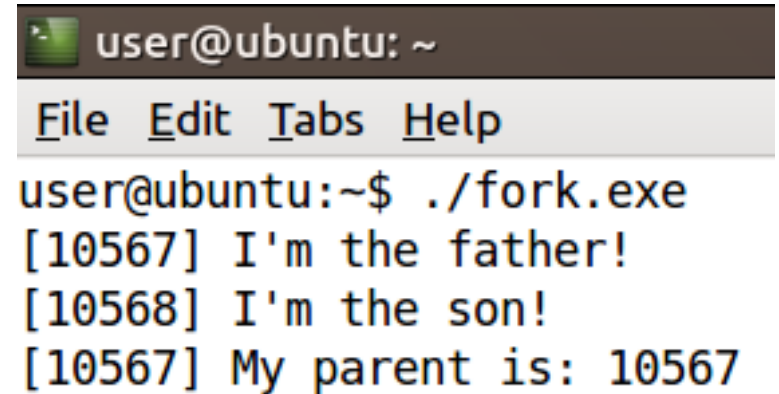
```
    { /* code only executed by the son process */
        printf("[%d] I'm the son!\n", getpid());
        printf("[%d] My parent is: %d\n", getpid(), getppid());
    }
```

```
    else if (id > 0 )
```

```
    { /* code only executed by the parent process */
        printf("[%d] I'm the father!\n", getpid());
        wait(NULL);
    }
```

```
    return 0;
```

```
}
```



```
user@ubuntu: ~
File Edit Tabs Help
user@ubuntu:~$ ./fork.exe
[10567] I'm the father!
[10568] I'm the son!
[10567] My parent is: 10567
```


Example – `fork` in a for loop

- ✓ How many processes are created by the following code?

```
#include <...>
int main(void){
    int i;
    for(i=0;i<3;i++){
        if( fork() == 0 ){
            printf("PID=%u\n", getpid());
            fflush(stdout);
        }
    }
    return 0;
}
```

Only newly created processes print their PID

Answer: 7


$2^n - 1$, with $n=3$

- ✓ Process ID (PID)
 - Integer identifier of a process
- ✓ The Linux kernel allocates process IDs to processes in a strictly linear fashion.
 - If pid 37 is the highest number currently allocated, pid 38 will be allocated next, even if the process last
- ✓ For compatibility with old UNIX, the max value for PID is 32768 (16-bit signed int)
- ✓ This value can be changed
 - `/proc/sys/kernel/pid_max`

Process ID (PID) - #2

- ✓ Within a C program, the PID of the calling process is returned with `getpid()`
 - `pid_t getpid(void);`
- ✓ The PID of the parent process is available through `getppid()`
 - `pid_t getppid(void);`

```
printf ("My pid=%jd\n", (intmax_t) getpid ());  
printf ("Parent's pid=%jd\n", (intmax_t) getppid());
```



printf format specifiers (2)

✓ Source: http://www.pixelbeat.org/programming/gcc/int_types/

```
uint32_t uint32=0xffffffff;
uintmax_t uintmax=UINTMAX_MAX;
off_t offset=TYPE_MAX(off_t); /* Depends on _FILE_OFFSET_BITS */
time_t time=TYPE_MAX(time_t); /* May be float! */
size_t size=TYPE_MAX(size_t); /* Depends on int size */

printf("native int bits %20zu %16x\n"
      "native long bits%20zu %16lx\n"
      "uint32_t max      %20\"PRIu32\" %16\"PRIx32\"\n"
      "uintmax_t max      %20ju %16jx\n" /* try PRIuMAX if %ju unsupported */
      "off_t max         %20jd %16jx\n" /* try PRIdMAX if %jd unsupported */
      "time_t max        %20jd %16jx\n"
      "size_t max         %20zu %16zx\n",
      sizeof(int)*CHAR_BIT, UINT_MAX,
      sizeof(long)*CHAR_BIT, ULONG_MAX,
      uint32, uint32,
      uintmax, uintmax,
      (intmax_t)offset, (intmax_t)offset,
      (intmax_t)time, (intmax_t)time,
      size, size);
```

✓ But...

- If all new processes execute the code of their parents, how can new applications be run?
 - The fork system call creates a clone of the parent process

✓ How do we run an application?

- vim, ps, ls, find, firefox,...

✓ Answer

- The “exec” family of system calls
 - These syscalls replace the image of the calling process

The *exec* family of system calls

✓ “exec” system calls

```
int execl(const char *path, const char *arg, ...);  
int execlp(const char *file, const char *arg, ...);  
int execl_e(const char *path, const char *arg, ..., char *const envp[]);  
int execv(const char *path, char *const argv[]);  
int execvp(const char *file, char *const argv[]);
```

✓ Usage of an exec system call is

- exec...(application_to_be_run)

✓ “exec”

- Replaces the image of the current process by another one from a given executable
 - Functions with “p” are “PATH”-aware
 - Functions with “v” get their parameters from a vector of strings
 - Functions with “l” get their parameters from a list, where itens are separated by “,” and the list ends with NULL

✓ Example: **execl("/bin/ps", "ps", "aux",NULL);**

Example - running “ls” (1)

- ✓ Running “ls -a” resorting to “execvp”
- ✓ Question
 - ✓ Why the “This cannot happen!” message?

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
int main(){
    if (execvp("ls", "ls", "-a", NULL) == -1)
        perror("Error executing ls: ");
    else
        printf("This cannot happen!\n");
    return 0;
}
```

Example – running “ls” (2)

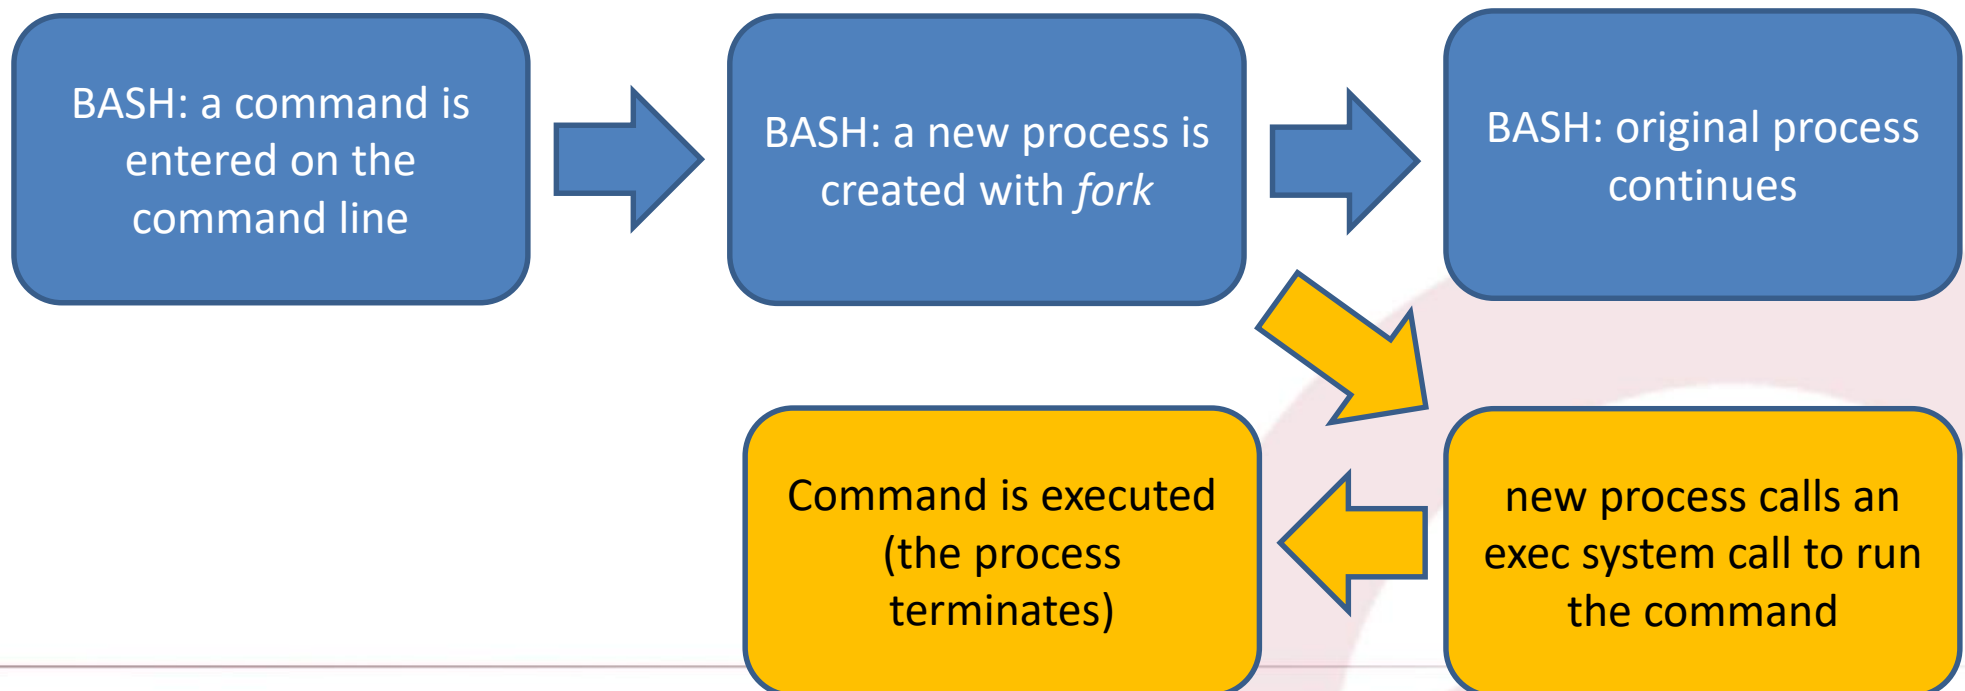
- ✓ Process launches “ls” via `execvp`
 - The exec system calls **NEVER** return when the execution is successful
 - The calling process image is replaced by the image of the executable called via “exec”
 - The calling process runs the executable
 - “ls” in our example
 - Therefore, *printf(“This cannot happen!”)* is removed from memory (as well as all the code of the calling process)
 - The code is replaced by the code of “ls -a”

Executing a command

✓ Executing a command

– Example with bash

- Applies to other shells (sh, zsh, etc.)
- 1st – fork
- 2nd – system call from the exec family



✓ wait and waitpid

```
pid_t wait(int *status);
```

```
pid_t waitpid(pid_t pid, int *status, int options);
```

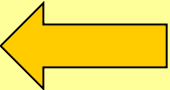
- system calls used to synchronize a parent process with its children
- Wait for state changes on their children processes
 - Child is stopped (SIGSTOP) or terminates
 - Child is resumed by a signal (SIGCONT)
- Example
 - `wait(&status);` -- waits until a children process terminates
 - `waitpid(-1, &status, 0);` -- same as above

Zombie process

- ✓ A child that terminates, but has not been waited for becomes a "zombie"
- ✓ The kernel maintains a minimal set of information about the zombie process
 - PID, termination status, resource usage information
- ✓ As long as a zombie is not removed from the system via `wait`, it will consume a slot in the kernel process table
- ✓ If a parent process terminates, then its "zombie" children (if any) are adopted by *init*, which automatically performs a `wait` to remove the zombies

Creating zombies...

```
#include <sys/types.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
void worker() {
    printf("[%d] Hi, I'm a worker process! Going to die...\n",
        getpid());
}
int main()
{ int i;
  for (i=0; i<5; i++) {
    if (fork() == 0) {
        worker();
        exit(0);
    }
  }
  system("ps aux | grep -i defunct");
  printf("[%d] Big father is sleeping!\n", getpid());
  sleep(10);
  return 0;
}
```



Question: how many
processes are created?

Creating zombies...

✓ Results

```
user@ubuntu: ~/SO
File Edit Tabs Help
user@ubuntu:~/SO$ ./zombie.exe
[17512] Hi, I'm a worker process! Going to terminate...
[17513] Hi, I'm a worker process! Going to terminate...
[17514] Hi, I'm a worker process! Going to terminate...
[17511] Hi, I'm a worker process! Going to terminate...
[17510] Hi, I'm a worker process! Going to terminate...
user      17467  0.2  0.4  12348  4616 pts/16  S+   12:59   0:00 vim zombie.c
user      17509  0.0  0.0   2024   276 pts/5    S+   13:02   0:00 ./zombie.exe
user      17510  0.0  0.0      0      0 pts/5    Z+   13:02   0:00 [zombie.exe] <defunct>
user      17511  0.0  0.0      0      0 pts/5    Z+   13:02   0:00 [zombie.exe] <defunct>
user      17512  0.0  0.0      0      0 pts/5    Z+   13:02   0:00 [zombie.exe] <defunct>
user      17513  0.0  0.0      0      0 pts/5    Z+   13:02   0:00 [zombie.exe] <defunct>
user      17514  0.0  0.0      0      0 pts/5    Z+   13:02   0:00 [zombie.exe] <defunct>
user      17515  0.0  0.0   2268   552 pts/5    S+   13:02   0:00 sh -c ps aux | grep -i zombie
user      17517  0.0  0.0   4680   832 pts/5    S+   13:02   0:00 grep -i zombie
[17509] Big father is sleeping!
user@ubuntu:~/SO$
```


The system function

- ✓ In the zombie code, we have the following line of code

```
system("ps aux | grep -i zombie");
```
- ✓ system launches a shell that executes the command given as string
 - It executes `/bin/sh -c command_line` and waits for its termination
 - The Shell process is the one that actually executes the command line
- It is costly, since it has to i) `fork` a process and then ii) `exec` its image to execute the Shell

CPU affinity (#1)

- ✓ On a multicore system, the OS scheduler needs to decide which processes runs on each CPU
- ✓ Once a process is running on a CPU, the OS scheduler tries to keep it there
 - Avoid the “cold cache” effect of moving a process to another CPU/core
 - When a process moves to another CPU/core, the cache(s) of the CPU/core do not have content of the process



<http://bit.ly/256cAlr>

CPU affinity (#2)

- ✓ CPU affinity of a process can be controlled programatically
 - Hard affinity
- ✓ `int sched_setaffinity(pid_t pid, size_t setsize, const cpu_set_t *set);`
- ✓ `int sched_getaffinity(pid_t pid, size_t setsize, cpu_set_t *set);`
- ✓ `void CPU_SET(unsigned long cpu, cpu_set_t *set);`
- ✓ `void CPU_CLR(unsigned long cpu, cpu_set_t *set);`
- ✓ `int CPU_ISSET(unsigned long cpu, cpu_set_t *set);`
- ✓ `void CPU_ZERO(cpu_set_t *set);`

✓ Example

```
#define _GNU_SOURCE
#include <sched.h>
#include <stdio.h>
cpu_set_t set;
int ret, i;
CPU_ZERO (&set);
ret = sched_getaffinity(0, sizeof (cpu_set_t), &set);
if (ret == -1) {
    perror ("sched_getaffinity");
}
for (i=0; i < CPU_SETSIZE; i++) {
    int cpu;
    cpu = CPU_ISSET(i, &set);
    printf ("cpu=%i is %s\n", i, cpu?"set":"unset");
}
```

Termination of a process

- ✓ Reason for a process to terminate
 - Regular termination
 - exit, return of main function, etc.
 - Process has exceeded maximum CPU time (e.g., “ulimit” from bash)
 - Not enough memory
 - I/O failure
 - Invalid instruction (e.g., “divide by zero”)
 - OS action
 - Deadlock or OOM (Out of Memory Killer)
 - User action
 - Kill -9 PID or killall -9 process_name
 - ...

ulimits (bash)

✓ The bash shell has an internal set of limits

– `ulimit`

- internal command

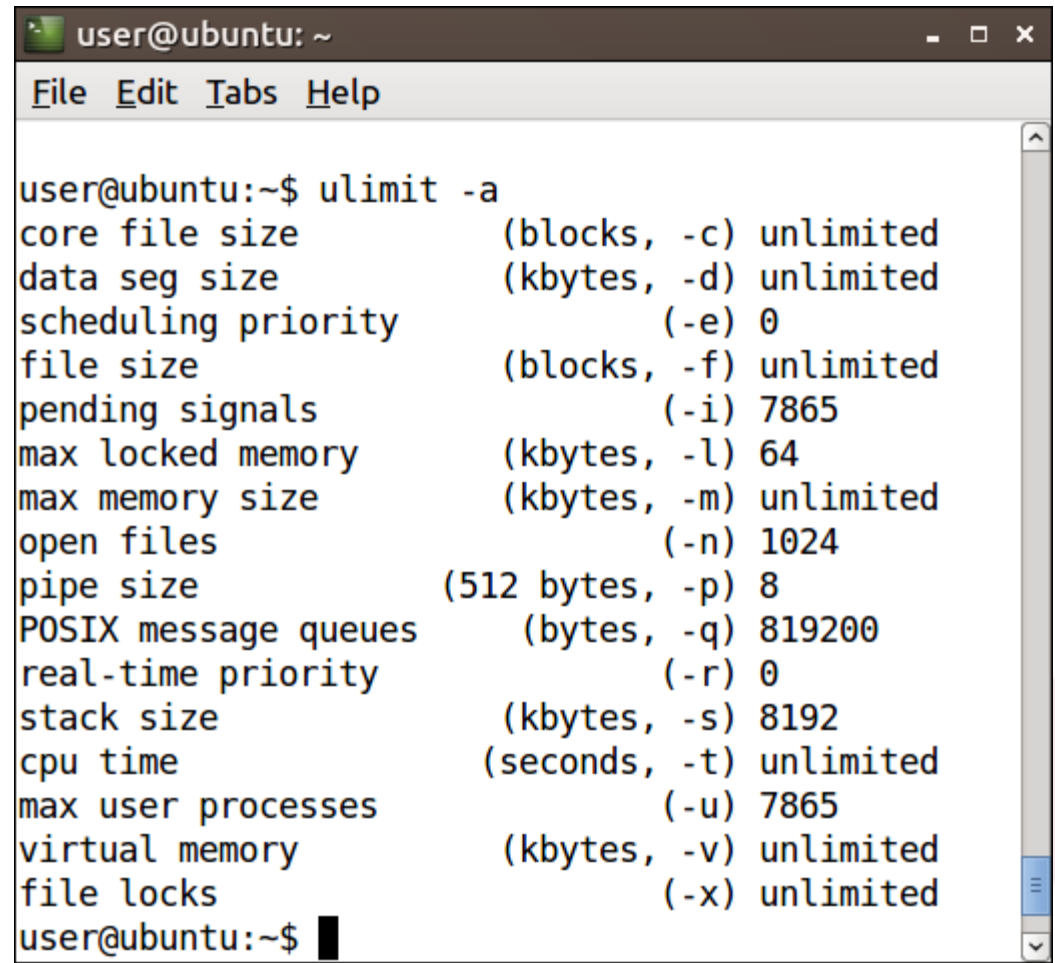
- Processed by bash, there is no `ulimit` executable

- There is no man for `ulimit`

- `help ulimit`

– `ulimit -a`

- List all limits for the current session



```
user@ubuntu: ~  
File Edit Tabs Help  
  
user@ubuntu:~$ ulimit -a  
core file size          (blocks, -c) unlimited  
data seg size           (kbytes, -d) unlimited  
scheduling priority     (-e) 0  
file size               (blocks, -f) unlimited  
pending signals         (-i) 7865  
max locked memory       (kbytes, -l) 64  
max memory size         (kbytes, -m) unlimited  
open files              (-n) 1024  
pipe size               (512 bytes, -p) 8  
POSIX message queues    (bytes, -q) 819200  
real-time priority      (-r) 0  
stack size              (kbytes, -s) 8192  
cpu time                (seconds, -t) unlimited  
max user processes      (-u) 7865  
virtual memory          (kbytes, -v) unlimited  
file locks              (-x) unlimited  
user@ubuntu:~$
```

ulimits and core files

- ✓ ulimits has a parameter that controls the size of core file
 - `ulimit -c 999999999` → unlimited
- ✓ With unlimited, the crash of an executable generates a core dump, that is, a file named “core”
- ✓ This core can be examined with a debugger, namely gdb
 - `gdb a.exe core`
 - The program needs to be compiled with the right options
 - `-g`
 - `-ggdb`

- Man pages
 - man 2 fork
 - man 2 exec
 - man system
 - man bash
 - help ulimit
- *Chapter 5 – Process management, “Linux System Programming”, Robert Love, 2013*
- printf format in C99 and C11
<http://bit.ly/1OvGzGI>

